

Distributed Systems

Spring Semester 2020

Lecture 8: Primary-Backup Replication

John Liagouris
liagos@bu.edu

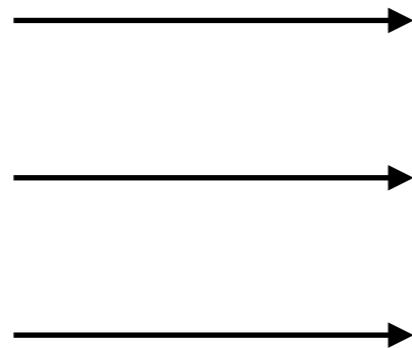
Replication

- Replication for Fault Tolerance
 - Ideally : System keeps work even when things fail
 - Consider Master of Map Reduce
 - Single point of Failure — 100's of workers useless if Master fails
 - Like to build systems without single points of failure
 - Several approaches but perhaps easiest to consider — general purpose — replicated service



Basic idea

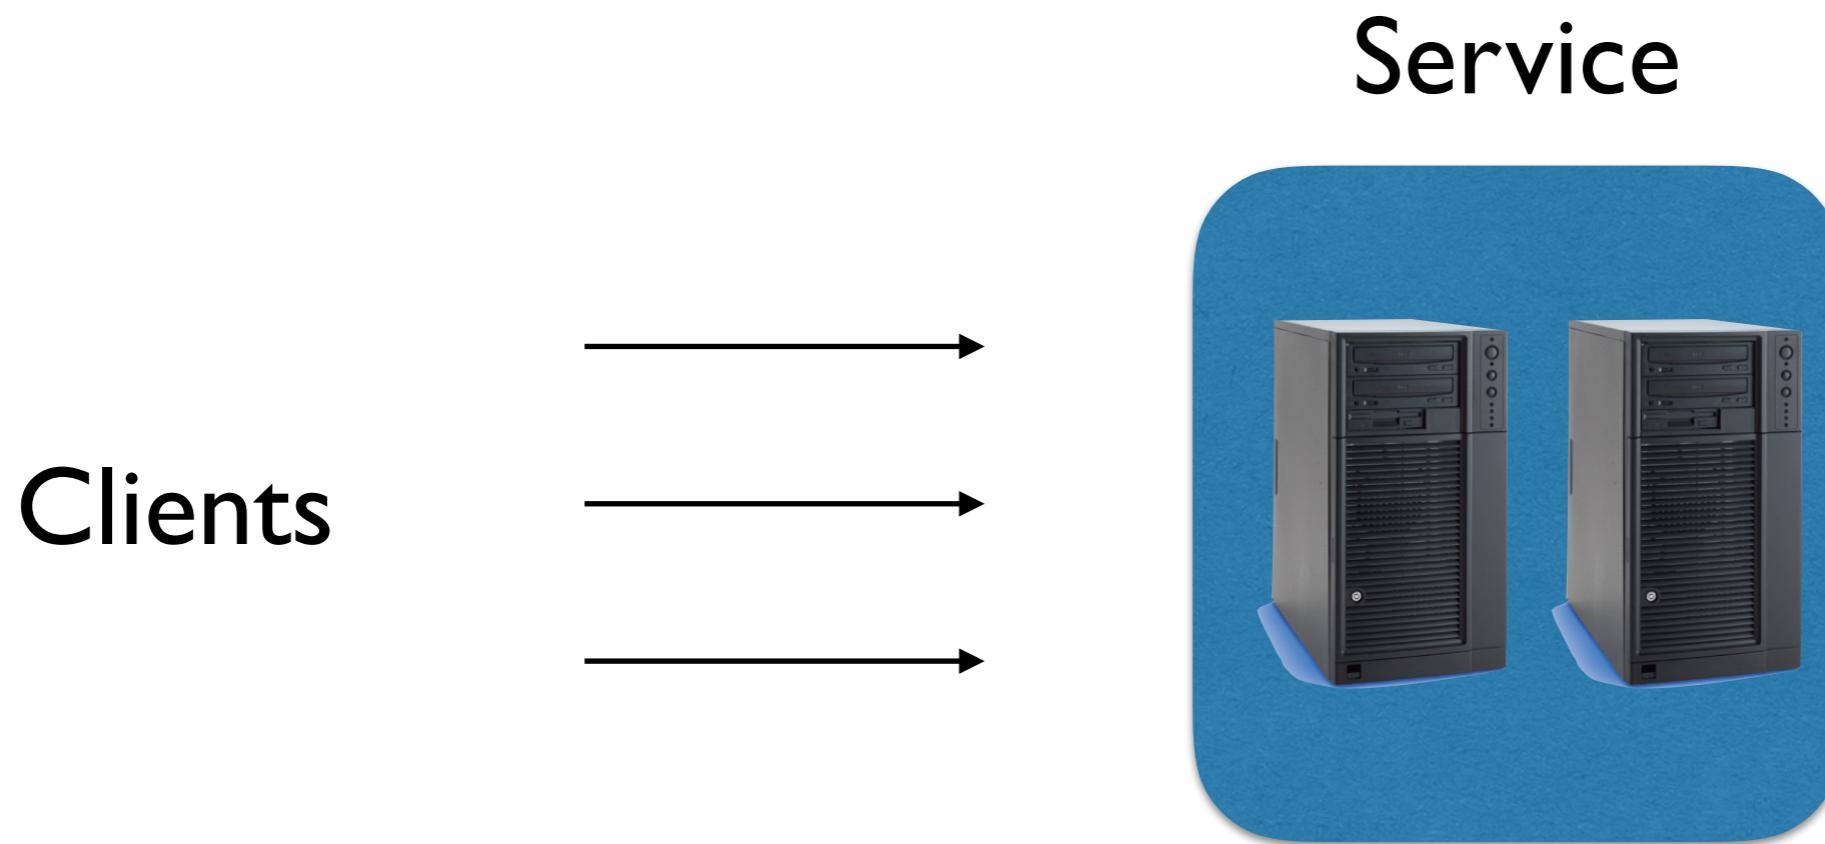
Clients



Service

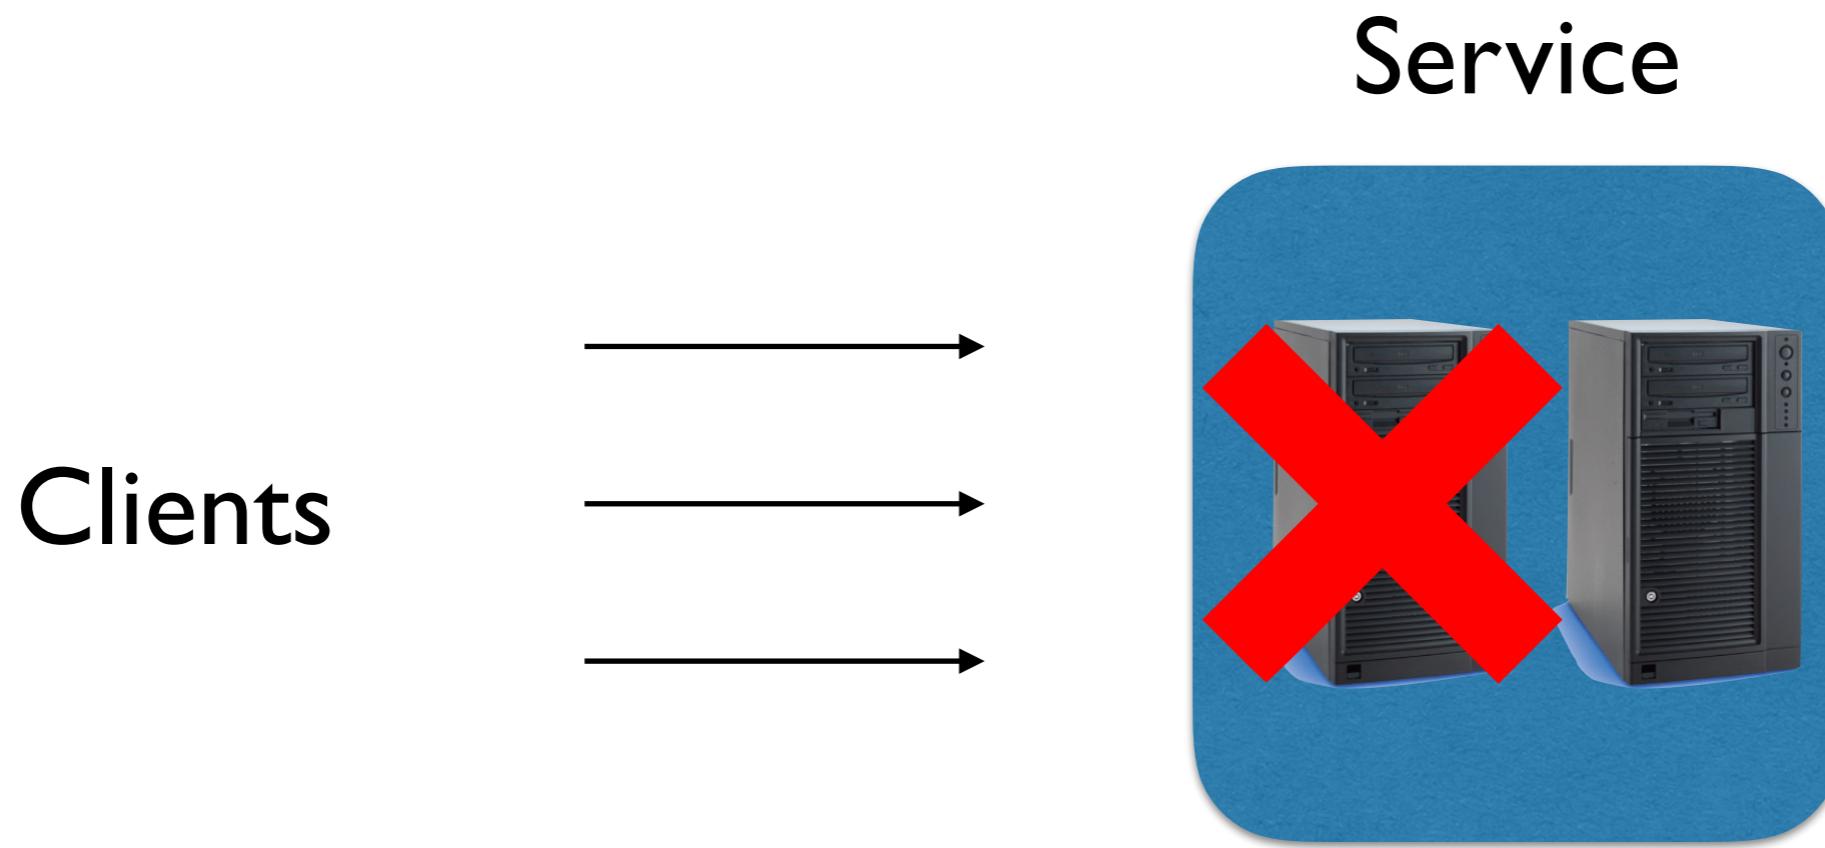


Basic Idea



Two computers run your service

Replication



Some how arrange for the second computer to take over from the first if it ever suffers a failure

High Level — Two things we are after

- Available — if first system fails would like the second to take over and pickup immediately while it make take hours to fix the first
- Correct in the face of replication — lots of difficulties around correctness

Correctness

- Basic Problem:
 - When we replicate — create a copy
 - very hard to avoid a window in which they do not match — ensure same ALL the time
 - So very hard to ensure that clients never notice that a failure has happened and cut-over occurred

Total transparency is very difficult — more often than not weaken interface to clients to make it easier

Failure Model

Never going to build a system that will run no matter what
Rather it will survive this particular set of faults and not
this other set

- Independent Fail Stop
 - Fail Stop: Assume computers execute correctly for awhile then halt — very different than considering incorrect execution and non-halting
 - Independent: The failure of one computer does not change the likelihood of the failure of another — necessary assumption for replication to be sensible

Are failures independent?

Perhaps the most common

correlate failure

— Site wide power failure —

Failure Model

The two we focus on

- Independent Fail Stop
- Site-wide power failures

Most systems have one story for Independent failures and another for Site-wide power failures usually involving disk drives

Big Questions to consider as we look at systems

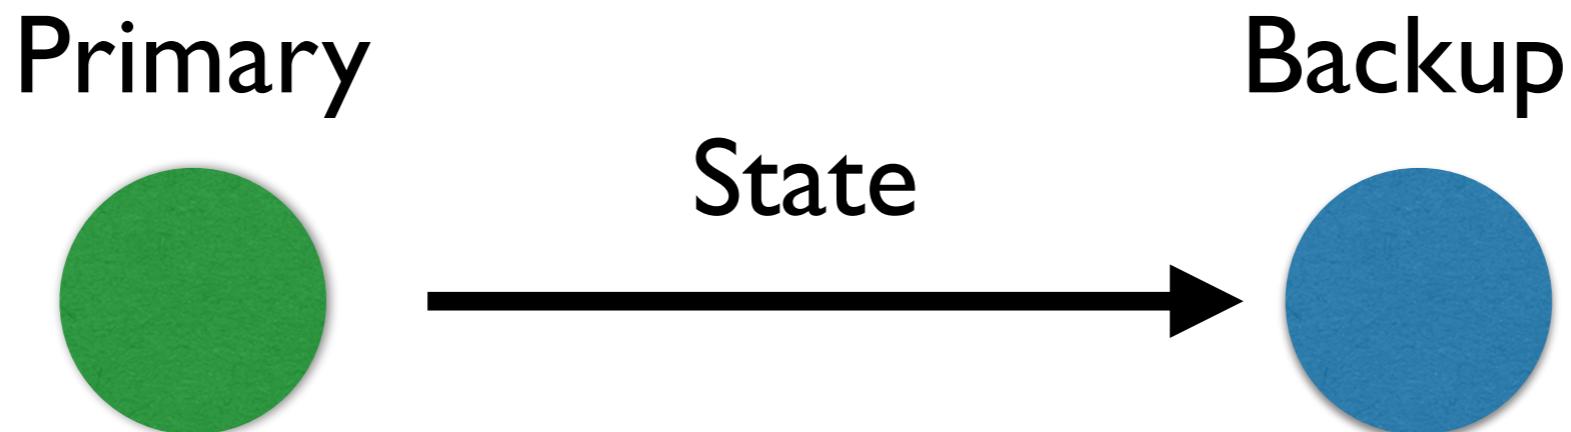
- 1.What state is replicated? — back up has to have some state so it can take over.
- 2.How do the replicas get state?
- 3.When to cut-over? can't directly inspect computers can only use network messages — failure detection — fraught with difficulty
- 4.What Anomalies are visible to the clients if cut-over occurs? — lost requests, duplicate requests, ...
- 5.How is re-integration handled? — restore a failed system after repair

We are going to
organize systems into
two main types

Types

I. State Transfer Systems

- Primary execute and updates its state and periodically transfers to backups



Types

2. Replicated State Machines

- Both Primary and Backup execute the requests/operations and their state evolves consistently



Clearly some assumptions : deterministic operations, operations are done in the same order, etc

Why might we prefer
Replicated State
Machines (RSM) at
times?

State transfer is simpler
But state maybe large
— slow to transfer

RSM can be more efficient

— if operations are small
compared to data

But complex to get right

Lots of computations is **NOT** Deterministic

Must take special care to restrict and or ensure
determinism

RSM can be more efficient

— if operations are small
compared to data

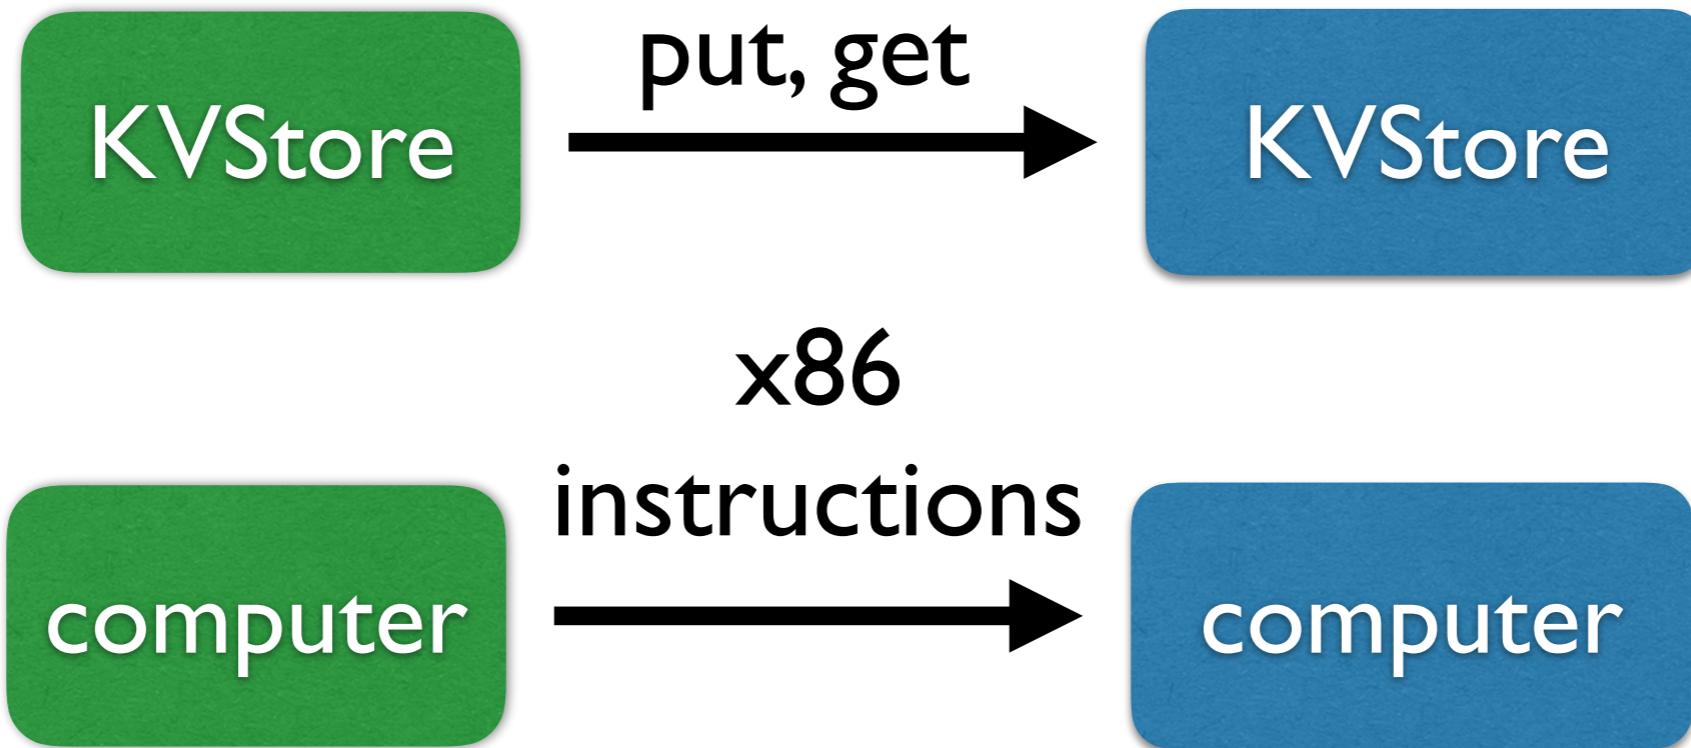
But complex to get right

Lots of computations is **NOT** Deterministic

Must take special care to restrict and or ensure
determinism

In Practice Replicated State
Machines are more
common approach —
lighter weight
vm-ft and Labs uses RSM

RSM

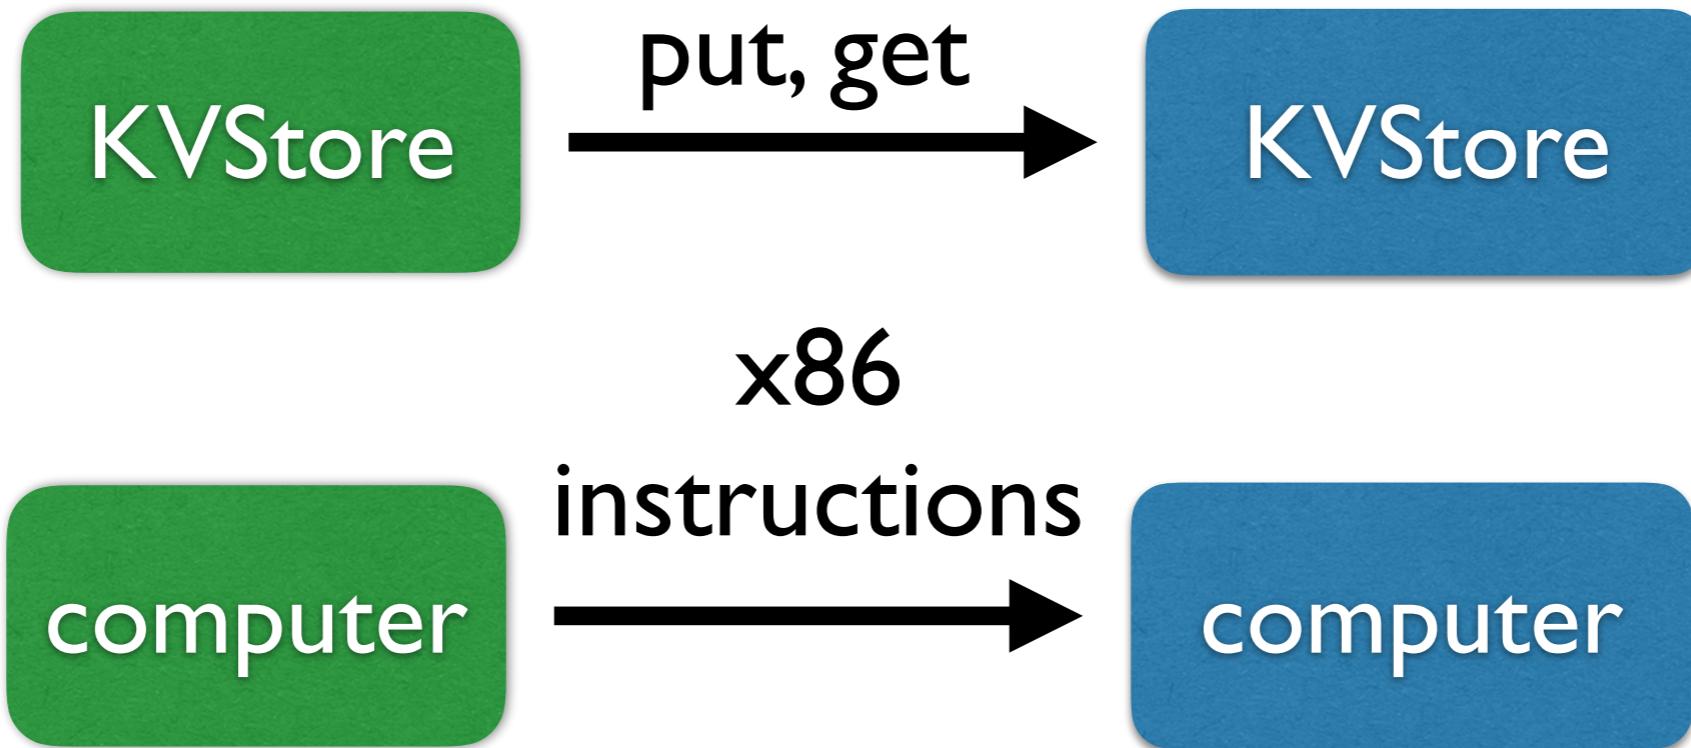


Choices impact 1) performance and 2) ease of implementation

How much interaction is necessary

Dealing with non-determinism is hard on x86 level but easier on put/get level

RSM



Choices impact 1) performance and 2) ease of implementation

How much interaction is necessary

Dealing with non-determinism is hard on x86 level but easier on put/get level

BUT higher-level RSM
works only for apps
written to higher-level
interface

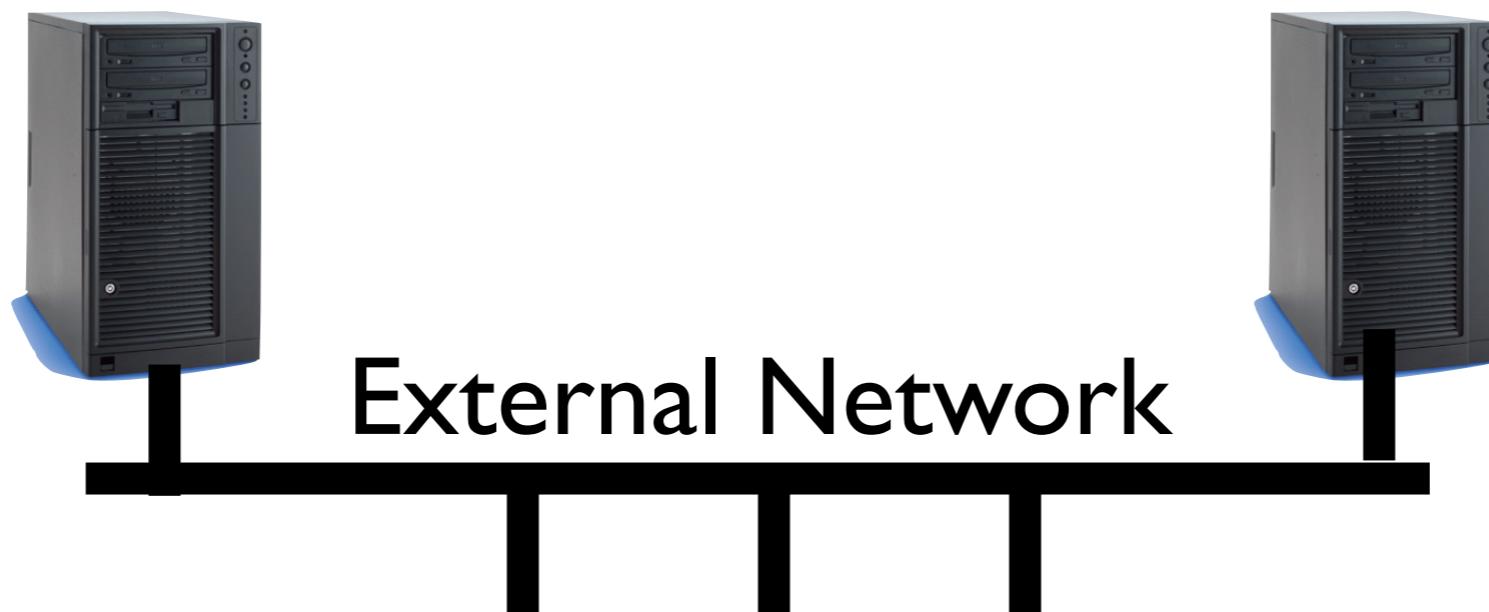
x86 RSM can make any
x86 program FT

VM-FT

- Super ambitious :
 - Take any software (OS and APP) and make it fault tolerant
 - Like a Magic Wand
- Costs are high but very cool given what it's goals are

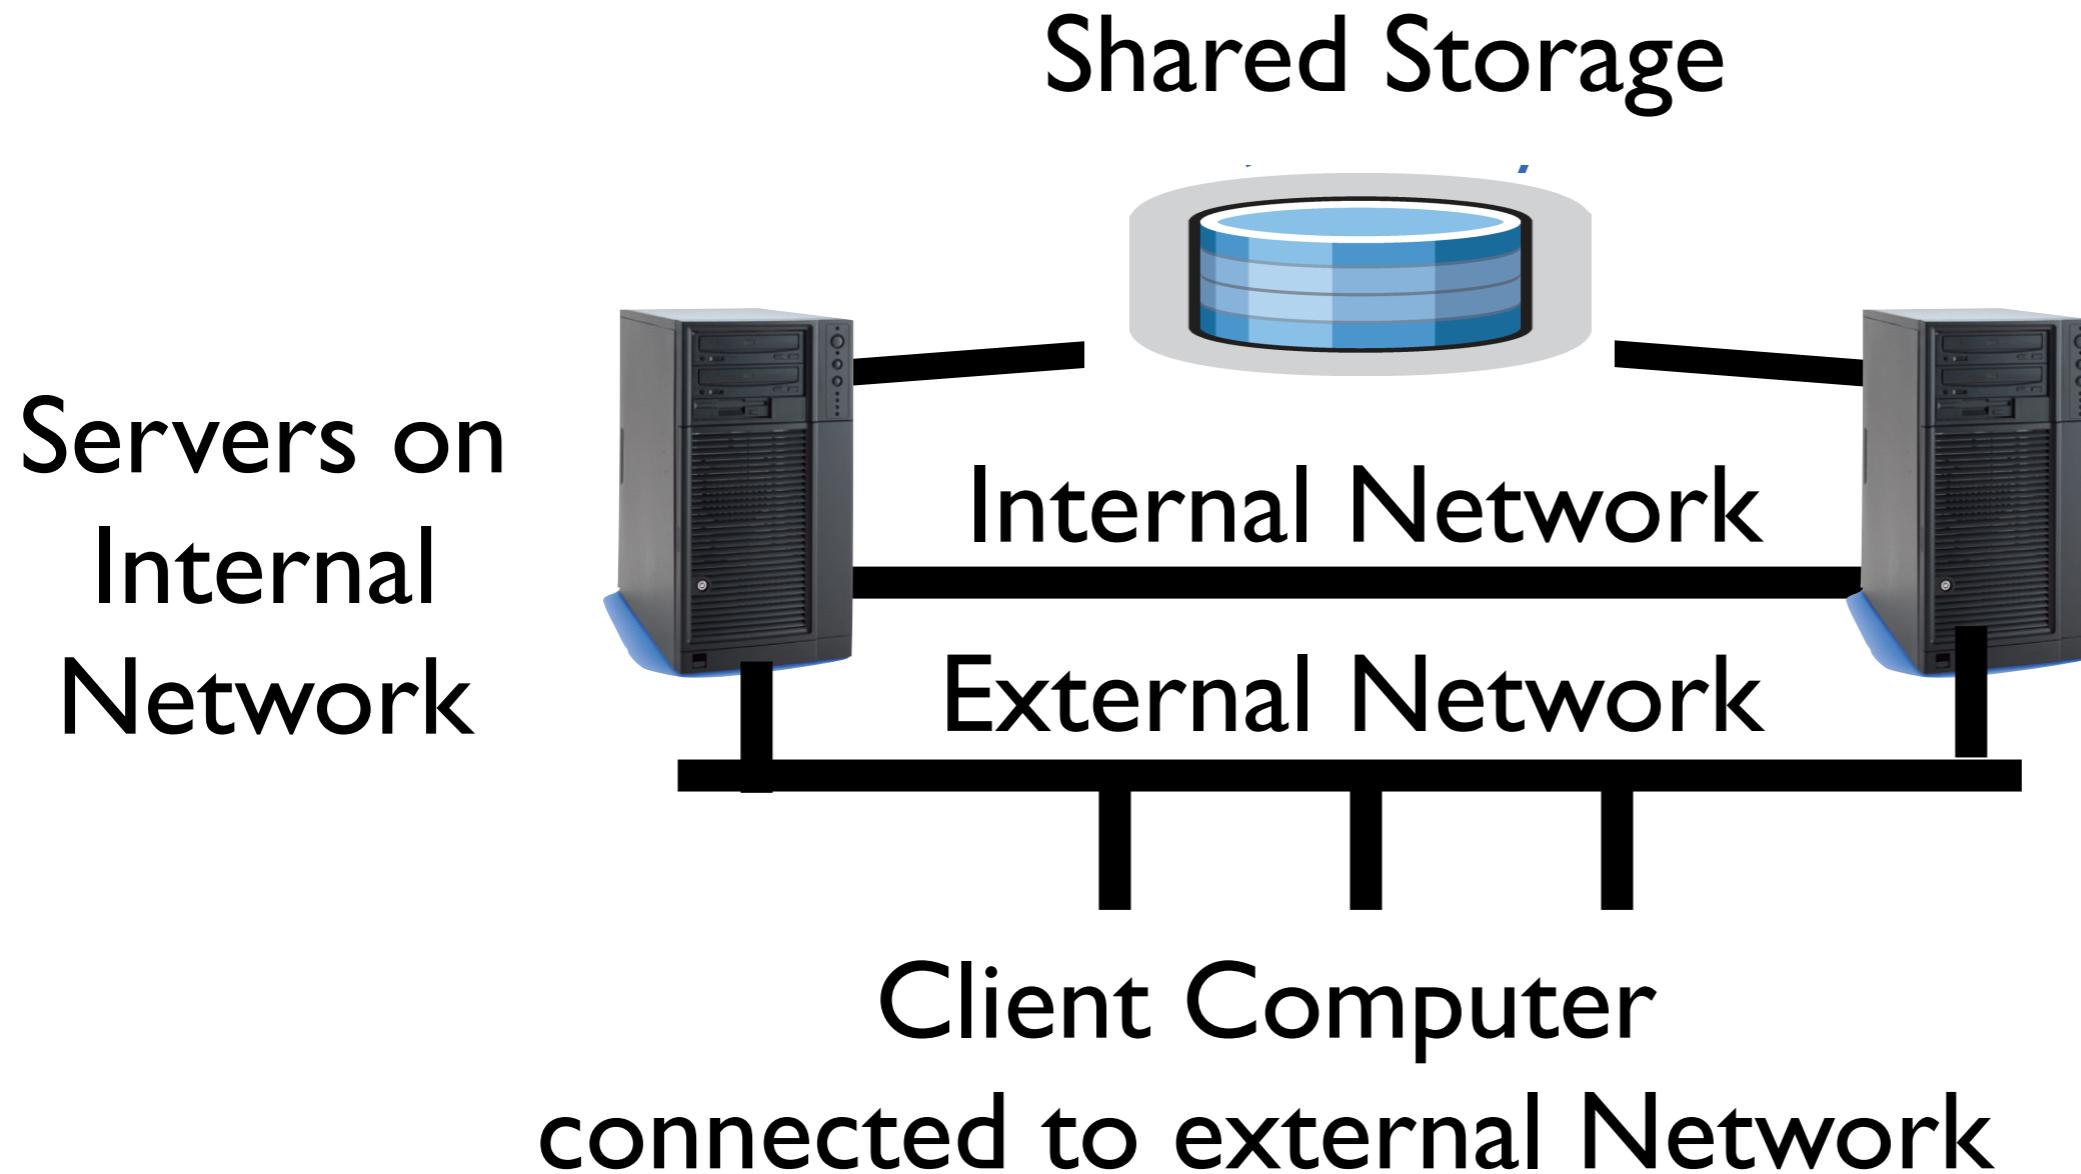
Overview

Servers on
Internal
Network

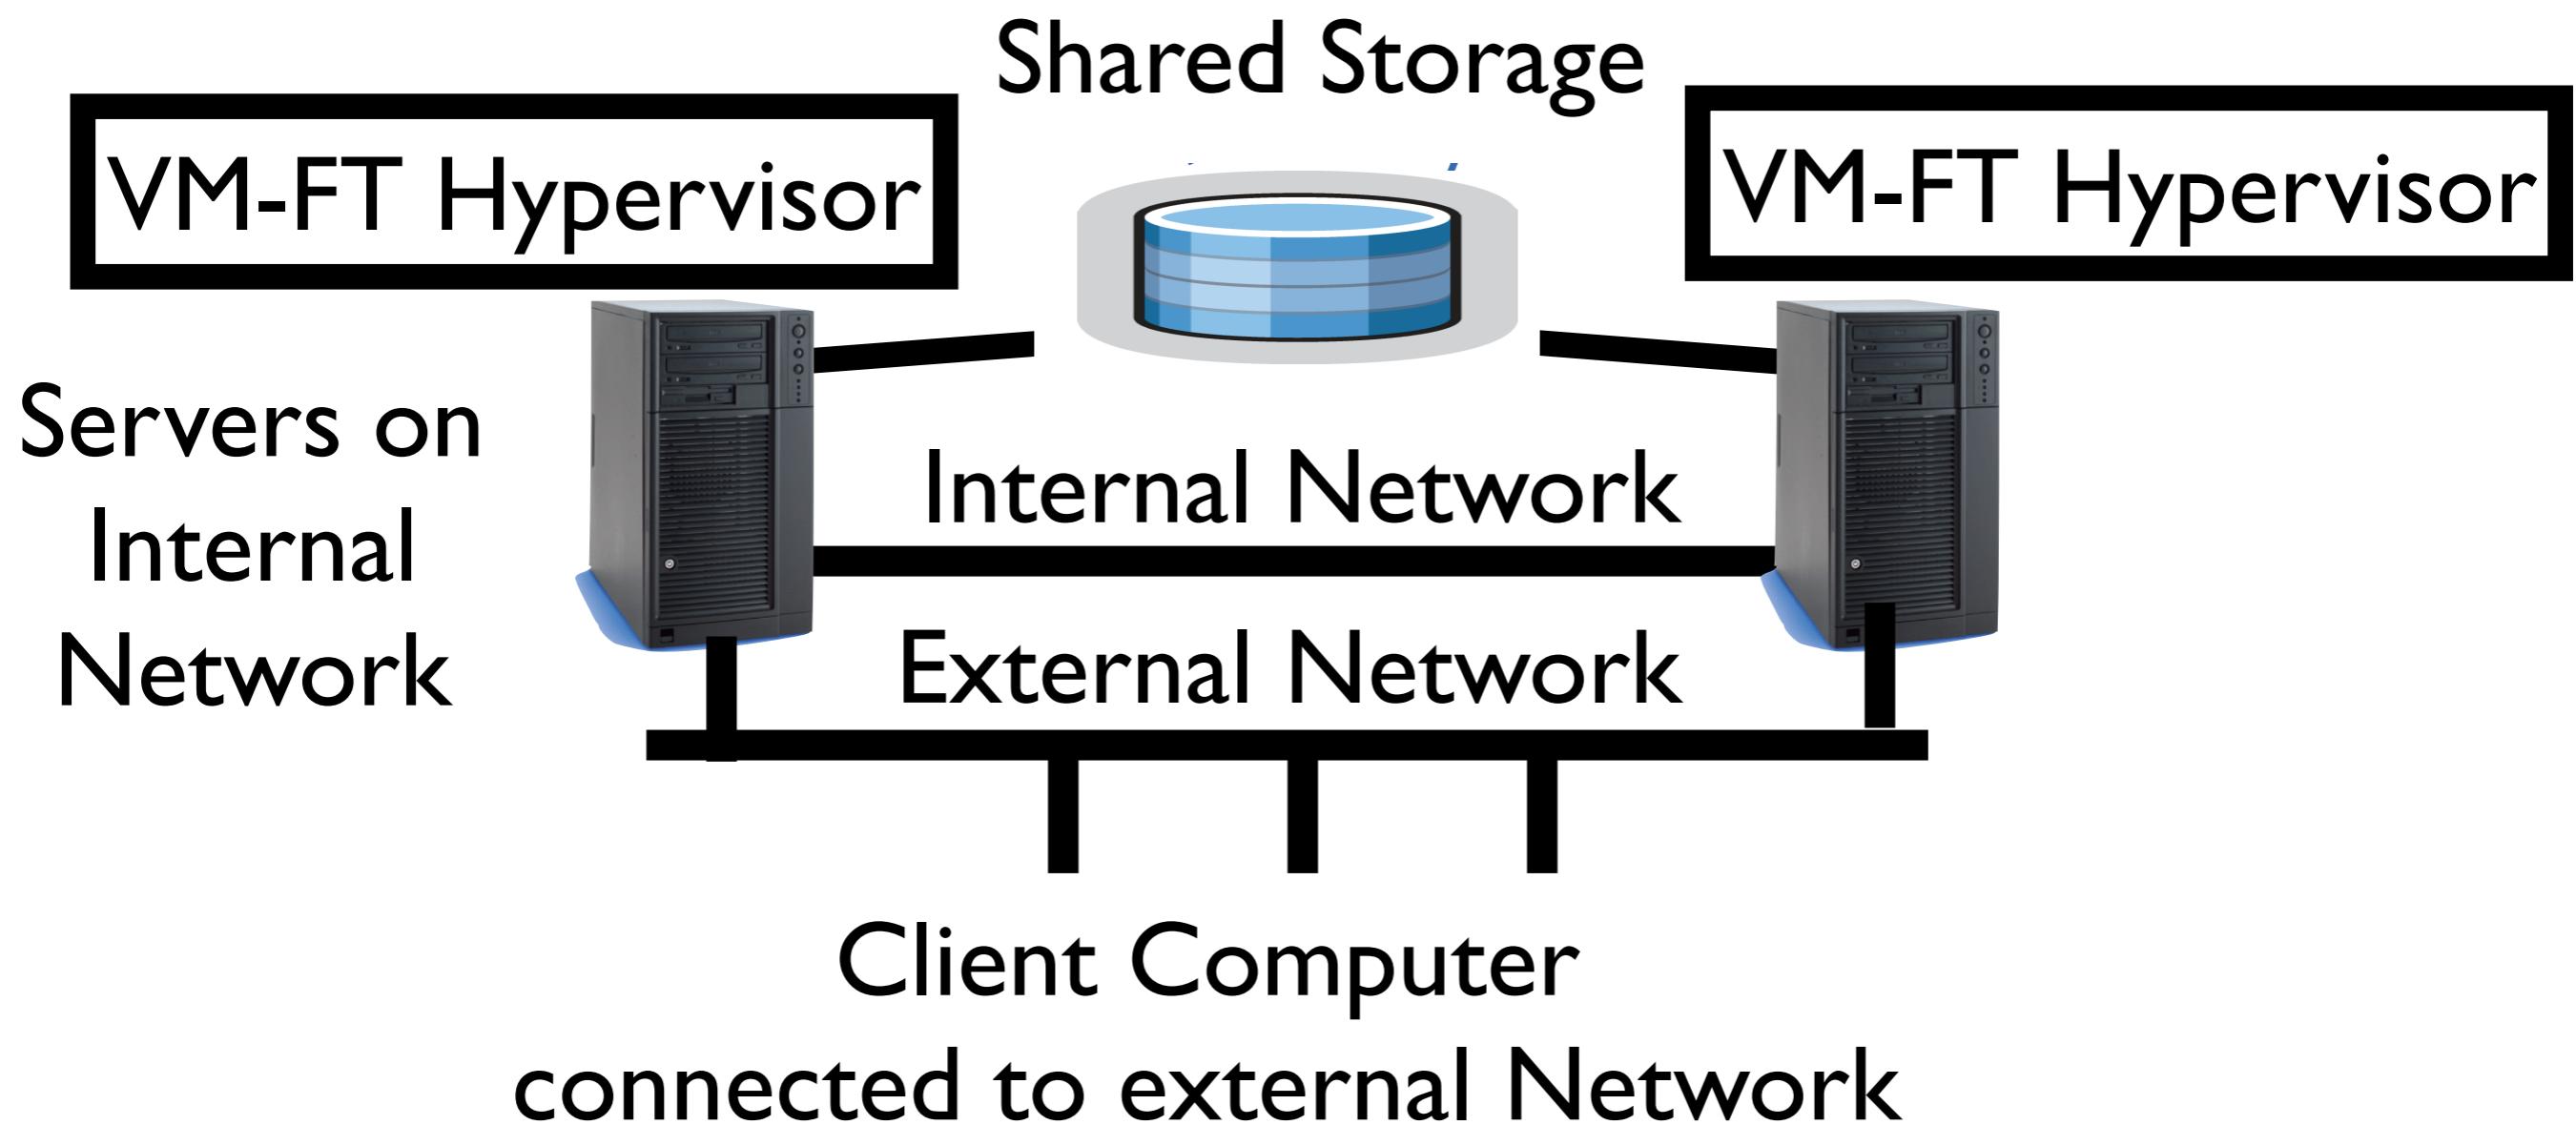


Client Computer
connected to external Network

Overview

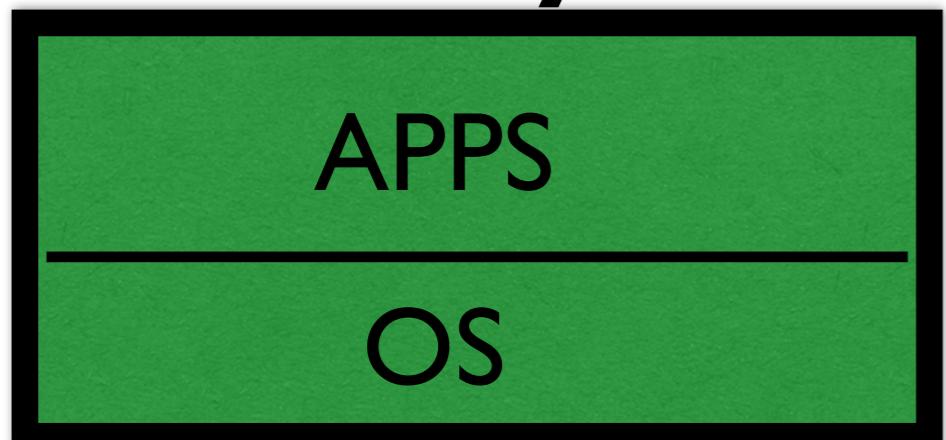


Overview

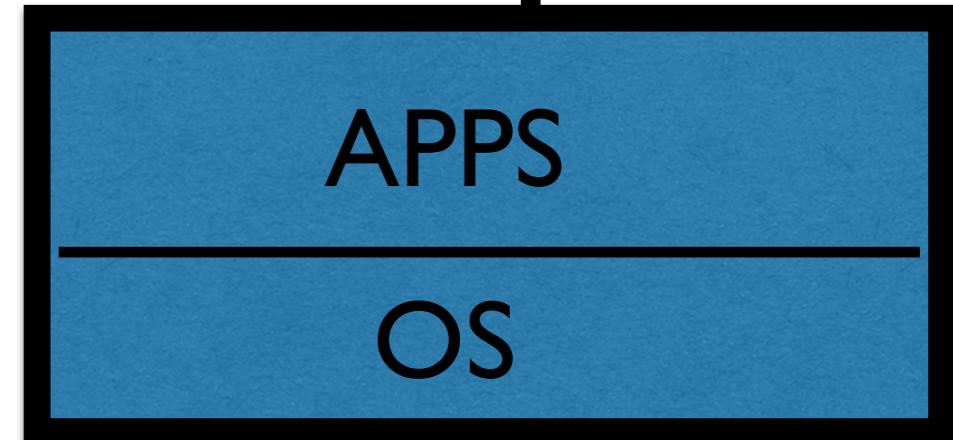


Overview

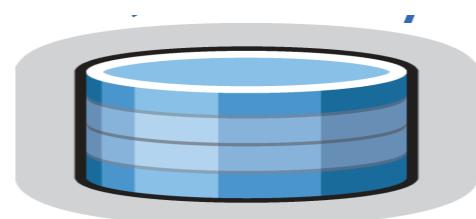
Primary VM



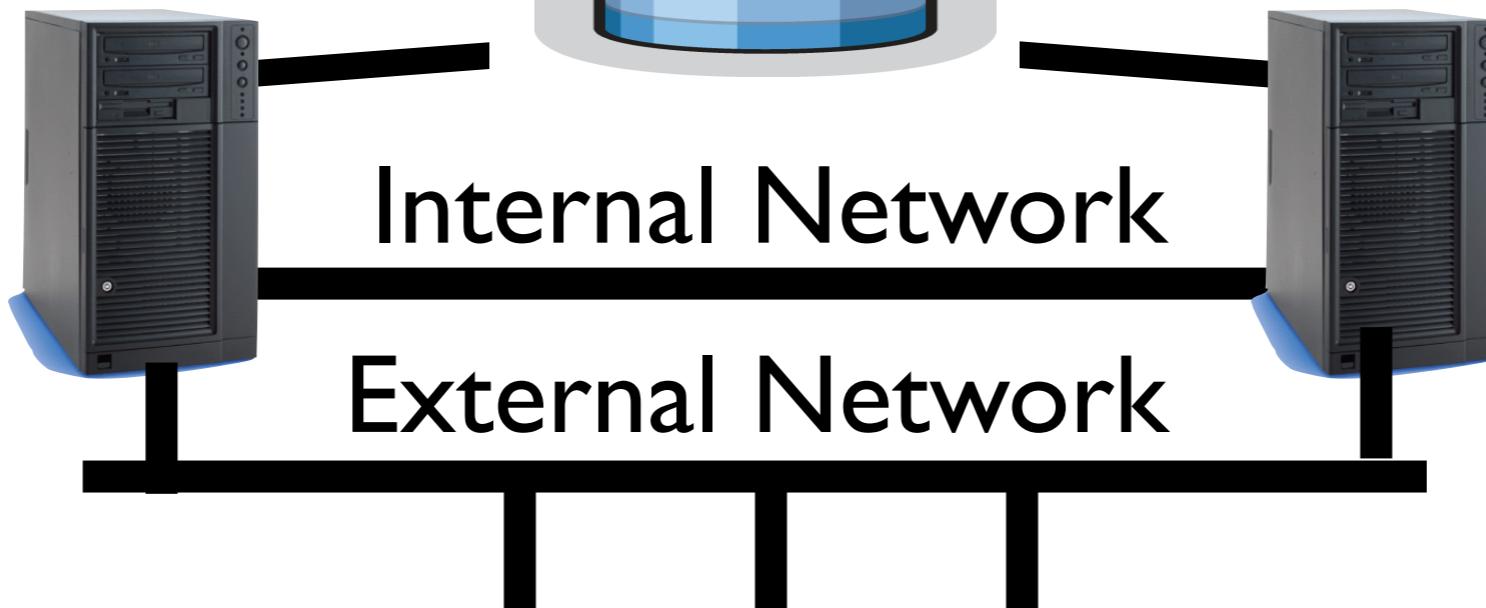
Backup VM



Shared Storage



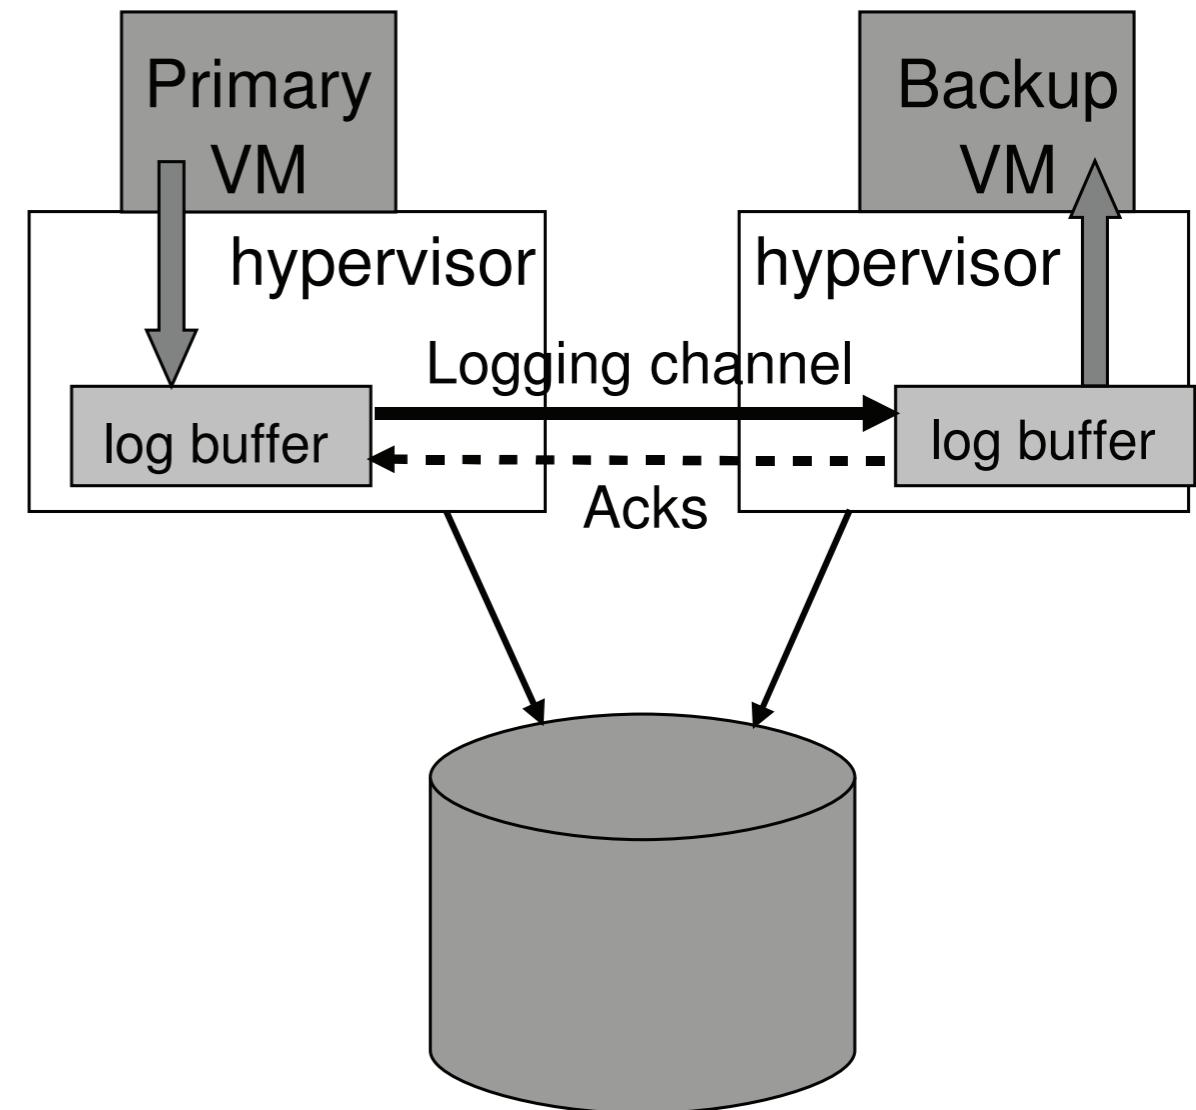
Servers on
Internal
Network



Client Computer
connected to external Network

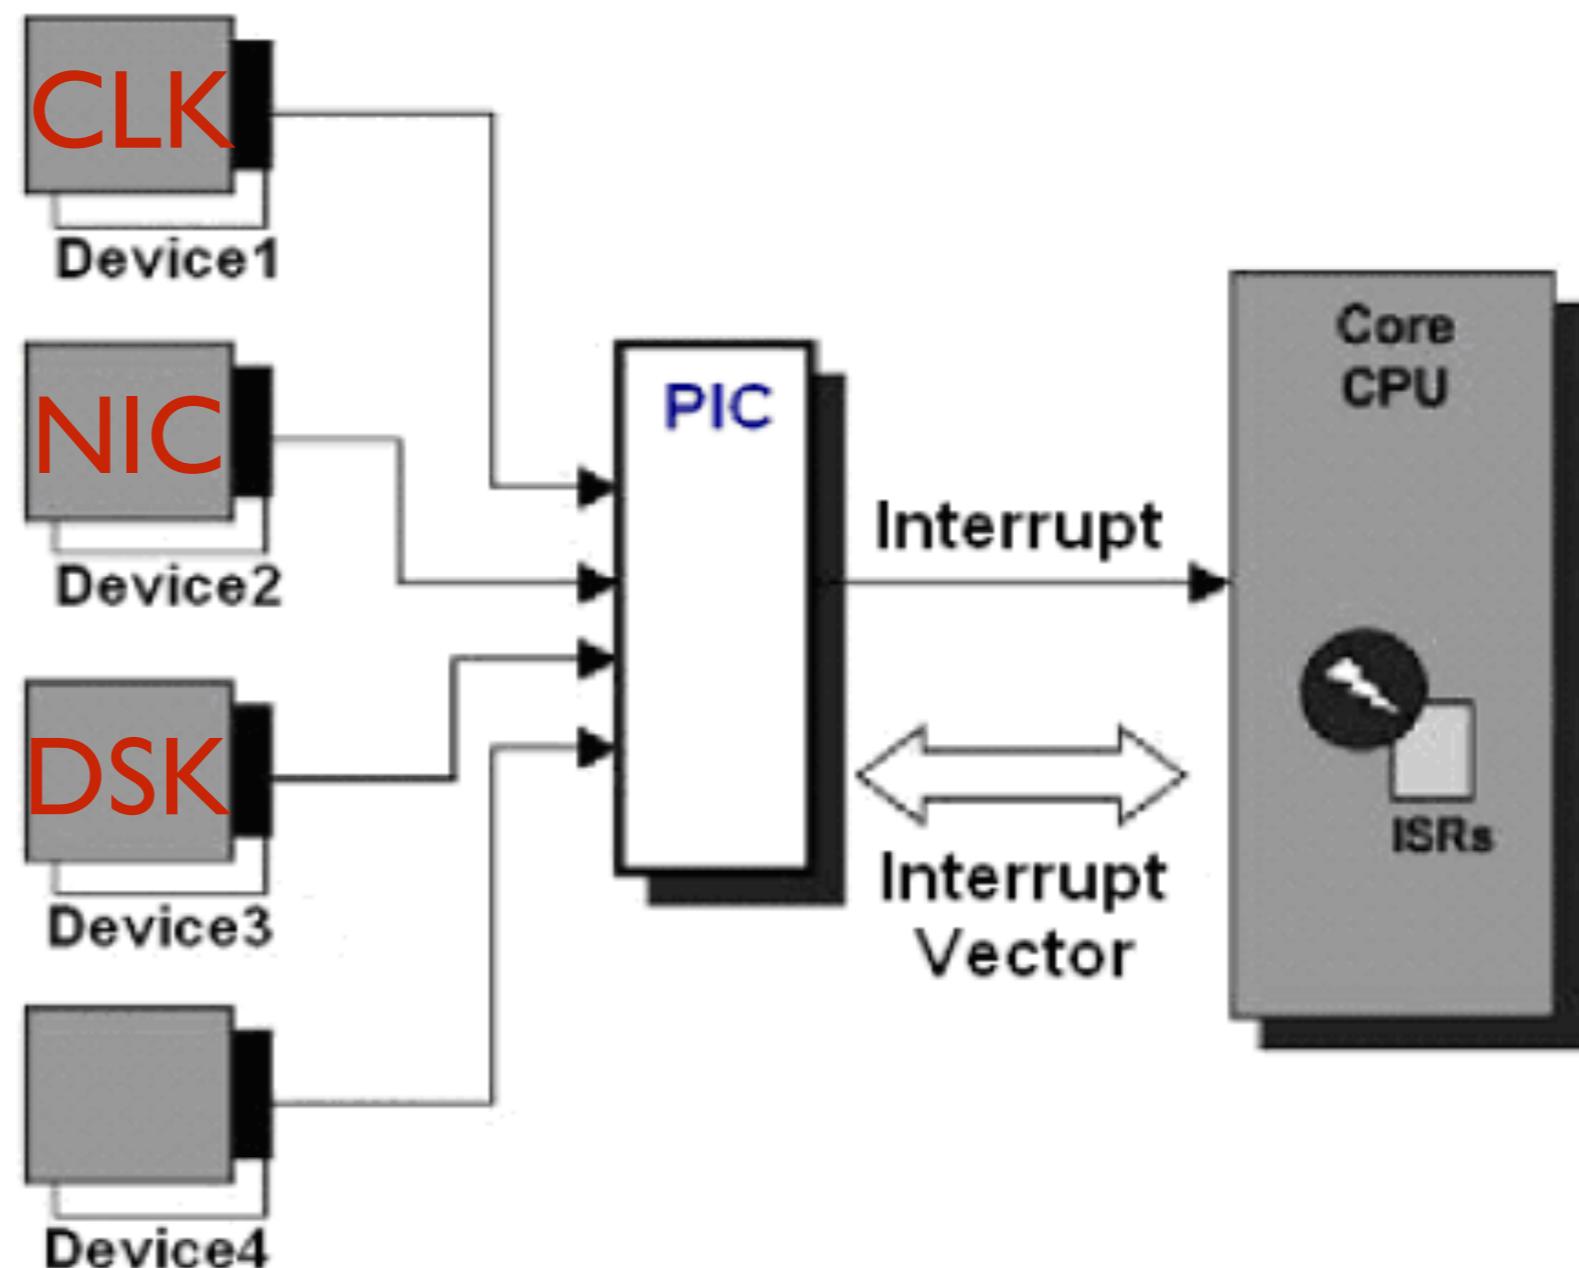
Overview

- Backup in “lock step” with primary
 - Primary sends all input to backup
 - output of backup is dropped
 - heart beats between primary and backup
 - if primary fails start backup executing



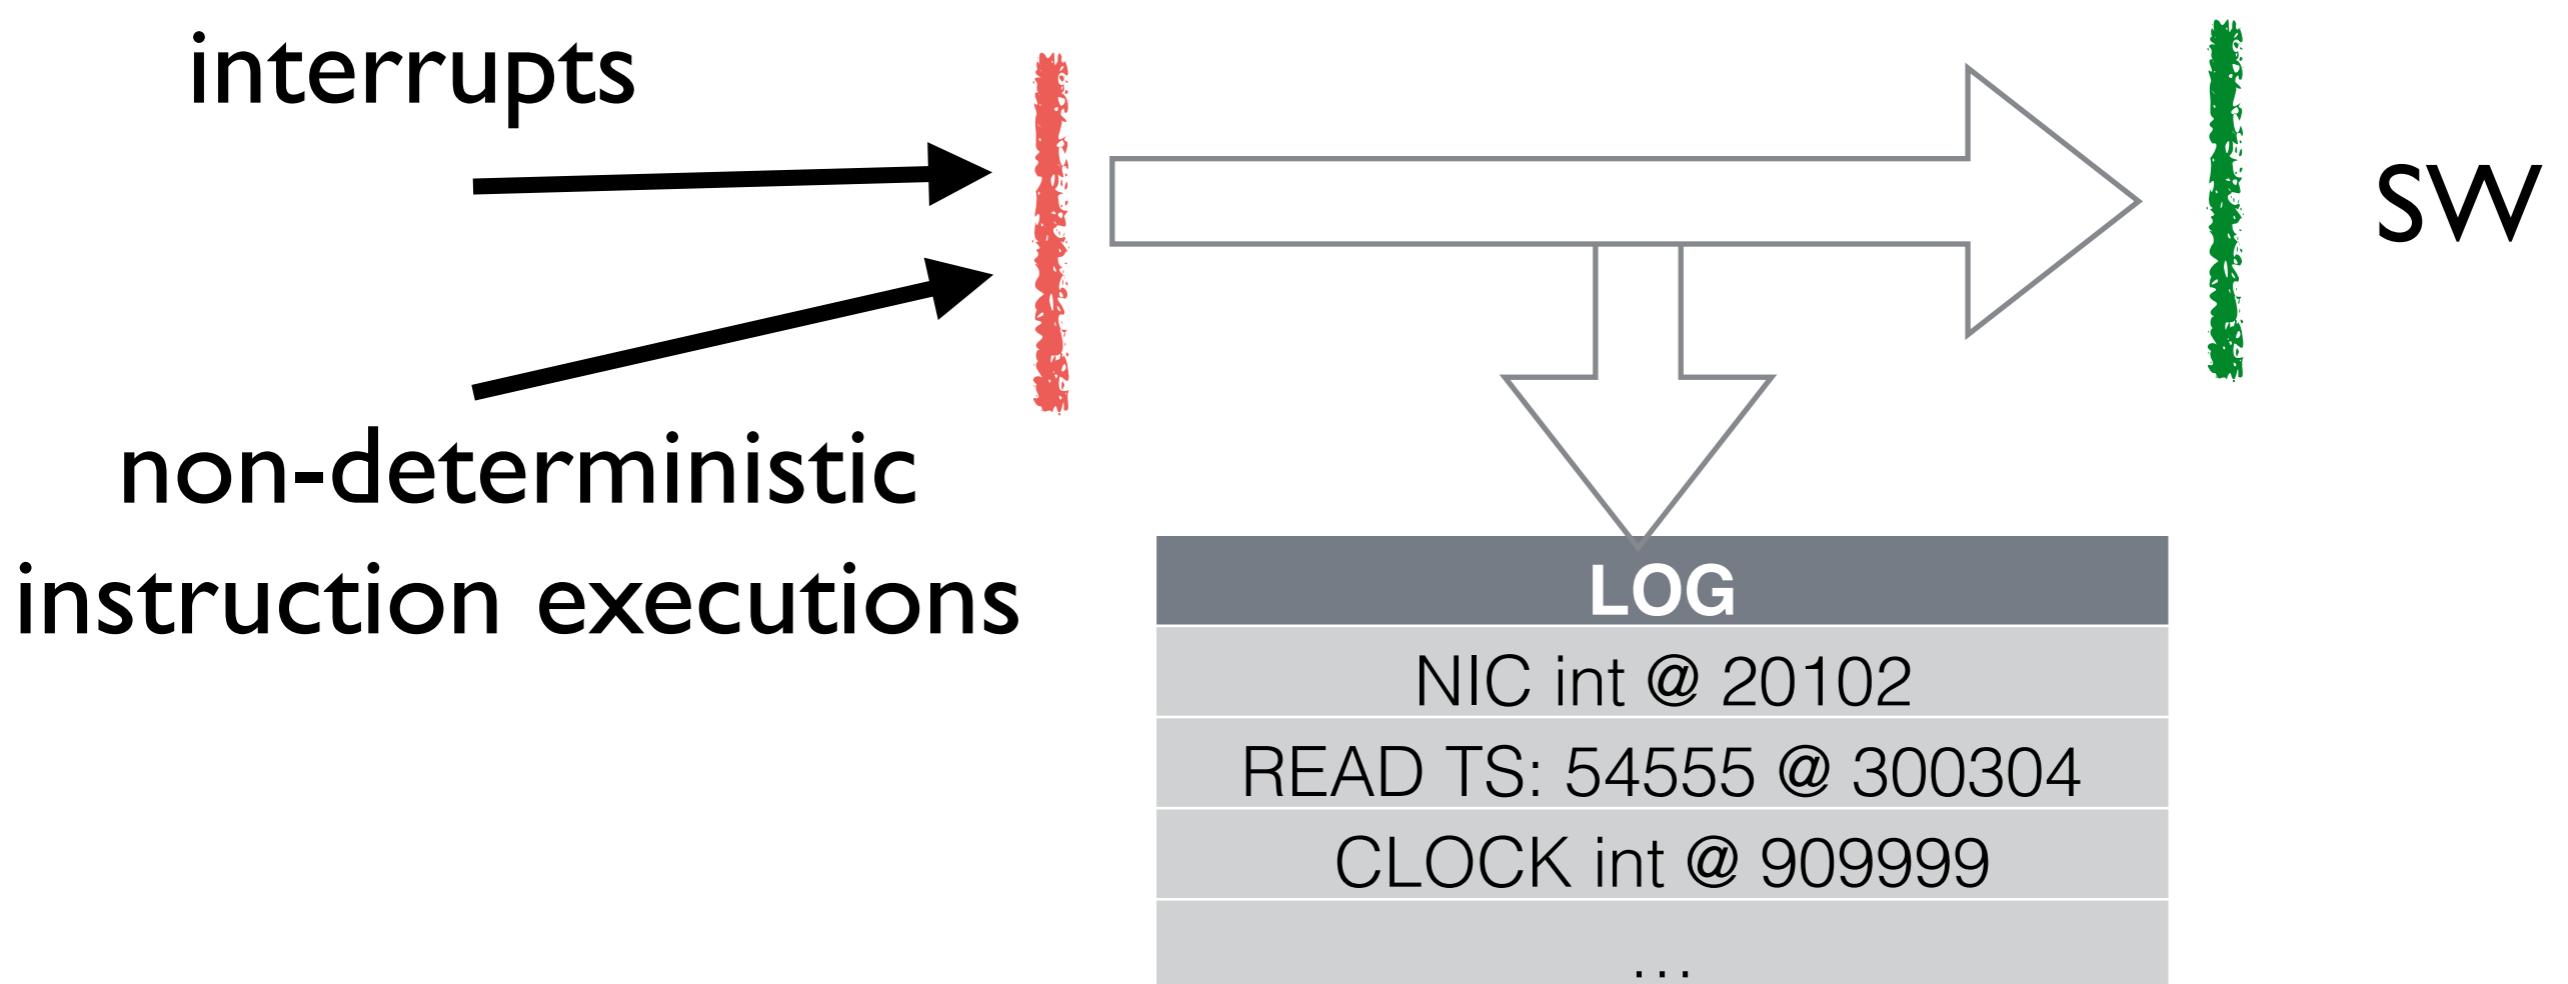
Inputs

- Clock interrupts
- Network interrupt
- Disk interrupts



Logging

- Log all hardware events into a log



EXAMPLE

APPS

OS

APPS

OS

VM-FT Hypervisor

NIC



INT

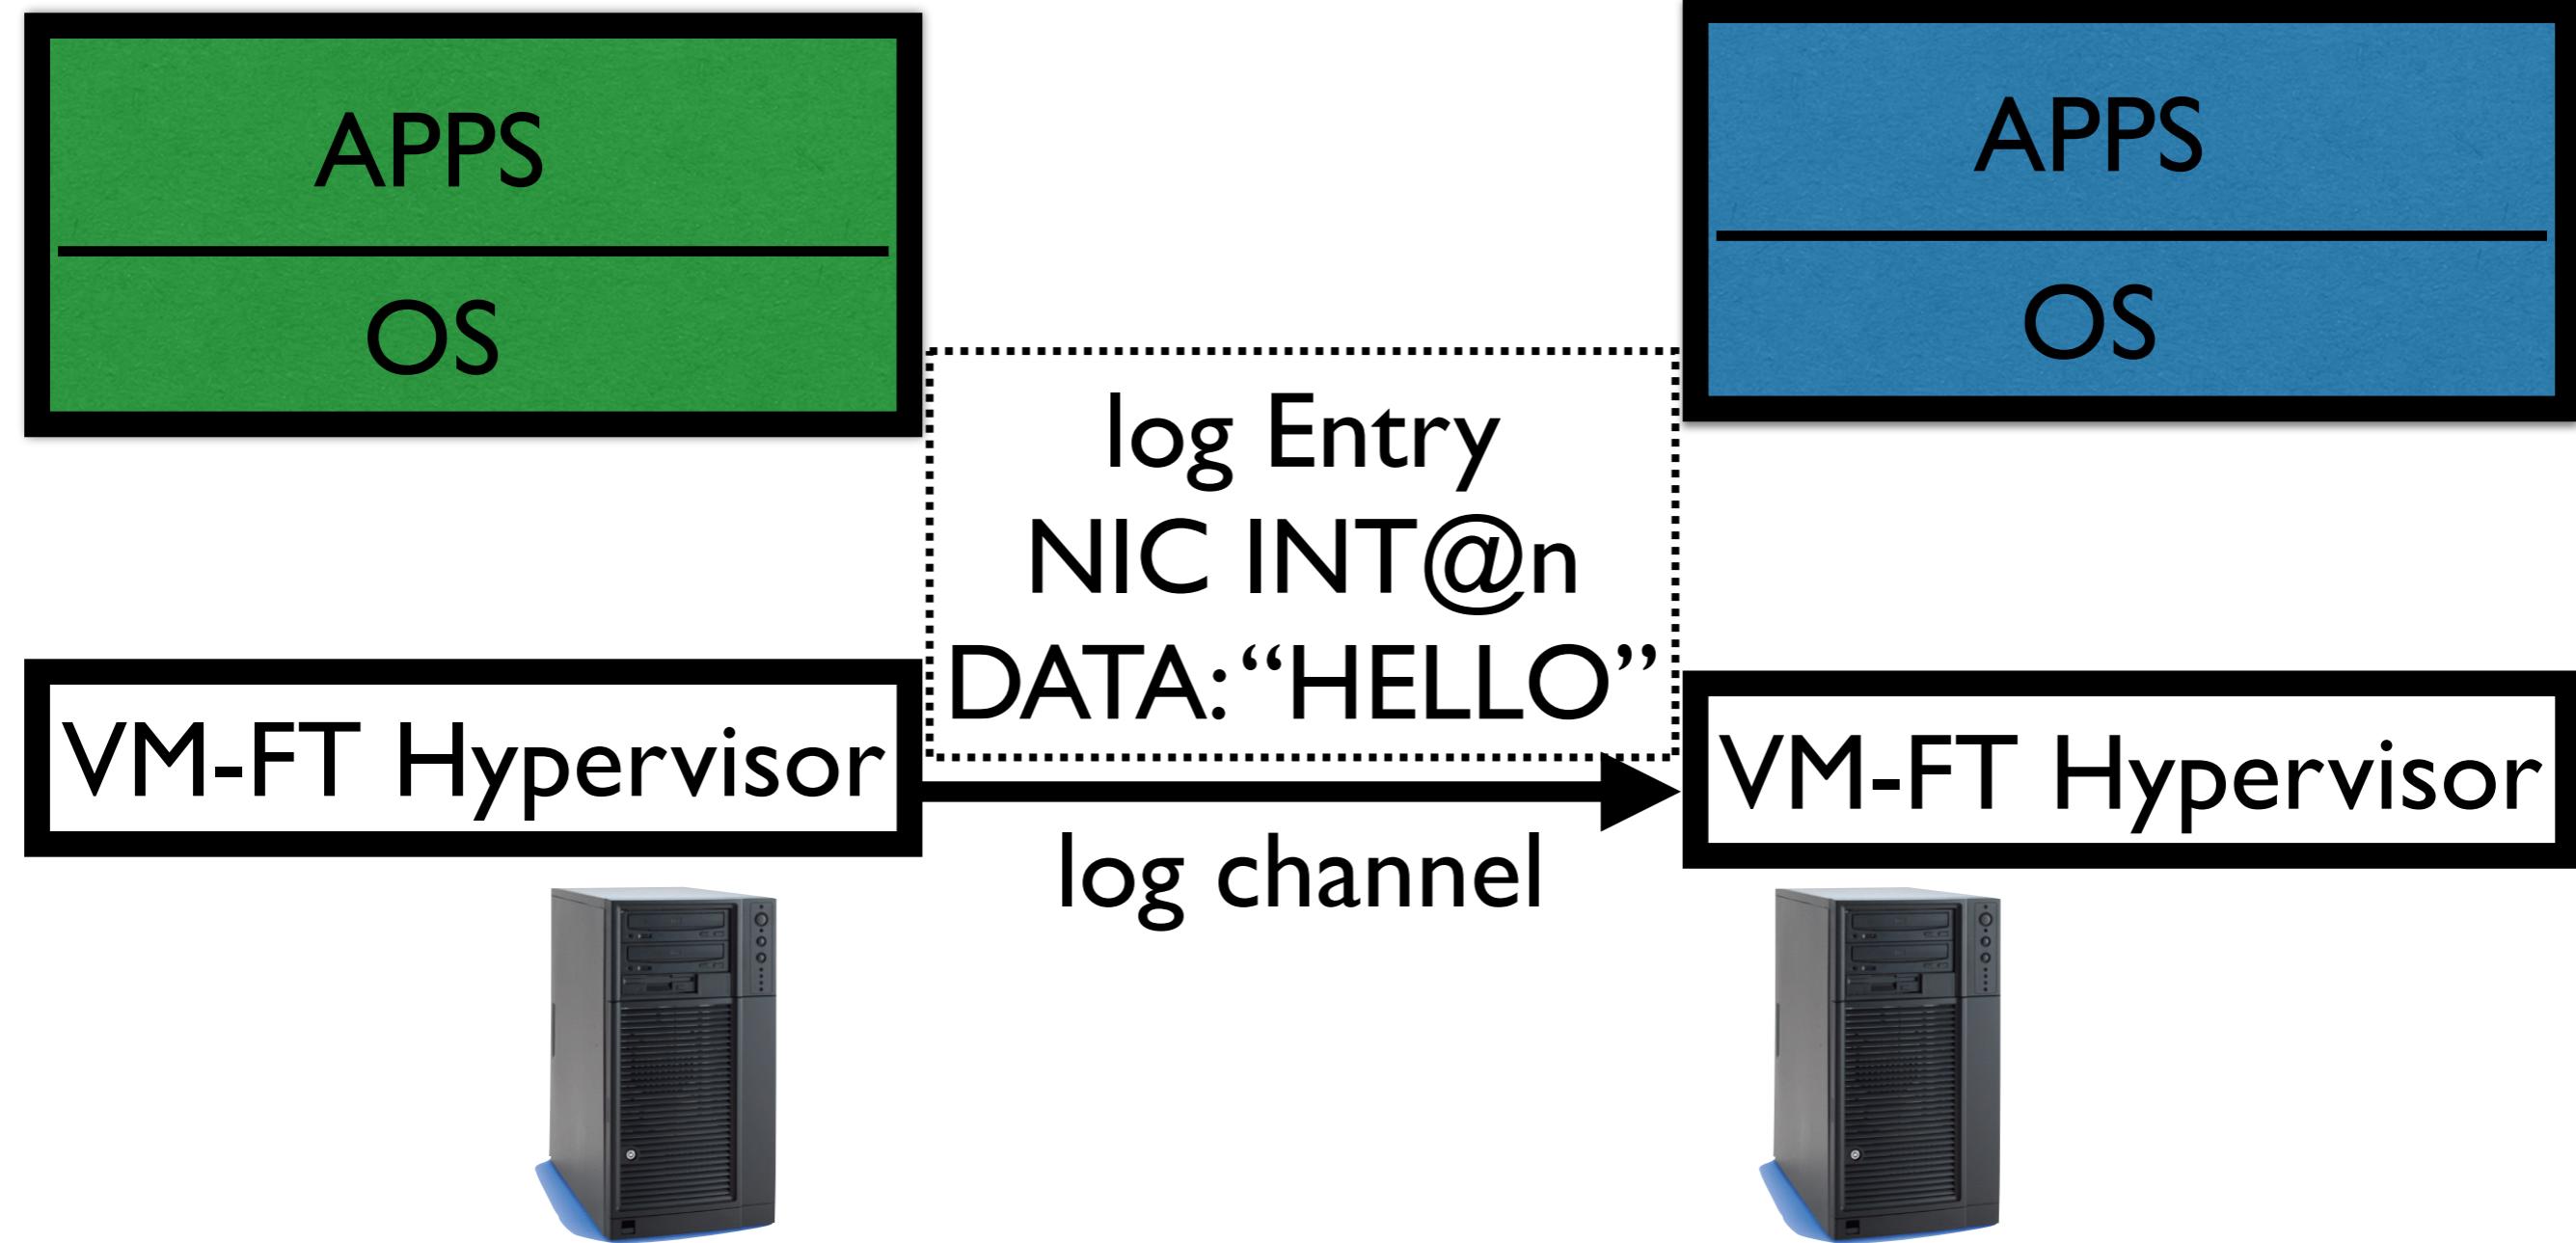


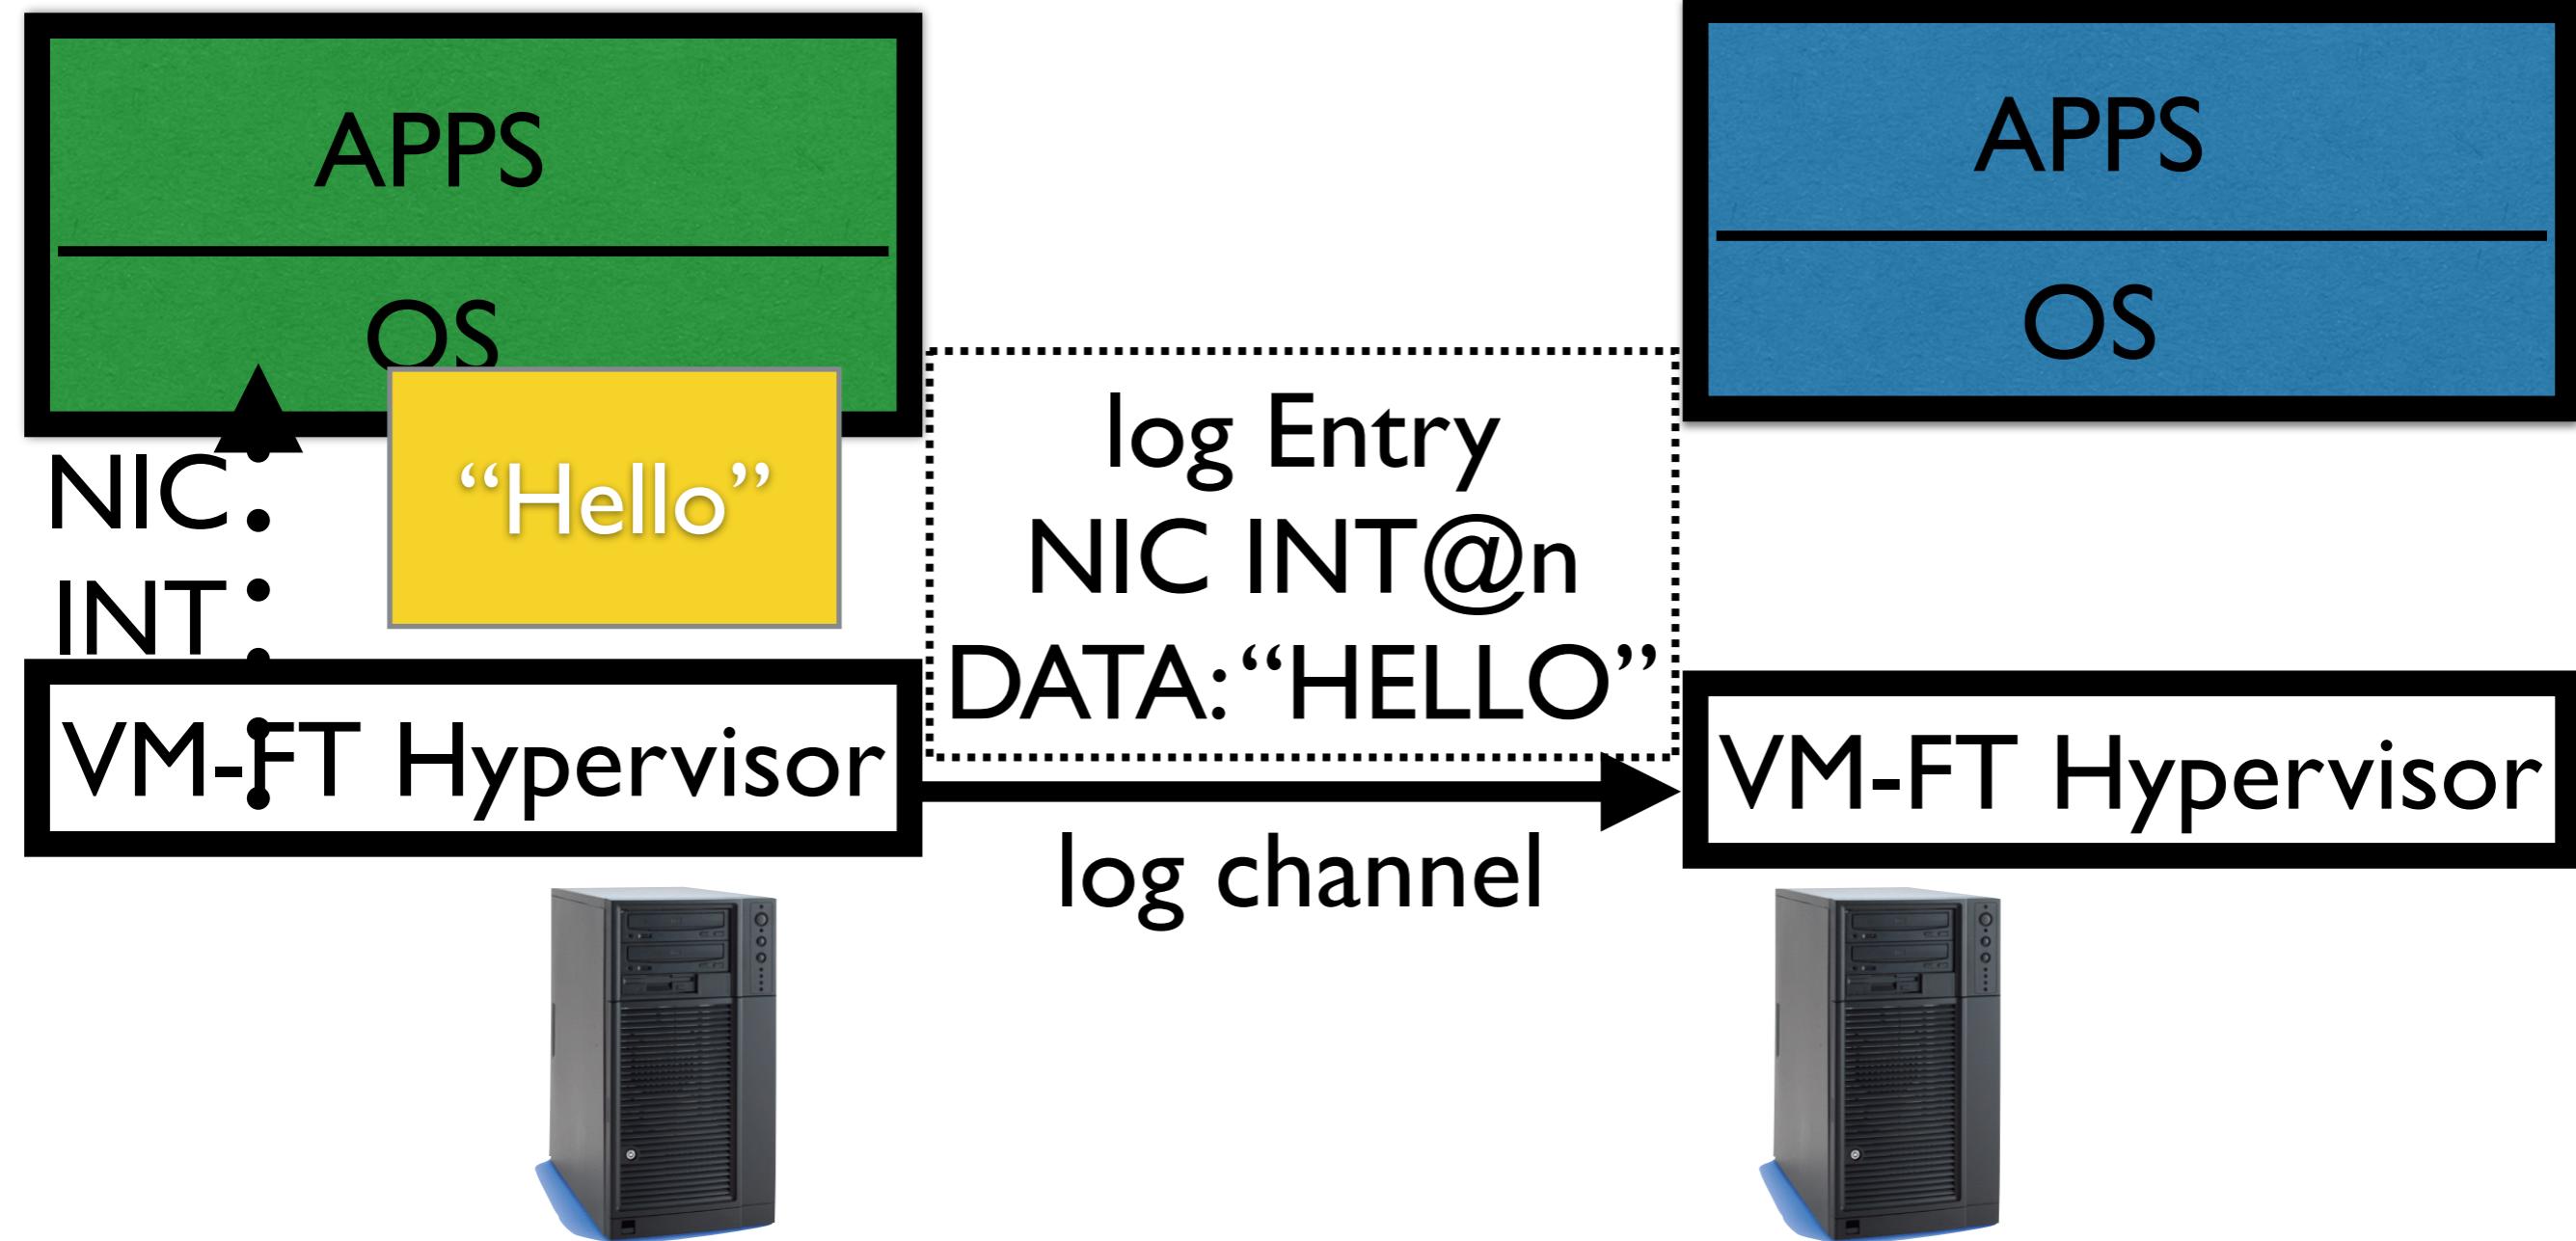
“Hello”

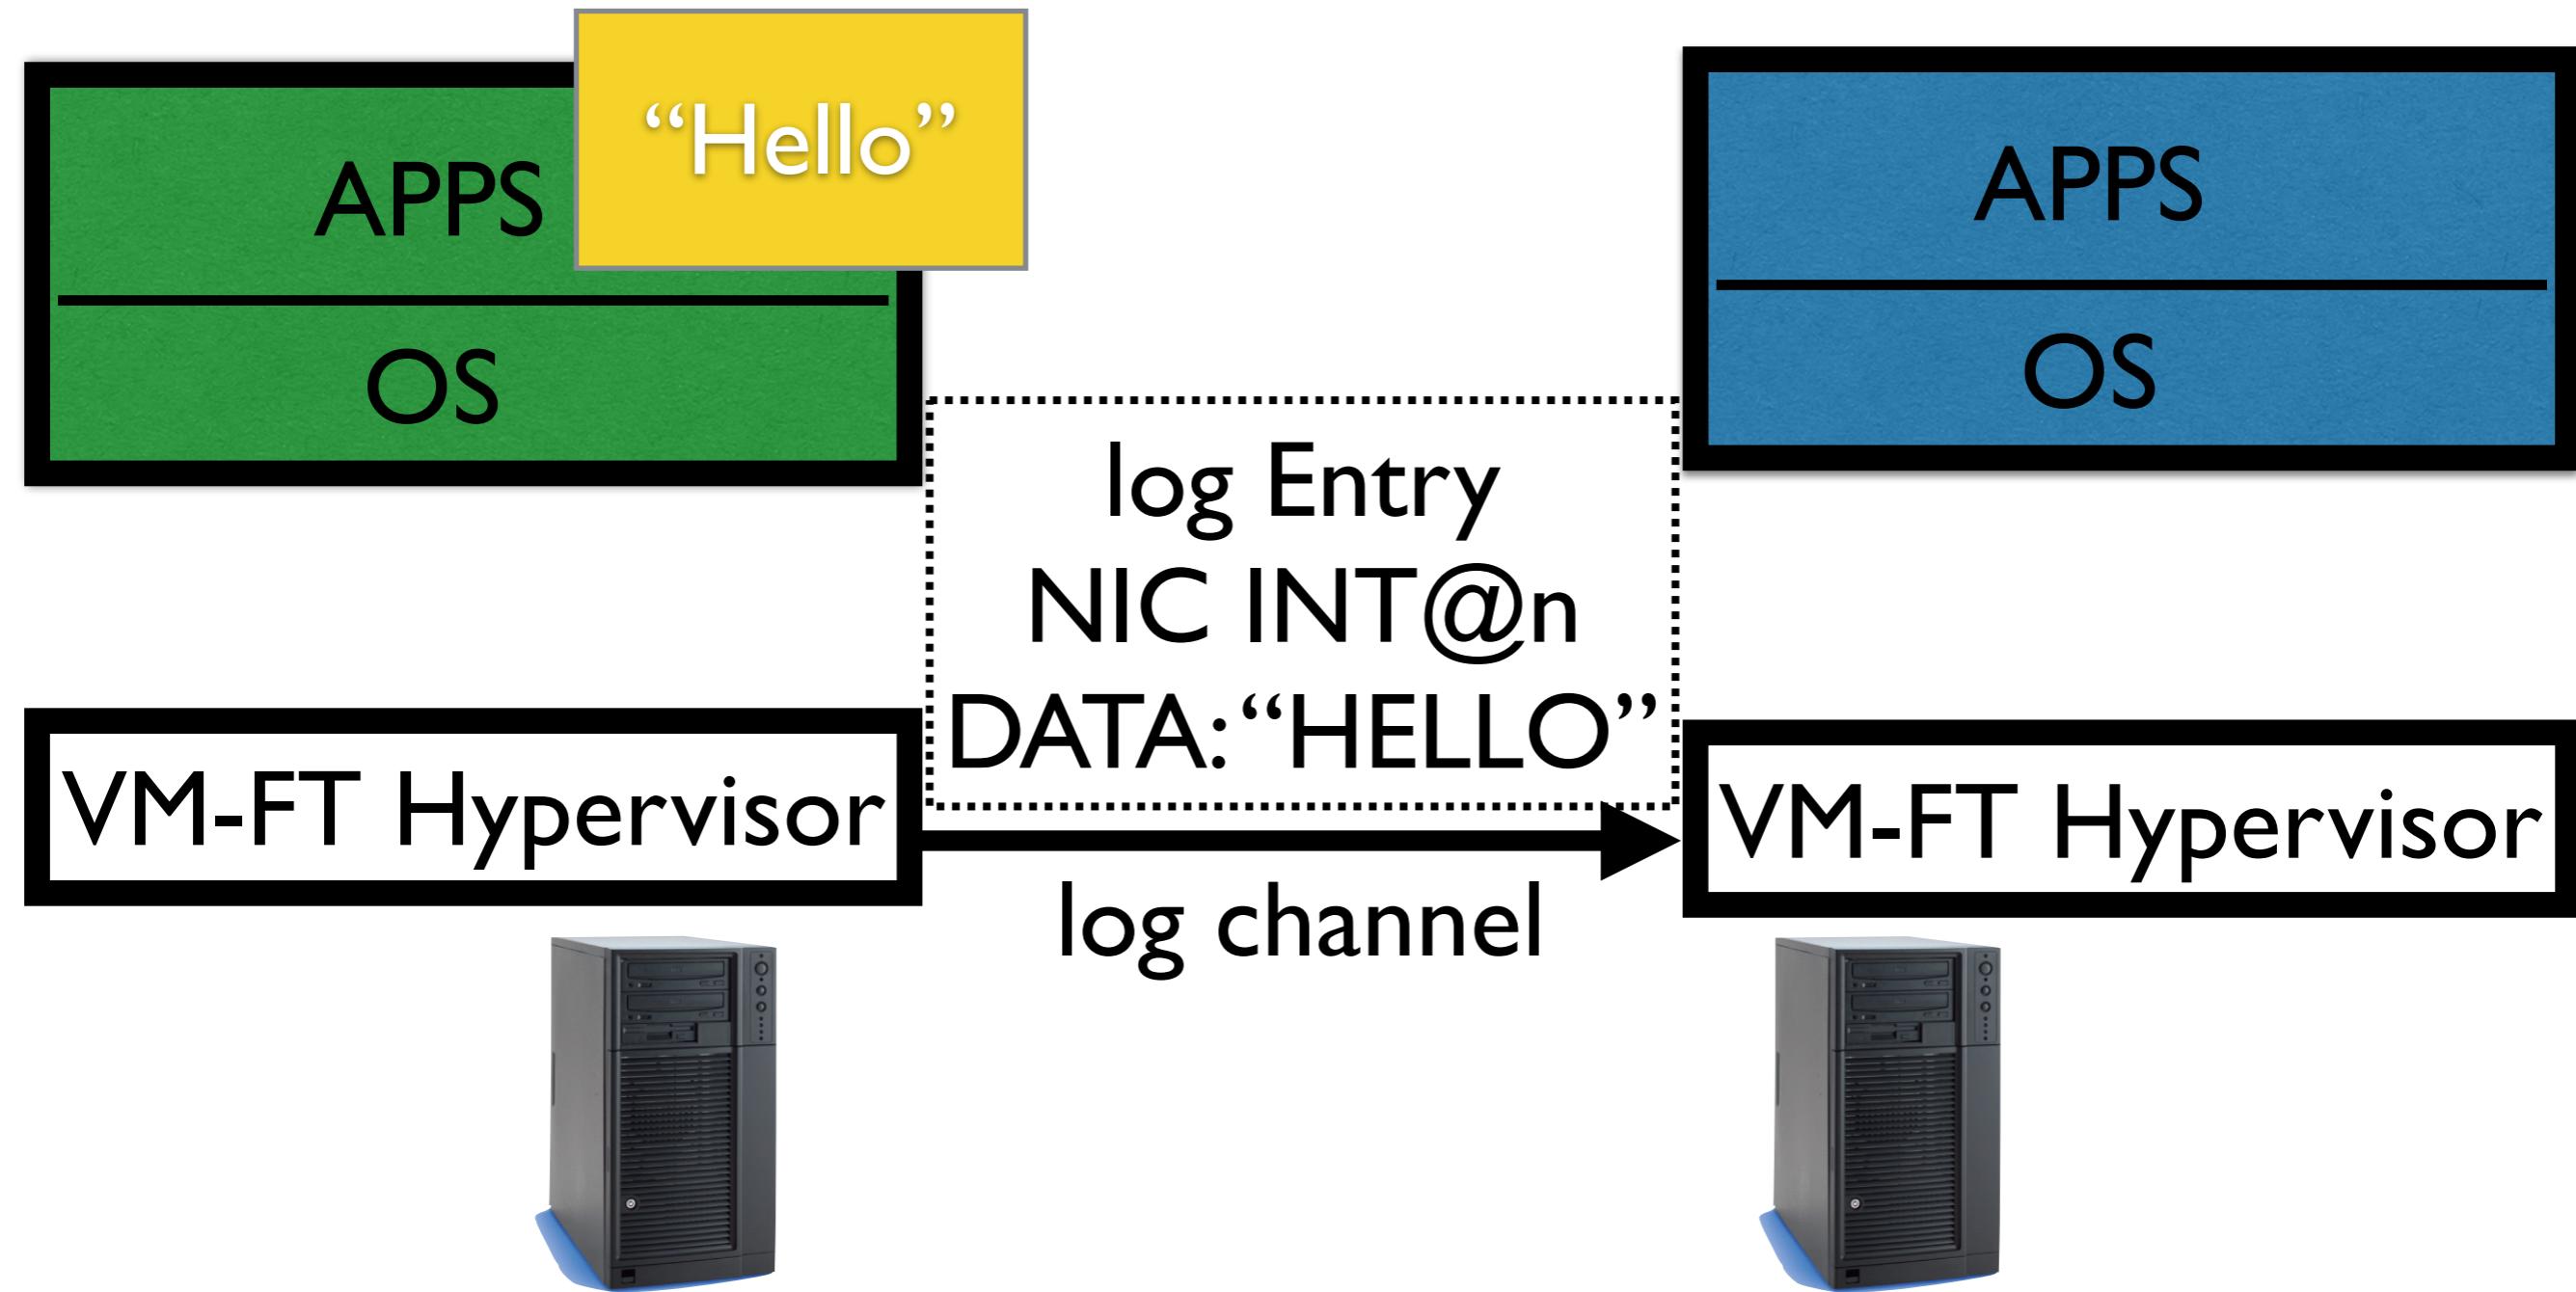
VM-FT Hypervisor

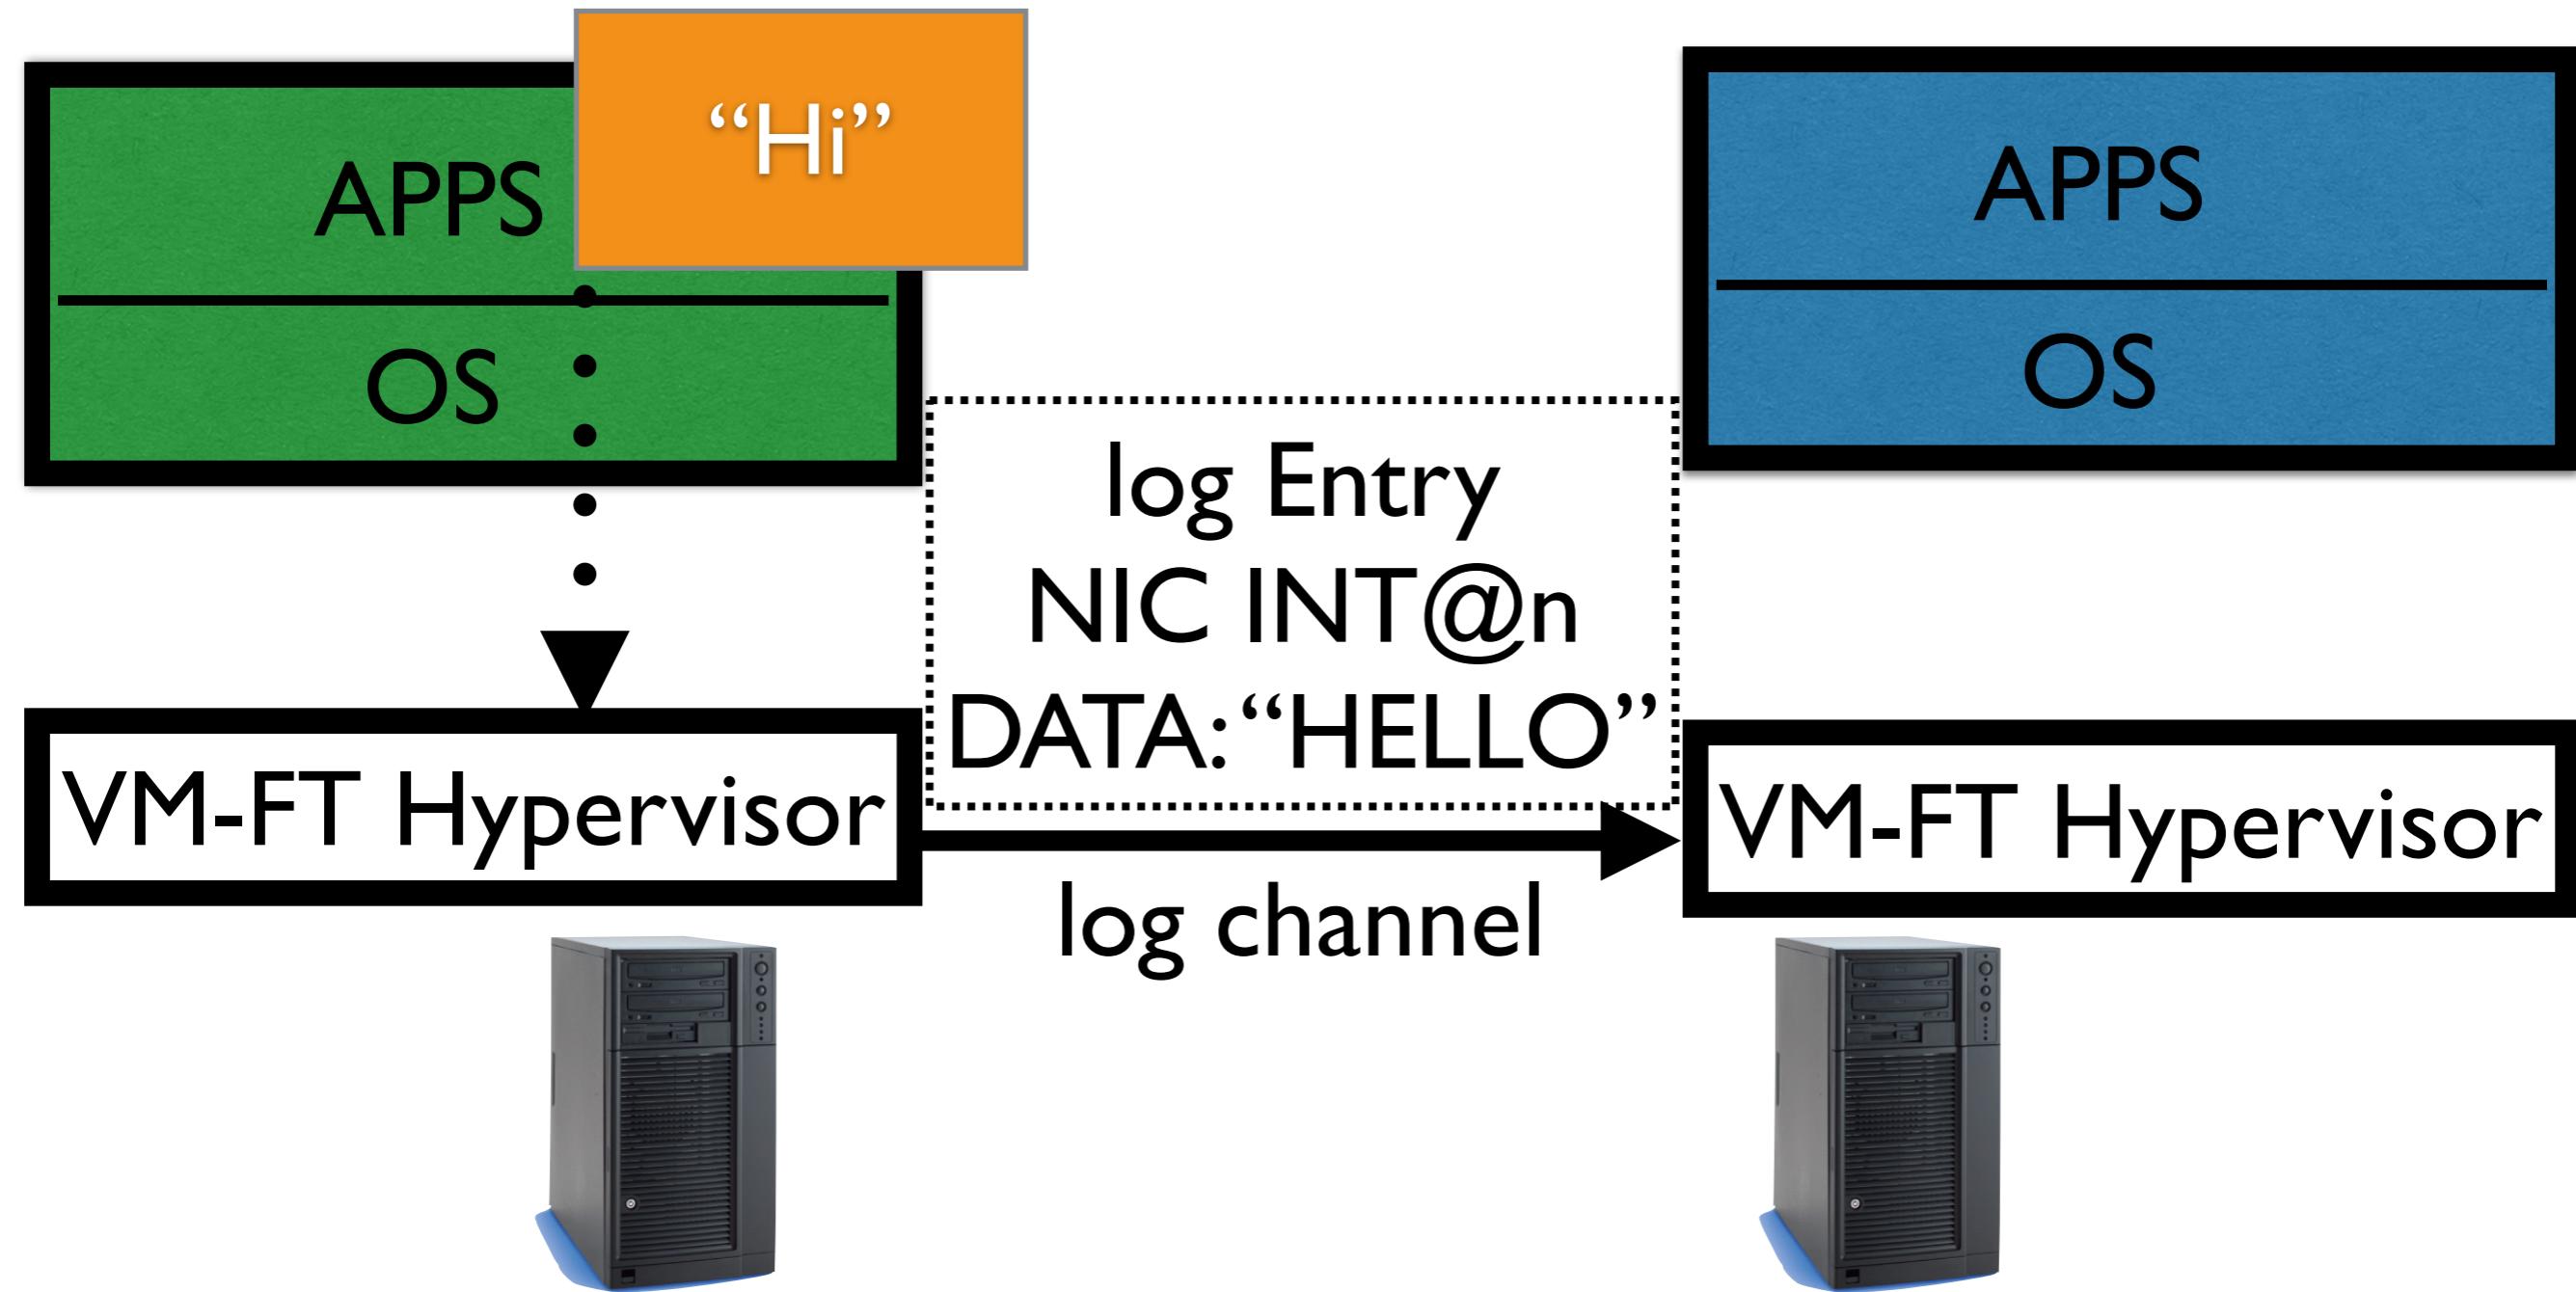


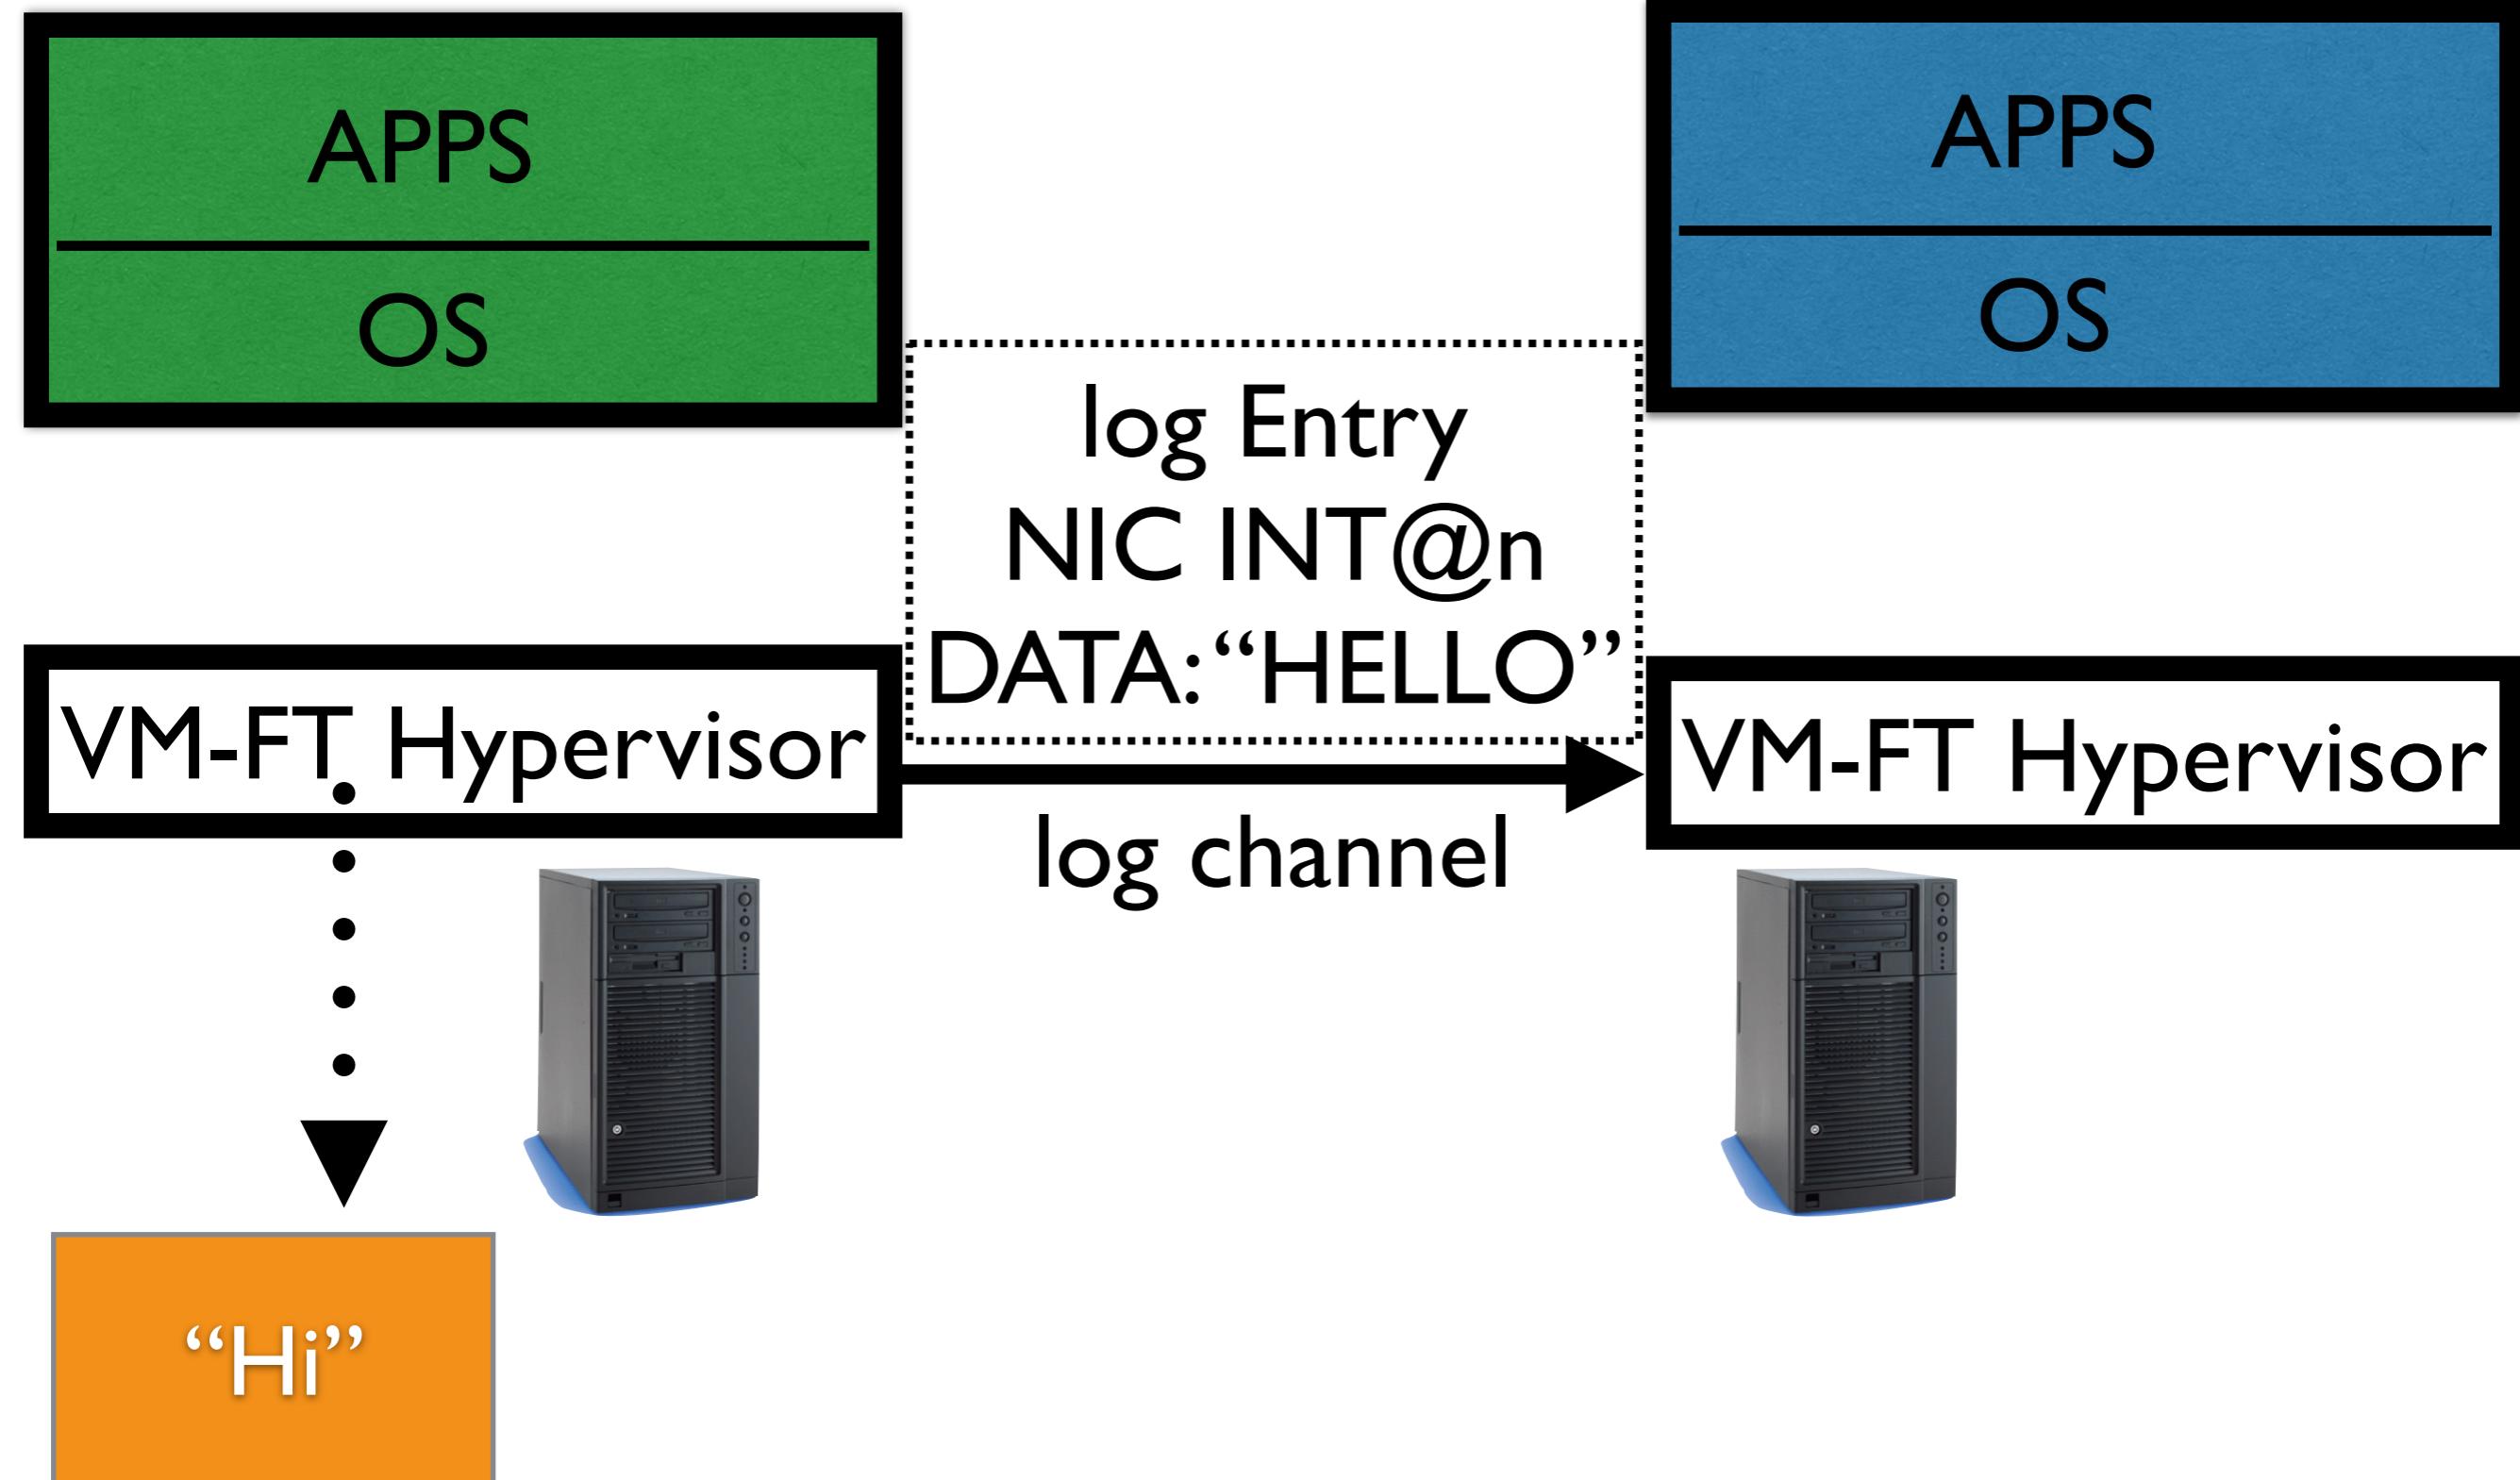
Network Packet

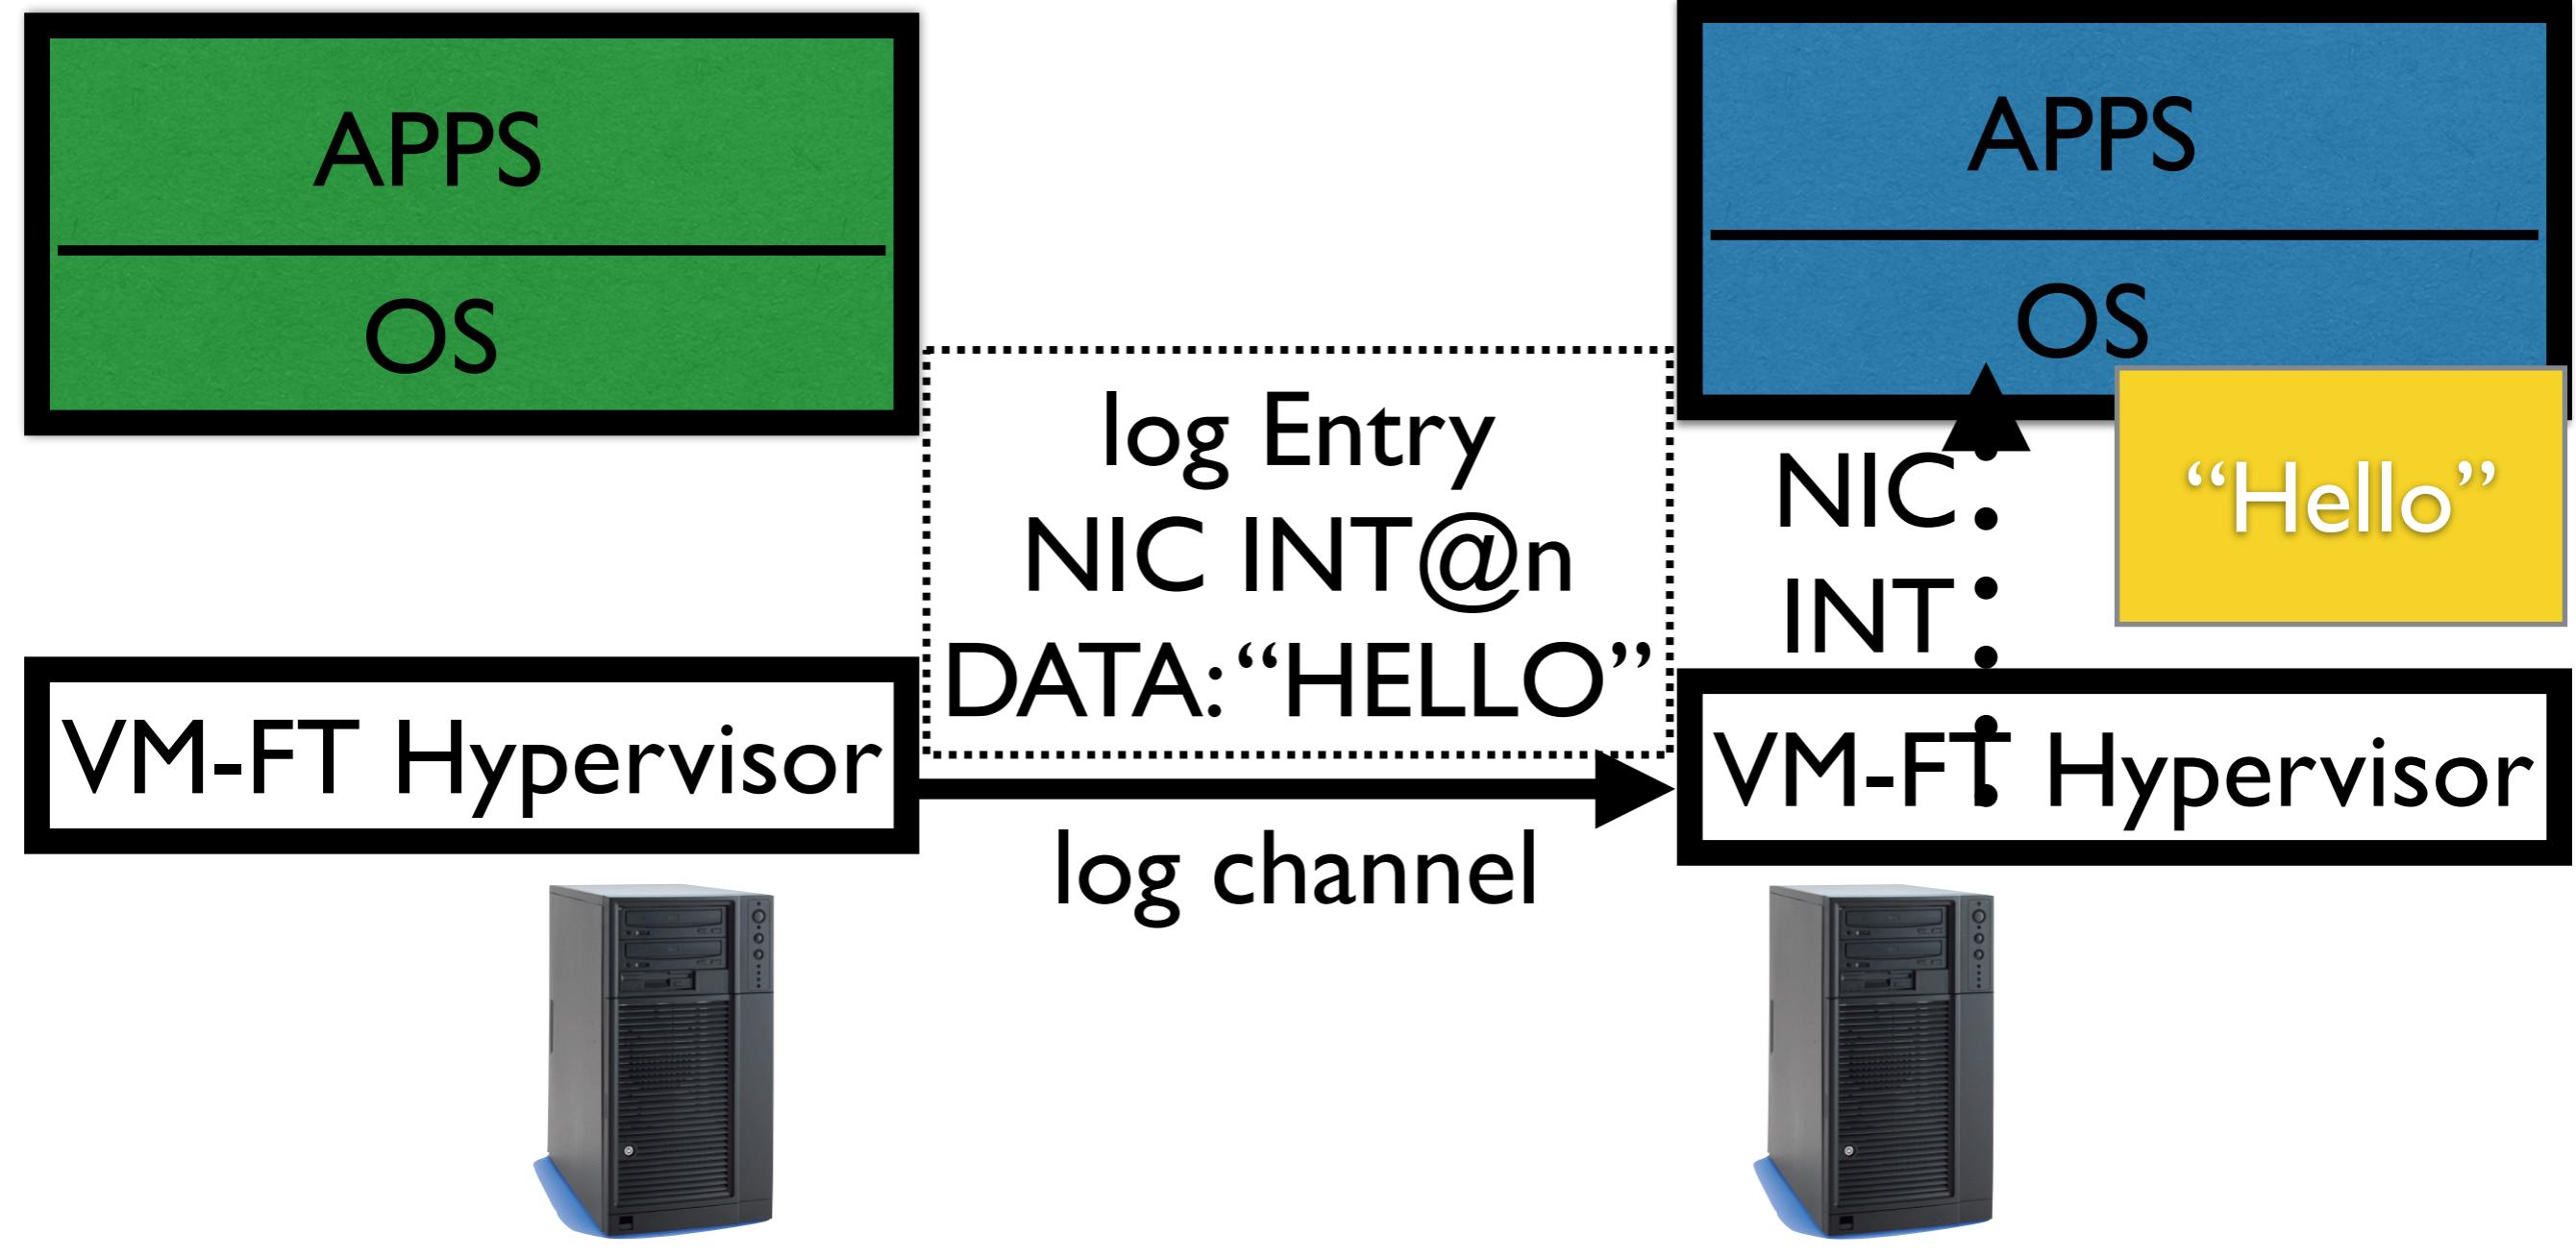




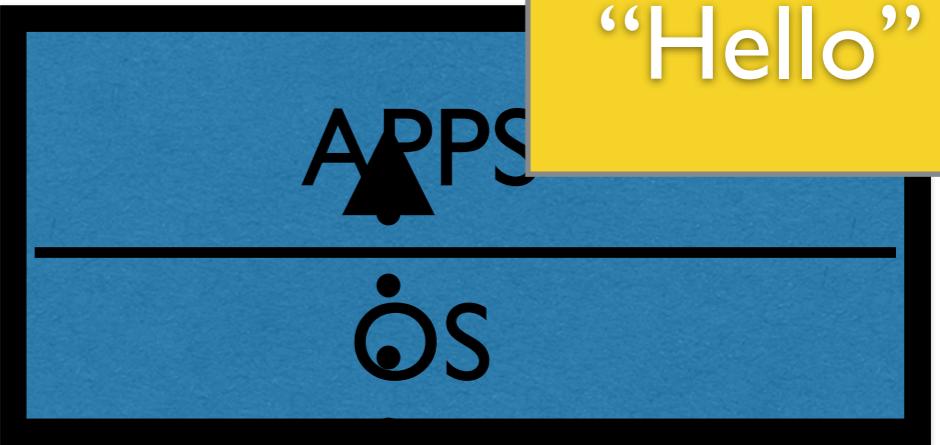
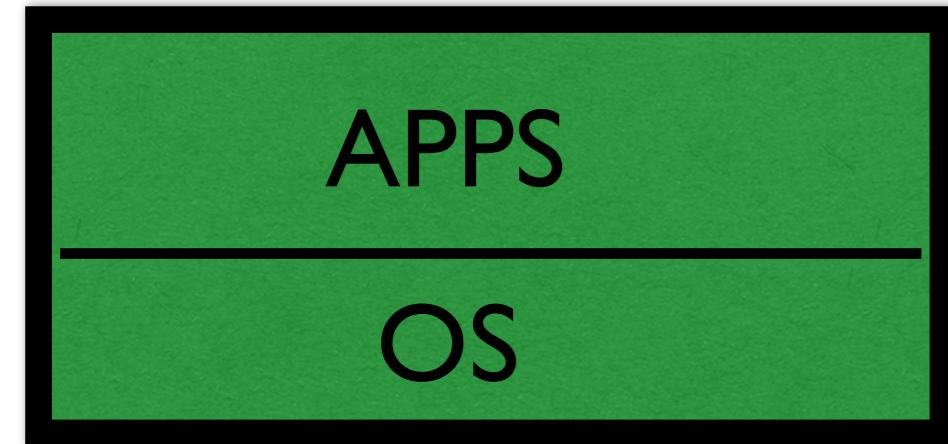




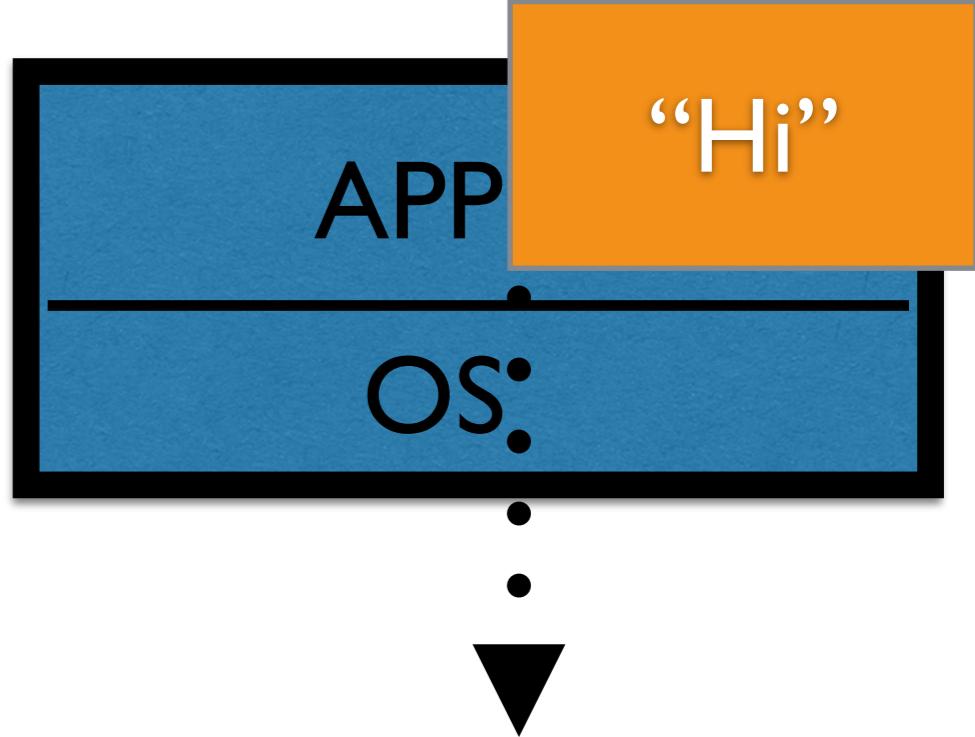
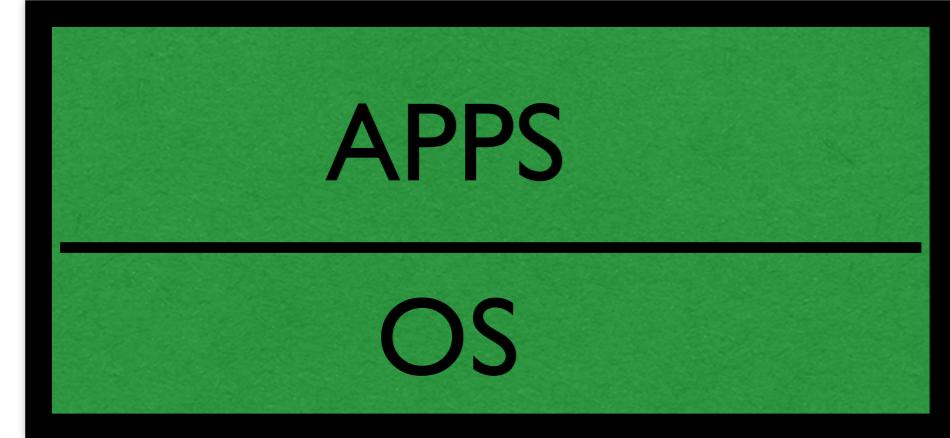




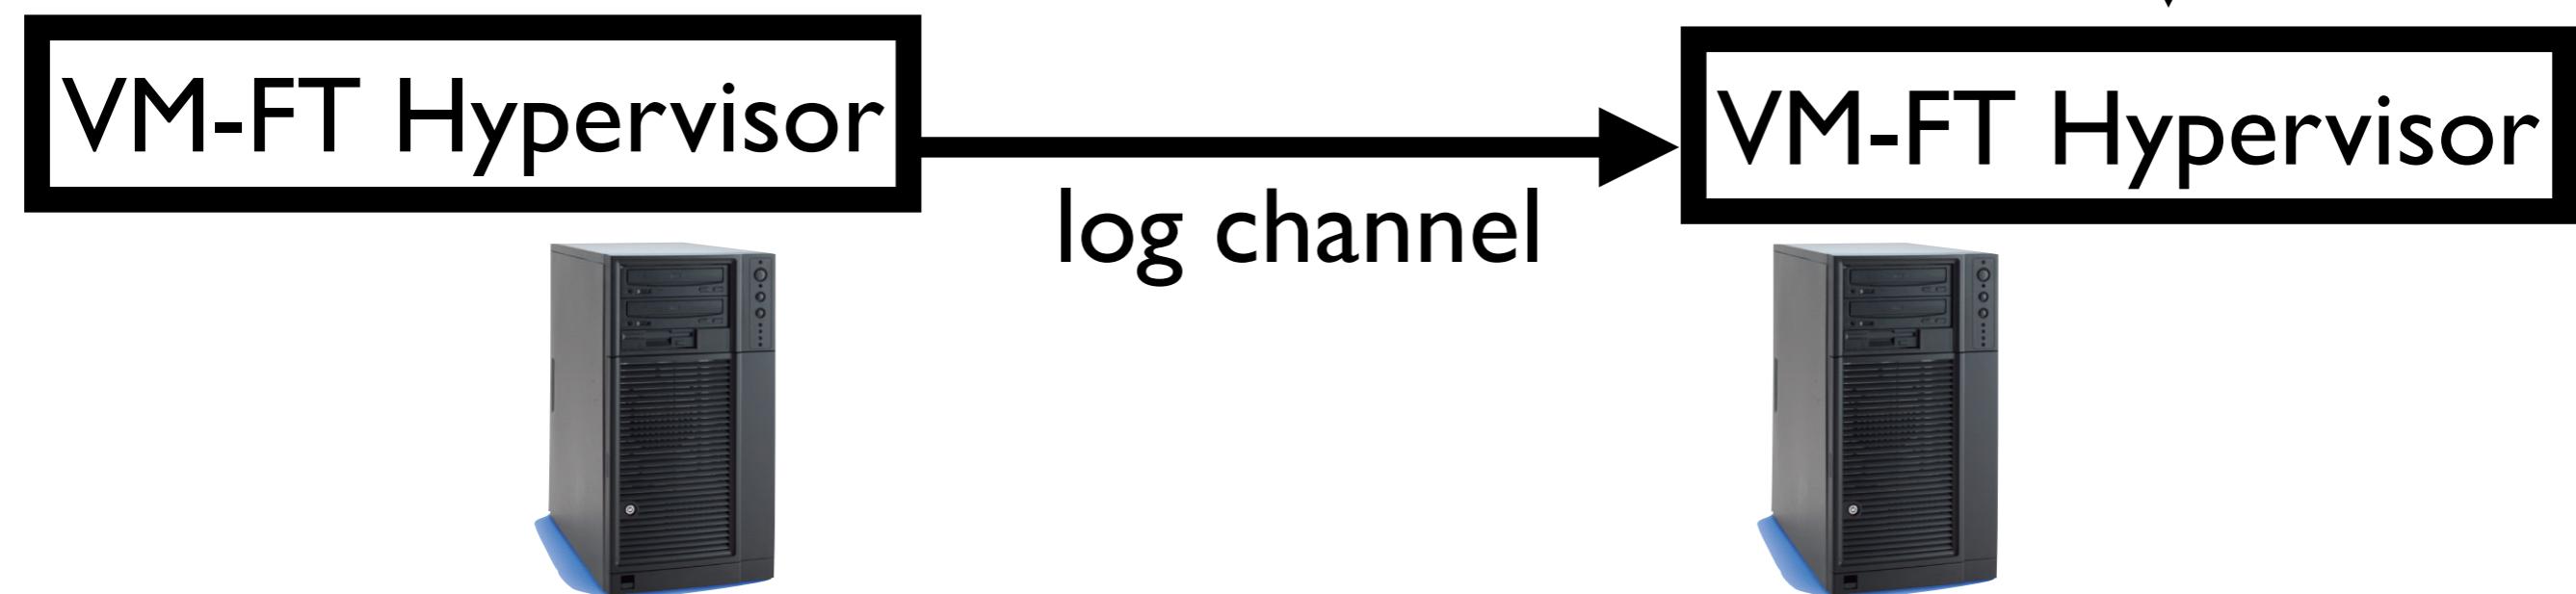
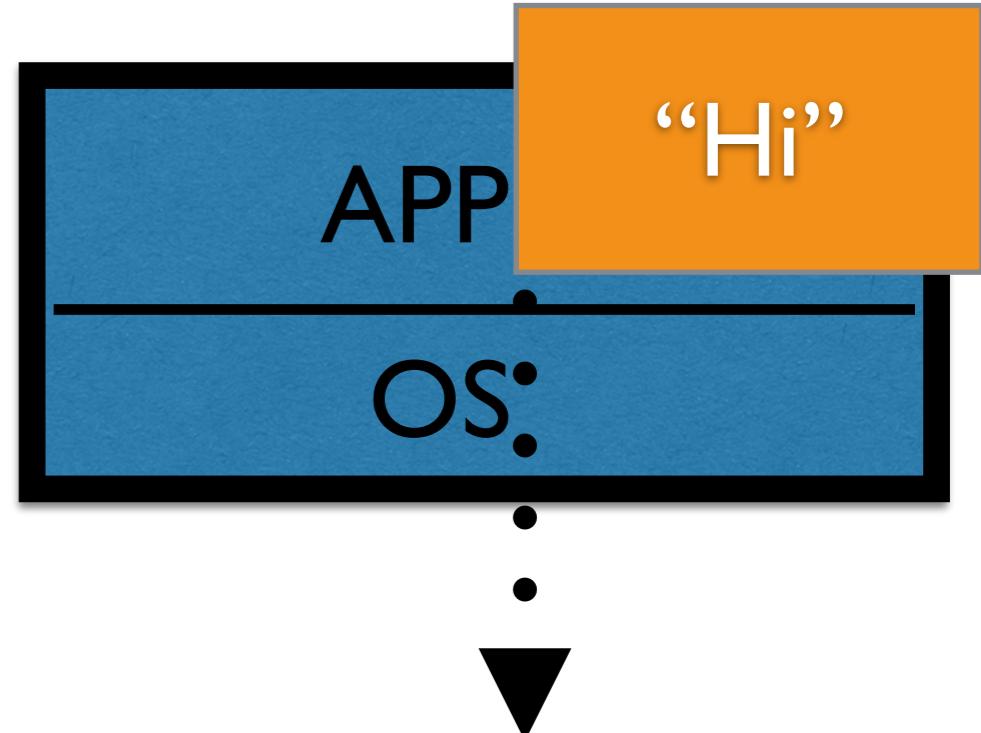
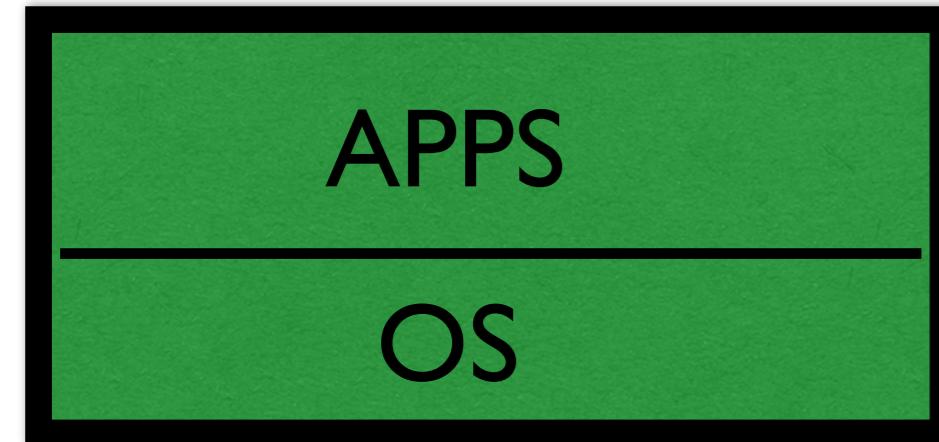
On Backup



On Backup

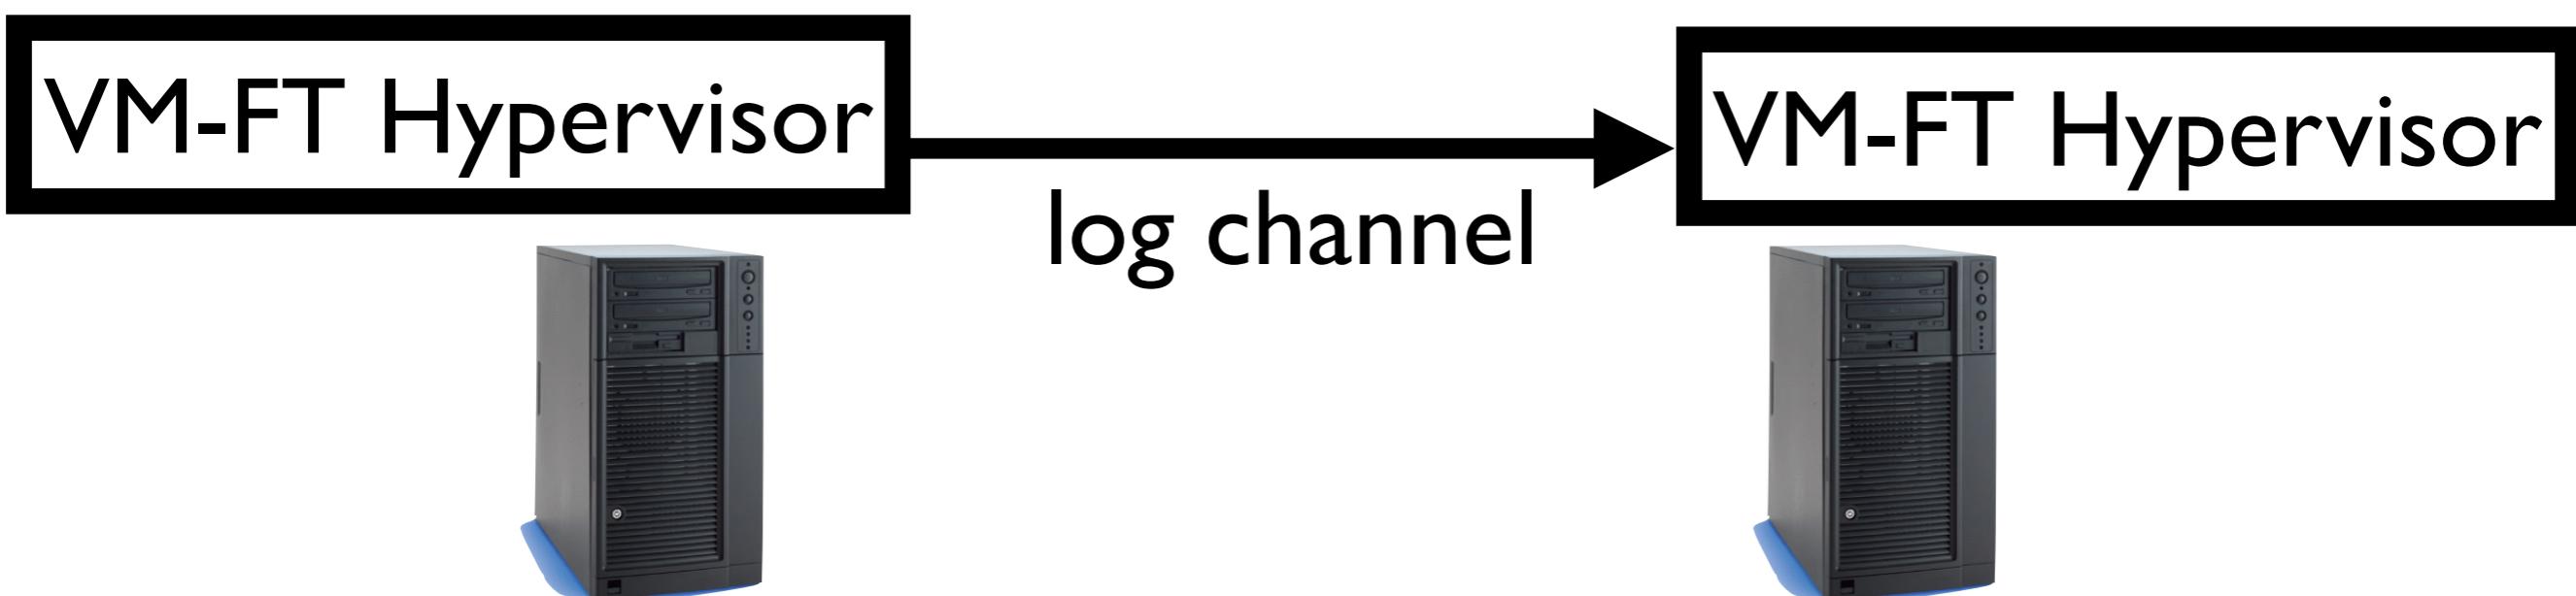
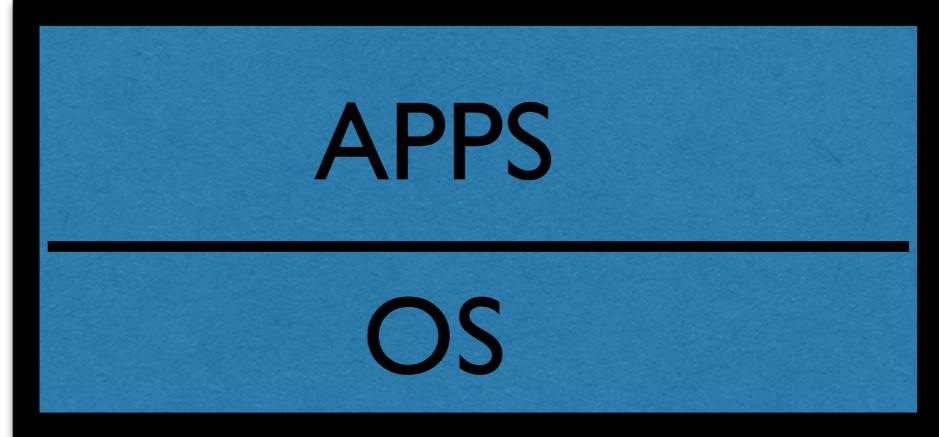
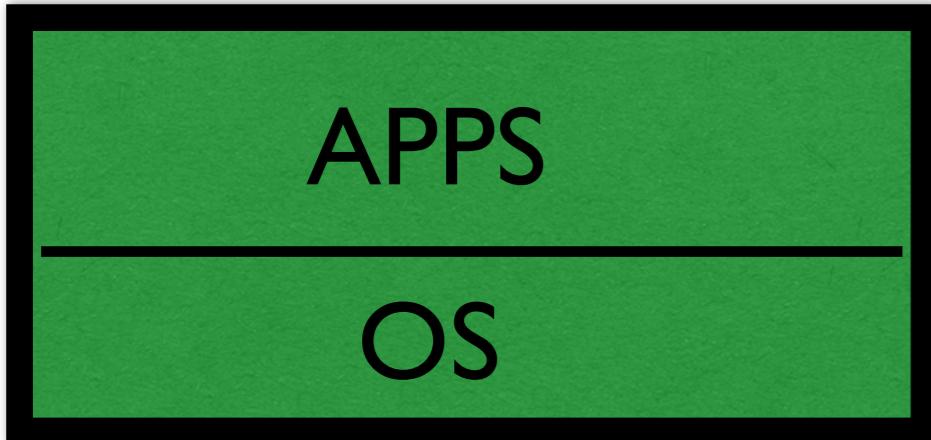


On Backup



does *not* put response on the wire

On Backup



does NOT put response on the wire

- ignores local clock gets from primary
- primary and backup same input end up in same state

Challenges

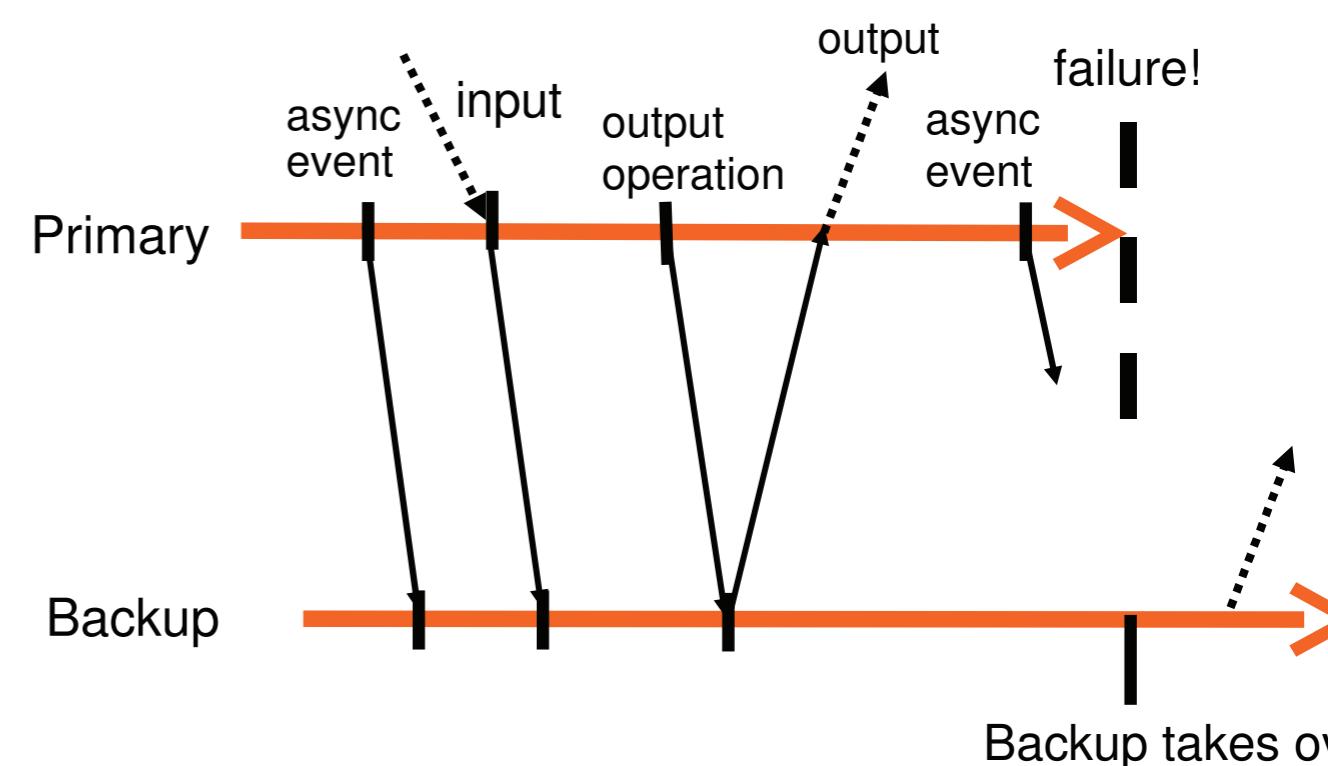
1. Making it look like a single reliable server from the outside
2. How to avoid two primaries?
3. How to make sure backup keeps up with primary?

Challenge I:

Making it look like a single reliable server from the outside

Primary delays output until backup acks

log each output op
only send after backup acked receiving output op



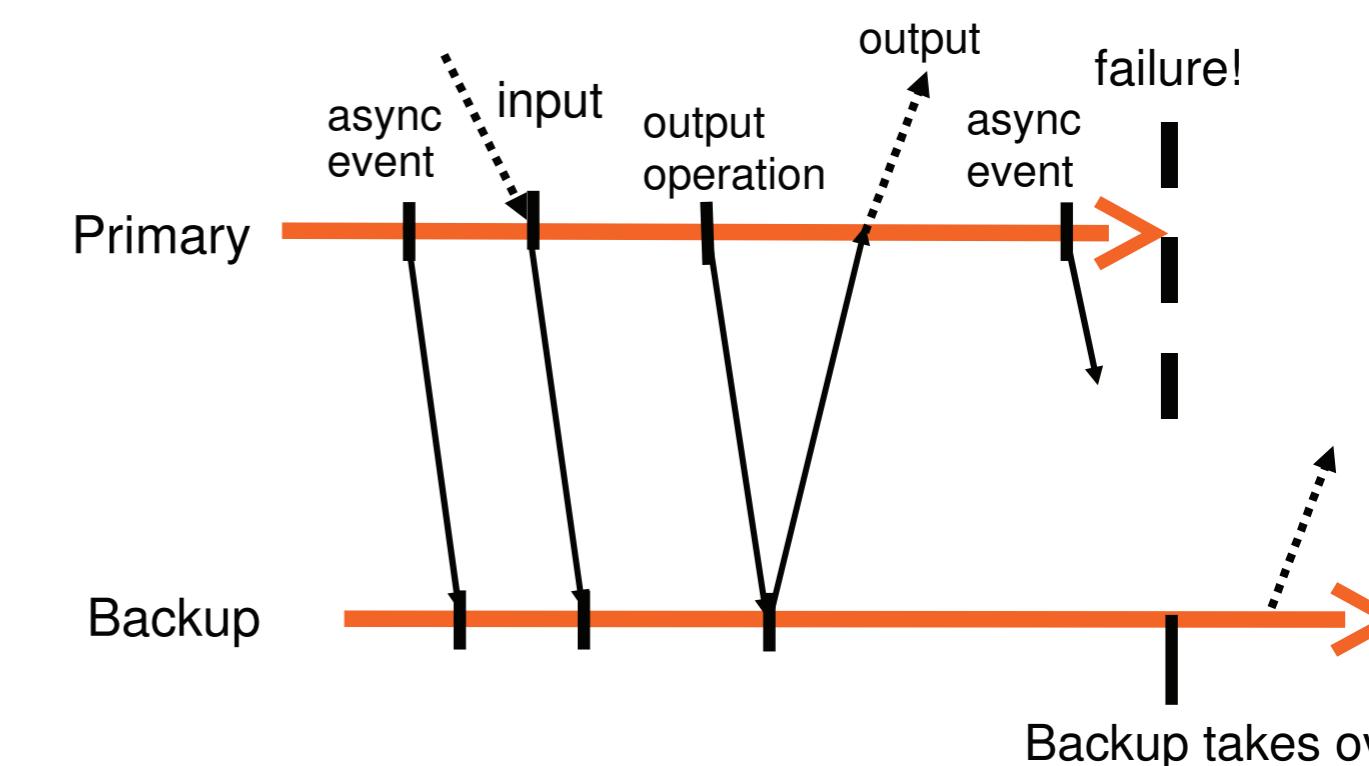
Challenge I:

Making it look like a single reliable server from the outside

Primary delays output until backup acks

log each output op

only send after backup asked receiving output op



perf opt: primary continues to execute after output buffers output till backup ack

What happens when
primary fails right after
sending the reply to the
client and the backup
takes over?

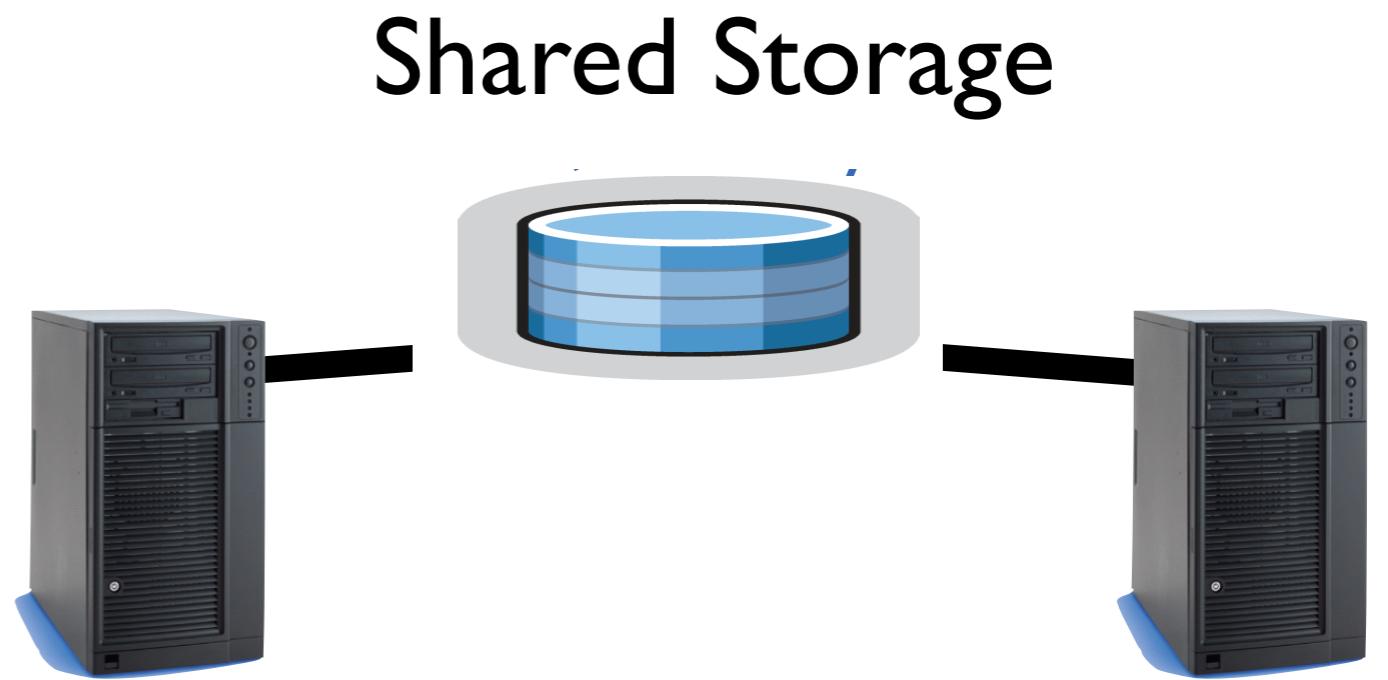
The backup doesn't know if client did receive a reply, so it will send the reply again when it becomes a primary

What happens when
primary fails after
receiving network input
but before
sending a
corresponding log
entry to backup?

Challenge 2: How to avoid two primaries? (Split Brain)

Split Brain

- Hard problem with only unreliable network!
- This is the point of things like Raft



Assume shared disk

Backup replays through last log entry

Backup atomically test-and-set “live” variable on disk

If set, primary is alive. Commit suicide

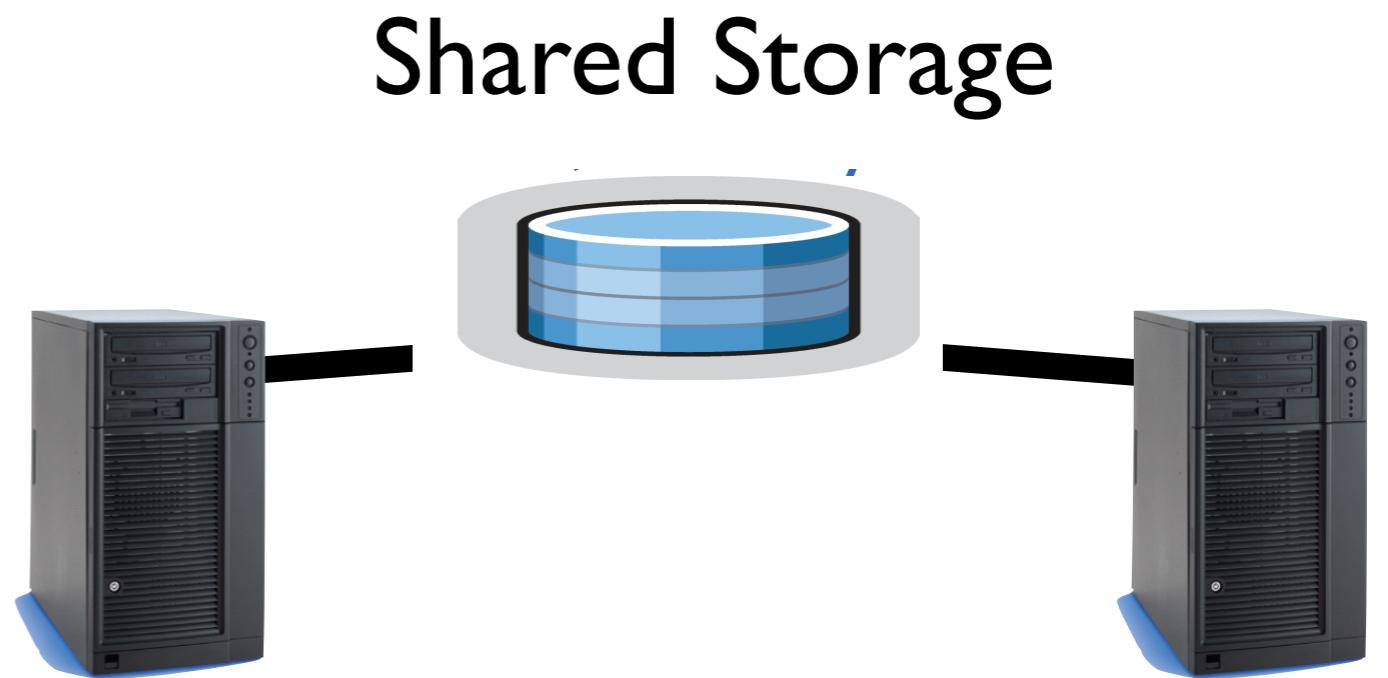
If not set, primary is dead. Become primary

If primary, create new backup from checkpoint

Using VMotion

Split Brain

- Hard problem with only unreliable network!
- This is the point of things like Raft



Assume shared disk

Backup replays through last log entry

Backup atomically test-and-set “live” variable on disk

If set, primary is alive. Commit suicide

If not set, primary is dead. Become primary

If primary, create new backup from checkpoint

Using VMotion

Challenge 3: Slow backup keeps up

- If Primary is much faster than backup
 - Log buffer will be full
 - Primary will block until there is free space in the buffer
 - VM FT can also slow down Primary proactively, until backup catches up

