

# Distributed Systems

**Spring Semester 2020**

Lecture 10: Optimistic Concurrency Control (Thor)

**John Liagouris**  
[liagos@bu.edu](mailto:liagos@bu.edu)

# Why this paper

- Focus on optimistic concurrency control (OCC)
  - Distributed OCC very interesting -- still active research
- Use of wall-clock time for transaction ordering to check serializability and provide external consistency
  - Rely on Network Time Protocol
  - Thor's use of time has been influential (e.g. Spanner)

<http://research.microsoft.com/en-us/um/people/mcastro/publications/liskov96safe.pdf>

<http://www.pmg.csail.mit.edu/papers/ecoop99.pdf>

[http://www.pmg.csail.mit.edu/thor/Thor.0\\_User\\_Guide.ps](http://www.pmg.csail.mit.edu/thor/Thor.0_User_Guide.ps)

# Why this paper

- Focus on optimistic concurrency control (OCC)
  - Distributed OCC very interesting -- still active research
- Use of wall-clock time for transaction ordering to check serializability and provide external consistency
  - Rely on Network Time Protocol
  - Thor's use of time has been influential (e.g. Spanner)

<http://research.microsoft.com/en-us/um/people/mcastro/publications/liskov96safe.pdf>

<http://www.pmg.csail.mit.edu/papers/ecoop99.pdf>

[http://www.pmg.csail.mit.edu/thor/Thor.0\\_User\\_Guide.ps](http://www.pmg.csail.mit.edu/thor/Thor.0_User_Guide.ps)

# Transaction (Databases)

- A unit of work (read-write operations) that is:
  - **A**tomic: Either the whole transaction succeeds or fails
  - **C**onsistent: Does not violate database restrictions after completion
  - **I**solated: Results of incomplete transactions are not visible to other transactions
  - **D**urable: After the transaction completes (commits) it cannot be lost

# Serializability

- Not to be confused with linearizability!
- A correctness condition
- Concurrent execution of transactions must be equivalent to **a** serial execution
- If no equivalent serial execution exists, a transaction have read results of other incomplete transactions
- If serial order respects the real time the transactions were committed by the client, then:
  - **External consistency**

# OCC vs Locking

## (Optimistic vs Pessimistic)

```
l.lock();  
val=read(x);  
val++;  
write(x);  
l.unlock();
```

```
tx.begin()  
val=read(x);  
val++;  
write(x);  
tx.end();
```

What is the difference between these two?


# OCC vs Locking

## (Optimistic vs Pessimistic)

```
l.lock();  
val=read(x);  
val++;  
write(x);  
l.unlock();
```

```
tx.begin()  
val=read(x);  
val++;  
write(x);  
tx.end();
```

**Assume  $x=0$**





# OCC vs Locking

## (Optimistic vs Pessimistic)

```
l.lock();  
val=read(x);  
val++;  
write(x);  
l.unlock();
```

```
tx.begin()  
val=read(x);  
val++;  
write(x);  
tx.end();
```

**Assume  $x=0$**





# OCC vs Locking



## (Optimistic vs Pessimistic)

```
l.lock();  
val=read(x);  
val++;  
write(x);  
l.unlock();
```

```
tx.begin()  
val=read(x);  
val++;
```

OCC : Fast but WRONG

Assume s



# OCC vs Locking



## (Optimistic vs Pessimistic)

```
l.lock();  
val=read(x);  
val++;  
write(x);  
l.unlock();
```

```
tx.begin()  
val=read(x);  
val++;
```

OCC :Must detect, abort one and retry — longer and more work


Assume s



# OCC vs Locking (Optimistic vs Pessimistic)

OCC better if few conflicts otherwise  
overheads could be larger than lock waits in  
case of conflicts

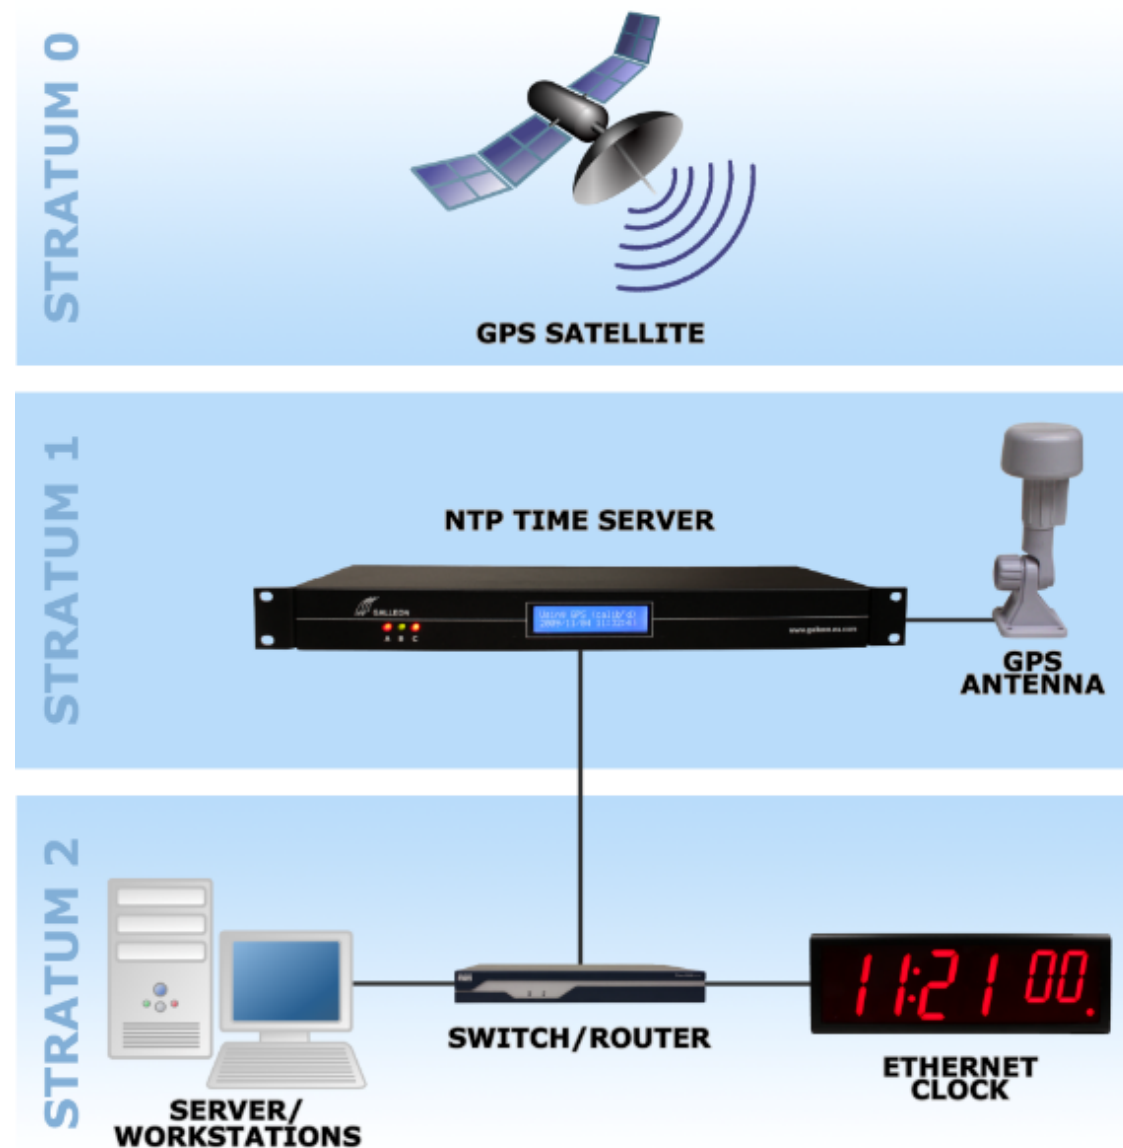
abort  
longer  
rk



# Network Time Protocol

- Clocks in different machines are not perfectly synchronized (skew, drift)
- Causes of clock skew: temperature variations, material imperfection, etc.
- Atomic clocks: Very accurate. Will not gain or lose a second in 300 million years!
- NTP uses an atomic clock as reference to allow machines in a network synchronise their clocks

# Network Time Protocol



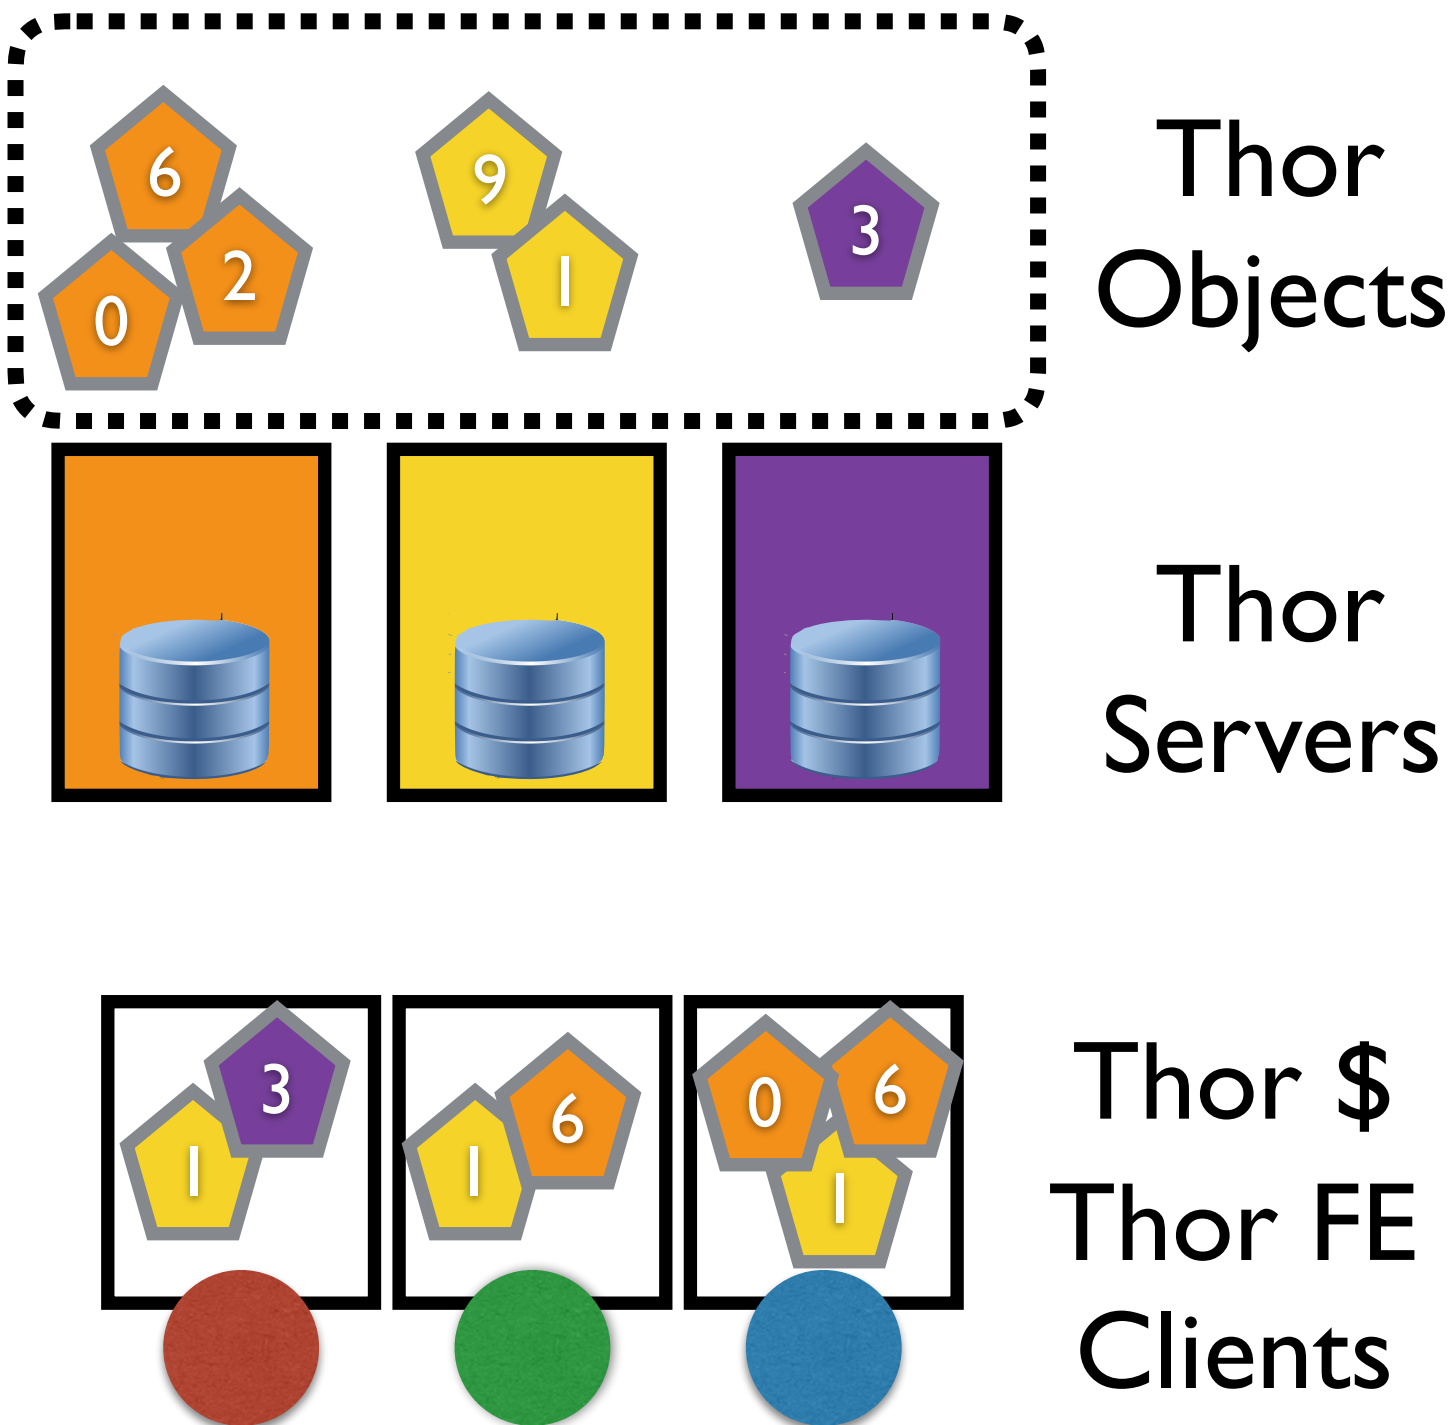
<https://www.galsys.co.uk/news/what-is-ntp-a-beginners-guide-to-network-time-protocol/#What-is-NTP?>

# Network Time Protocol

- With NTP clock differences are in the order of 10s of *milliseconds*
- Geng et al. presented an approach to reduce clock skew to 10-100s *nanoseconds*.
- Might have implications to the design and implementation of distributed systems algorithms


<https://www.usenix.org/system/files/conference/nsdi18/nsdi18-geng.pdf>

# Thor: Distributed DB



- Universe of persistent Thor Objects
- Data/Object state partitioned across servers
- methods/operations run on clients (not RPC)
- clients read/write object state to from server DB's
- client cache for fast access

# Client Caching and Trans



- Writes have to invalidate cached copies
- How to cope with reads of stale cached data?
- How to cope with read-modify-write races?
- Clients could lock before using each record
- but that's slow -- probably need to contact server
- wrecks the whole point of fast local caching in clients



# Thor uses optimistic concurrency control (OCC)

- R/W local copies directly
  - Pretend sequential — no-concurrency — until commit
- On Commit:
  - send R/W info to server for “validation”
  - validation decides if it is really “OK” — serializable
    - YES: Send invalidates to cached copies
    - NO: Abort — discard changes

# Thor uses optimistic concurrency control (OCC)

- Optimistic: Designed around the hope that no conflicts is common case
  - if true then Fast :-)
  - if false (validation can detect) but Slow :-)

Core of the paper is how to do Validation  
and what we will focus on


# What should validation do?

- Look at what was Read and Written by the executing transactions and
  - decide if there is a serial order that would have gotten the same result as concurrent execution
- Many OCC validation algorithms

We will work our way to Thor's

# Scheme I : Centralized

- Single Validation Server
- Clients send Read and Writes (including Values)
- Server must decide:
  - Would the results be “serializable” if we let the transactions commit?
  - Shuffle the transactions to see if we can find an order that ensures value READ are from most recent WRITES



# Scheme I : Centralized

## Example I

**x=0    y=0    z=0**


**T1:   Rx0   Wx1**

**T2:   Rz0   Wz9**

**T3:   Ry1   Rx1**

**T4:   Rx0   Wy1**

Validation needs to decide if this  
execution (reads, writes) is  
equivalent to some serial order



# Scheme I : Centralized

## Example I

**x=0    y=0    z=0**


**T1:   Rx0   Wx1**

**T2:   Rz0   Wz9**

**T3:   Ry1   Rx1**

**T4:   Rx0   Wy1**

Validation needs to decide if this execution (reads, writes) is equivalent to some serial order



**YES: T4, T1, T3, T2 – Say yes to all Transactions (T2 can go anywhere)**

# Scheme I : Centralized

## Example I

**x=0    y=0    z=0**

**T1:   Rx0   Wx1**


**T2:   Rz0   Wz1**

**T3:   Ry1   Rx1**

**T4:   Rx0   Wy1**

More permissive than  
Thor

Validation needs to decide if this execution (reads, writes) is equivalent to some serial order



**YES: T4 , T1 , T3 , T2 – Say yes to all Transactions (T2 can go anywhere)**