

Distributed Systems

Spring Semester 2020

Lecture 12: FaRM

John Liagouris
liagos@bu.edu

Why this paper

- 2015 — Data Center Scale Distributed Systems in light of Modern Hardware
- An outstanding achievement
 - 90 million Replicated, Persistent, Transactions per second! (TATP)
 - 1 million transactions/second per machine
 - ~10,000 tweets per second, ~3 million emails/s

Is this hard? Think of what we have seen to date.

How do they do it?

- Questions a very basic assumptions that the other papers have been predicated on
- The hardware! — Software carefully designed
 - EVERYTHING IN RAM
 - NO DISKS FOR PERSISTENCE
 - FASTER NETWORK
 - ELIMINATE PROCESSING OVERHEADS ON COMMUNICATON PATHS

The HW Trends

128-512GB RAM/machine

<\$12/GB —


2000 machines 1PetaByte

Lithium ION Batteries
and SSD's — Cheap NVM

Datacenter networking which integrates
supercomputer interconnect capabilities
2x56Gbps NICs full bisectional bandwidth switch
— cpu-less communication — RDMA

All aspects required

- Massive Ram required to remove disk bottleneck
- NVM required to have durability again to remove disk bottleneck
- Faster technology required to remove network bottleneck
- Advanced features to avoid CPU bottleneck



Must alleviate or
remove from hot
paths

Big Memories

FaRM writes go to RAM, not disk
-- eliminates a huge bottleneck

Can write RAM in 200 ns, but
takes 10 ms to write hard drive,
100 us for SSD

ns = nanosecond 10^{-9} , ms =
millisecond 10^{-3} , us = microsecond
 10^{-6}

But RAM loses content in power
failure! not persistent by itself.

Why not just write to RAM of f
machines, to tolerate f failures?



Non-Volatile RAM

Batteries power the machines for a while if main power fails

s/w knows when main power fails

s/w halts all transaction processing

s/w writes FaRM's RAM to SSD; may take a few minutes then machine shuts down cleanly

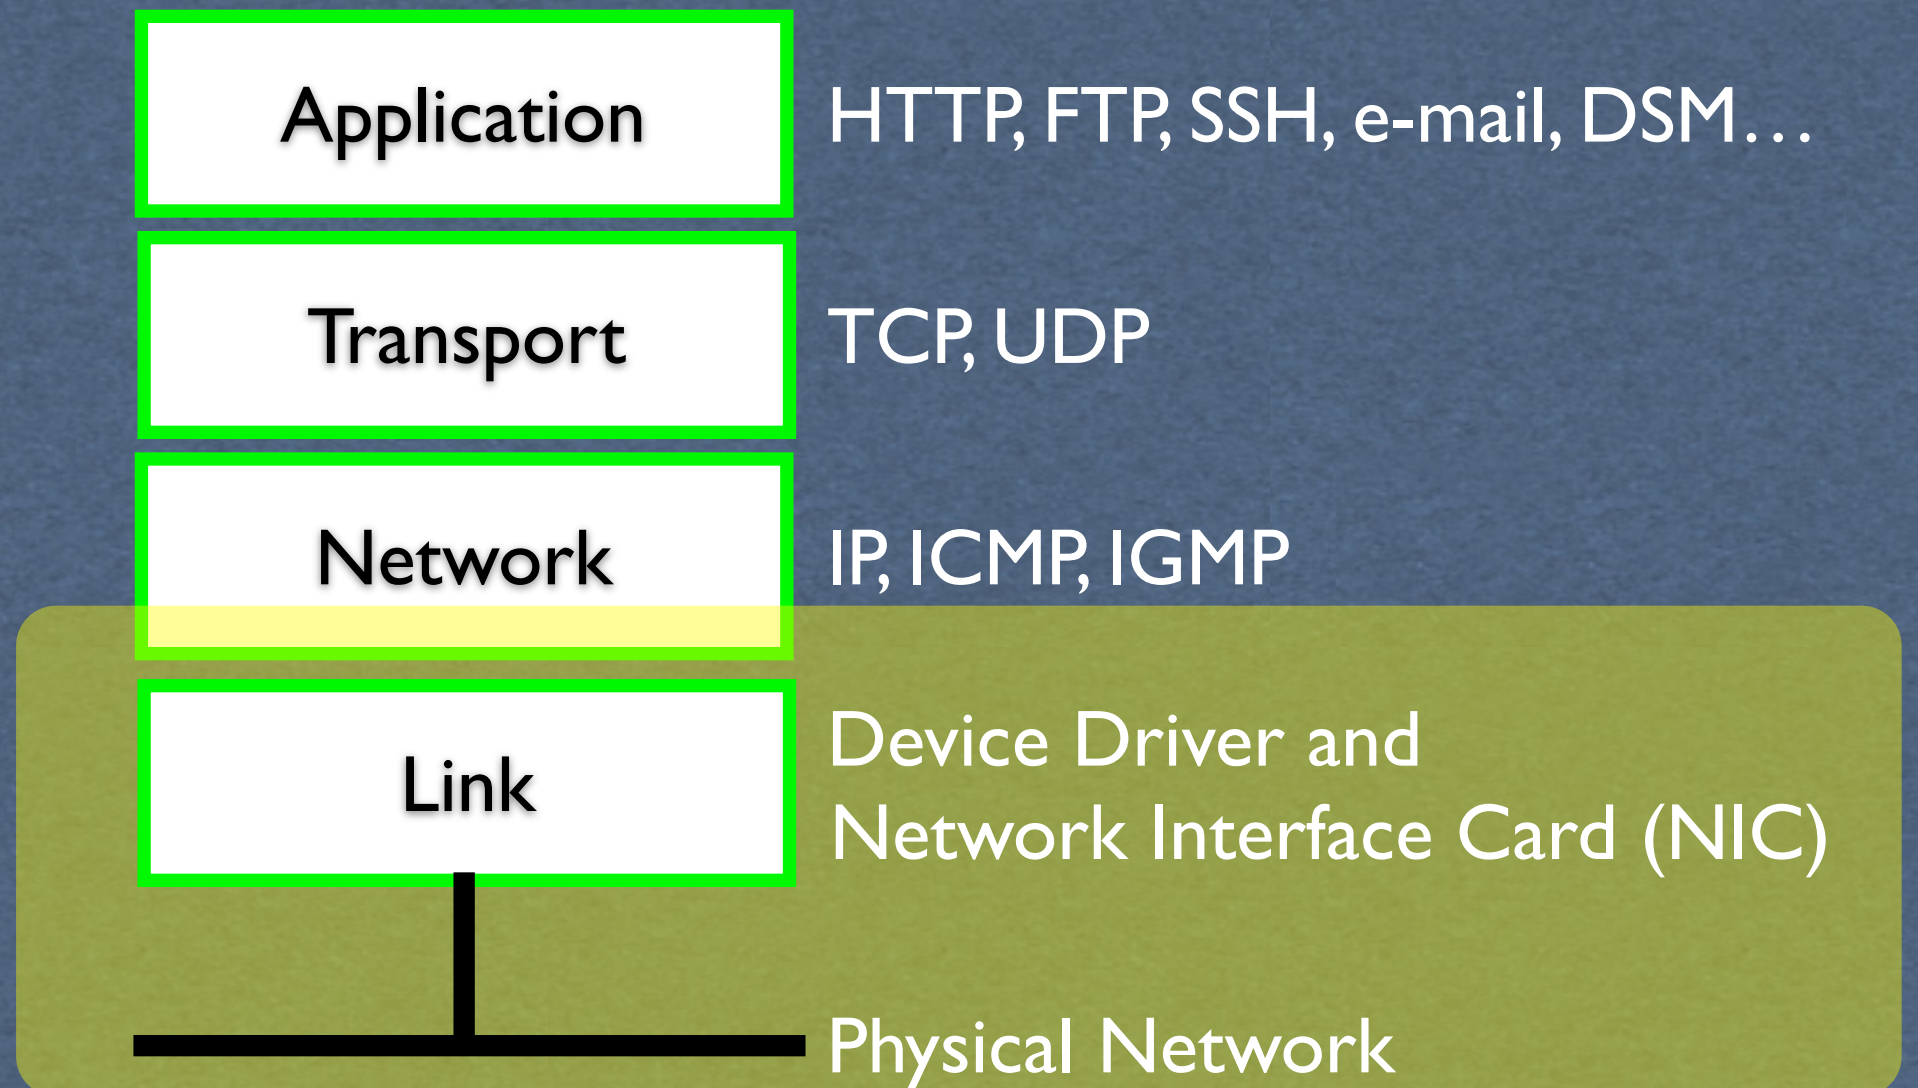
On re-start, FaRM reads saved memory image from SSD

What if crash prevents s/w from writing SSD? e.g kernel bug, or cpu/memory/hardware error

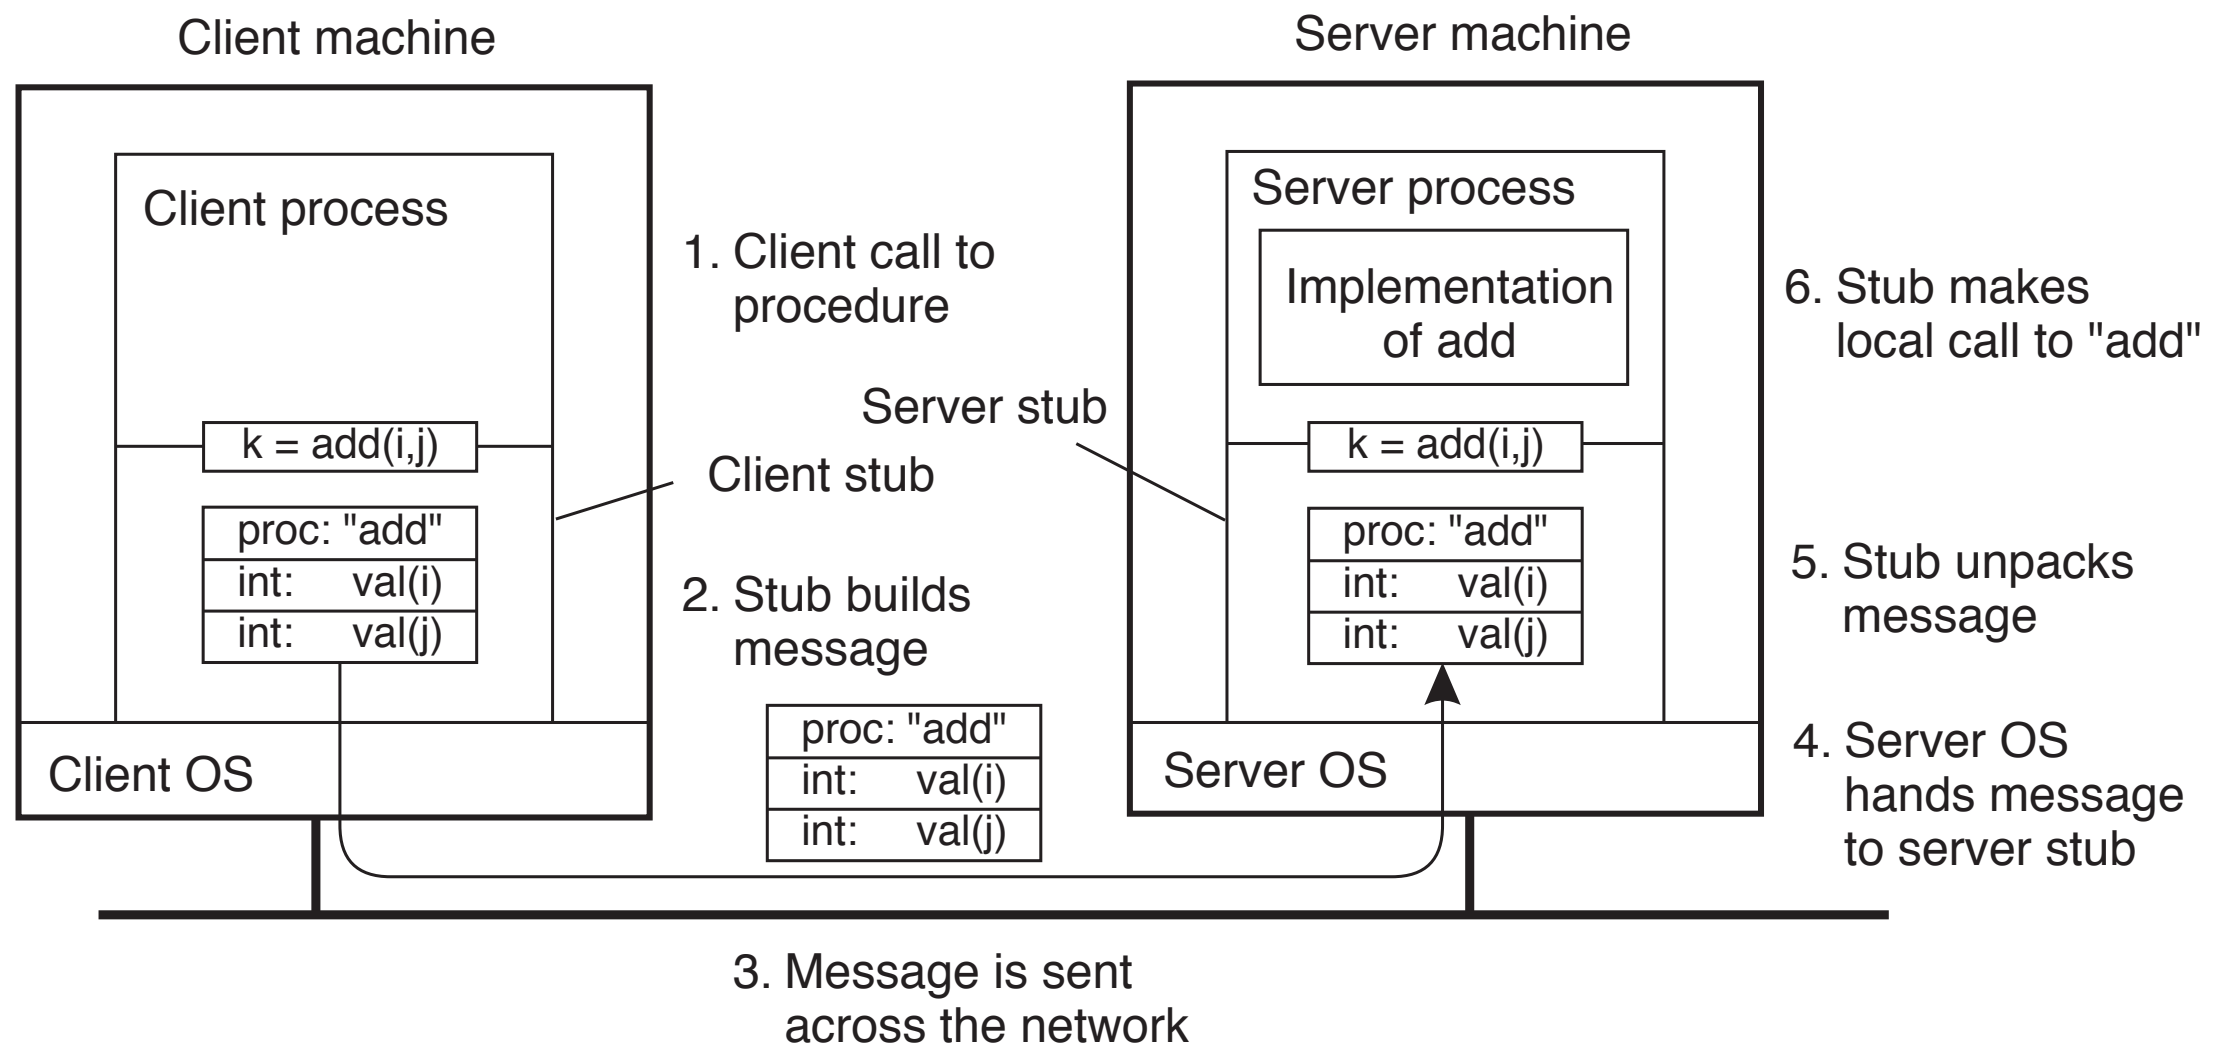


Networking

The Layers of the TCP/IP protocol suite.





RPC: Under the covers




- 1 Client procedure calls client stub.
- 2 Stub builds message; calls local OS.
- 3 OS sends message to remote OS.
- 4 Remote OS gives message to stub.
- 5 Stub unpacks parameters and calls server.
- 6 Server returns result to stub.
- 7 Stub builds message; calls OS.
- 8 OS sends message to client's OS.
- 9 Client's OS gives message to stub.
- 10 Client stub unpacks result and returns to the client.

What's wrong with this
story?

Memory Access



RDMA



One-Sided RDMA

NIC does "one-sided RDMA": memory read/write,
not packet delivery

Sender says "write this data at this address", or "read
this address"

NIC *hardware* executes at the far end
returns a "hardware acknowledgement"

FaRM

Distributed Transactions — Designed to exploit modern data center hardware

Distributed Transactions — single machine that executes operations on memory objects.

Optimistic Concurrency Control

Strict Serialisability

Fault-tolerance

Transaction

- Start transaction
- Read and Write memory objects
- Commit either succeeds or fails


```
TxBegin()  
    a = TxRead(oidA);  
    b = TxRead(oidB);  
    a.value -= amt;  
    b.value += amt;  
    TxWrite(oidA, a);  
    TxWrite(oidB, b);  
TxCommit();
```

Real FaRM API more subtle — and Event Driven

3 Principles

- Use one-sided RDMA operations
- Reduce message counts
- Exploit parallelism — to increase availability

FaRM (NSDI '14)




General Purpose Platform : KV Store, graph Store, OLTP databases, etc. — small objects, irregular access



All Machines clients and servers

Transactions: Read, Write, Alloc, Free, ...



FaRM



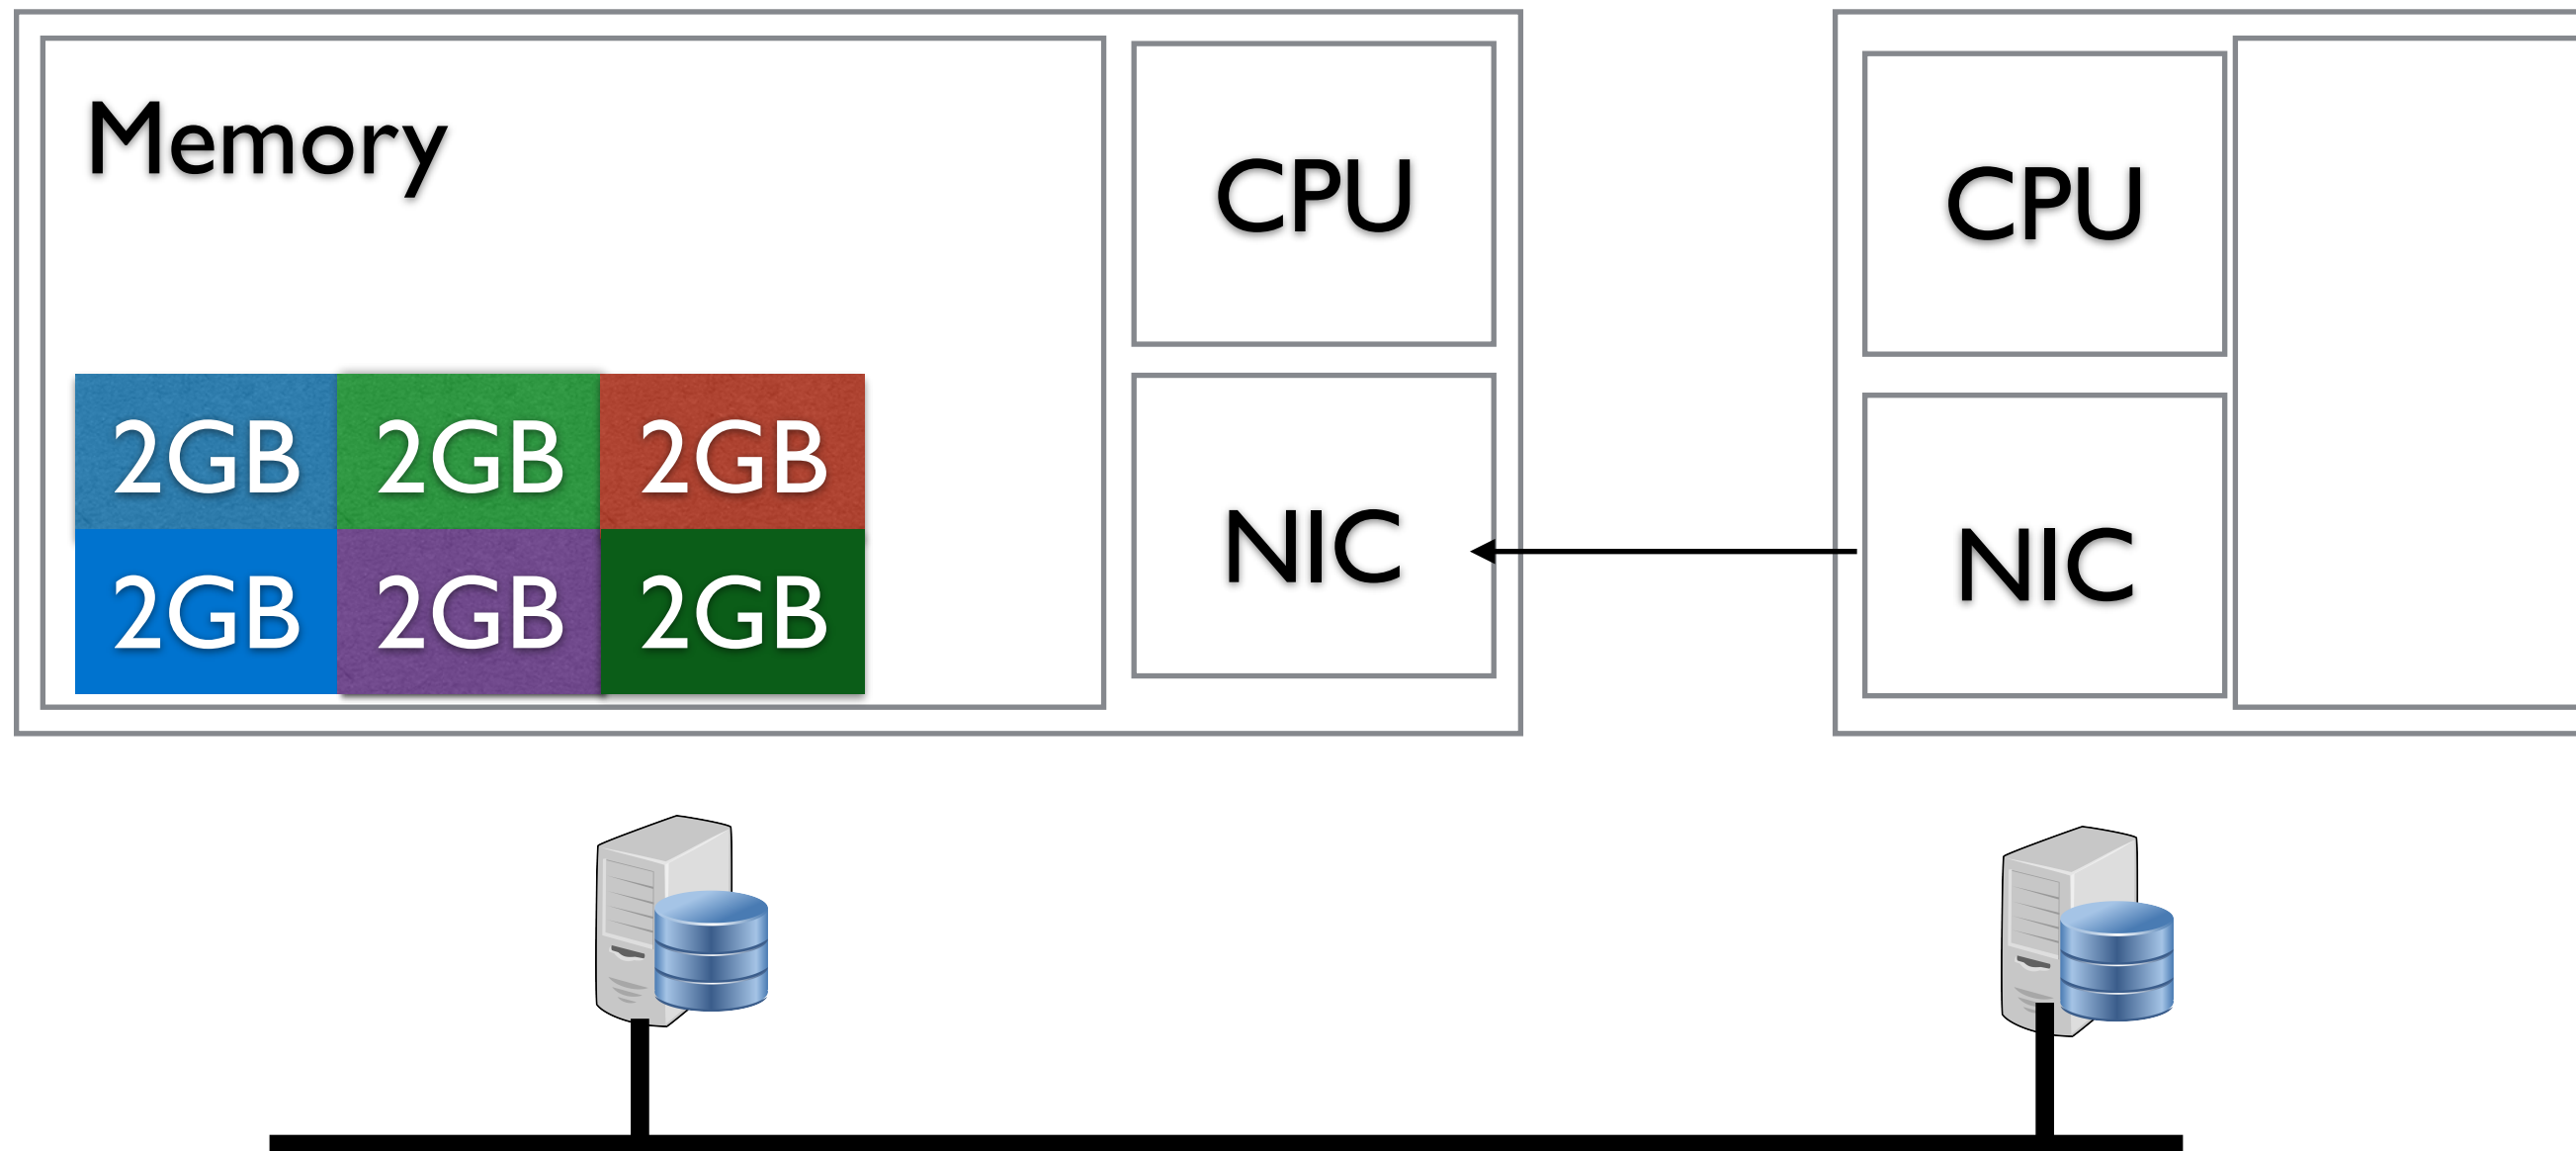
FaRM



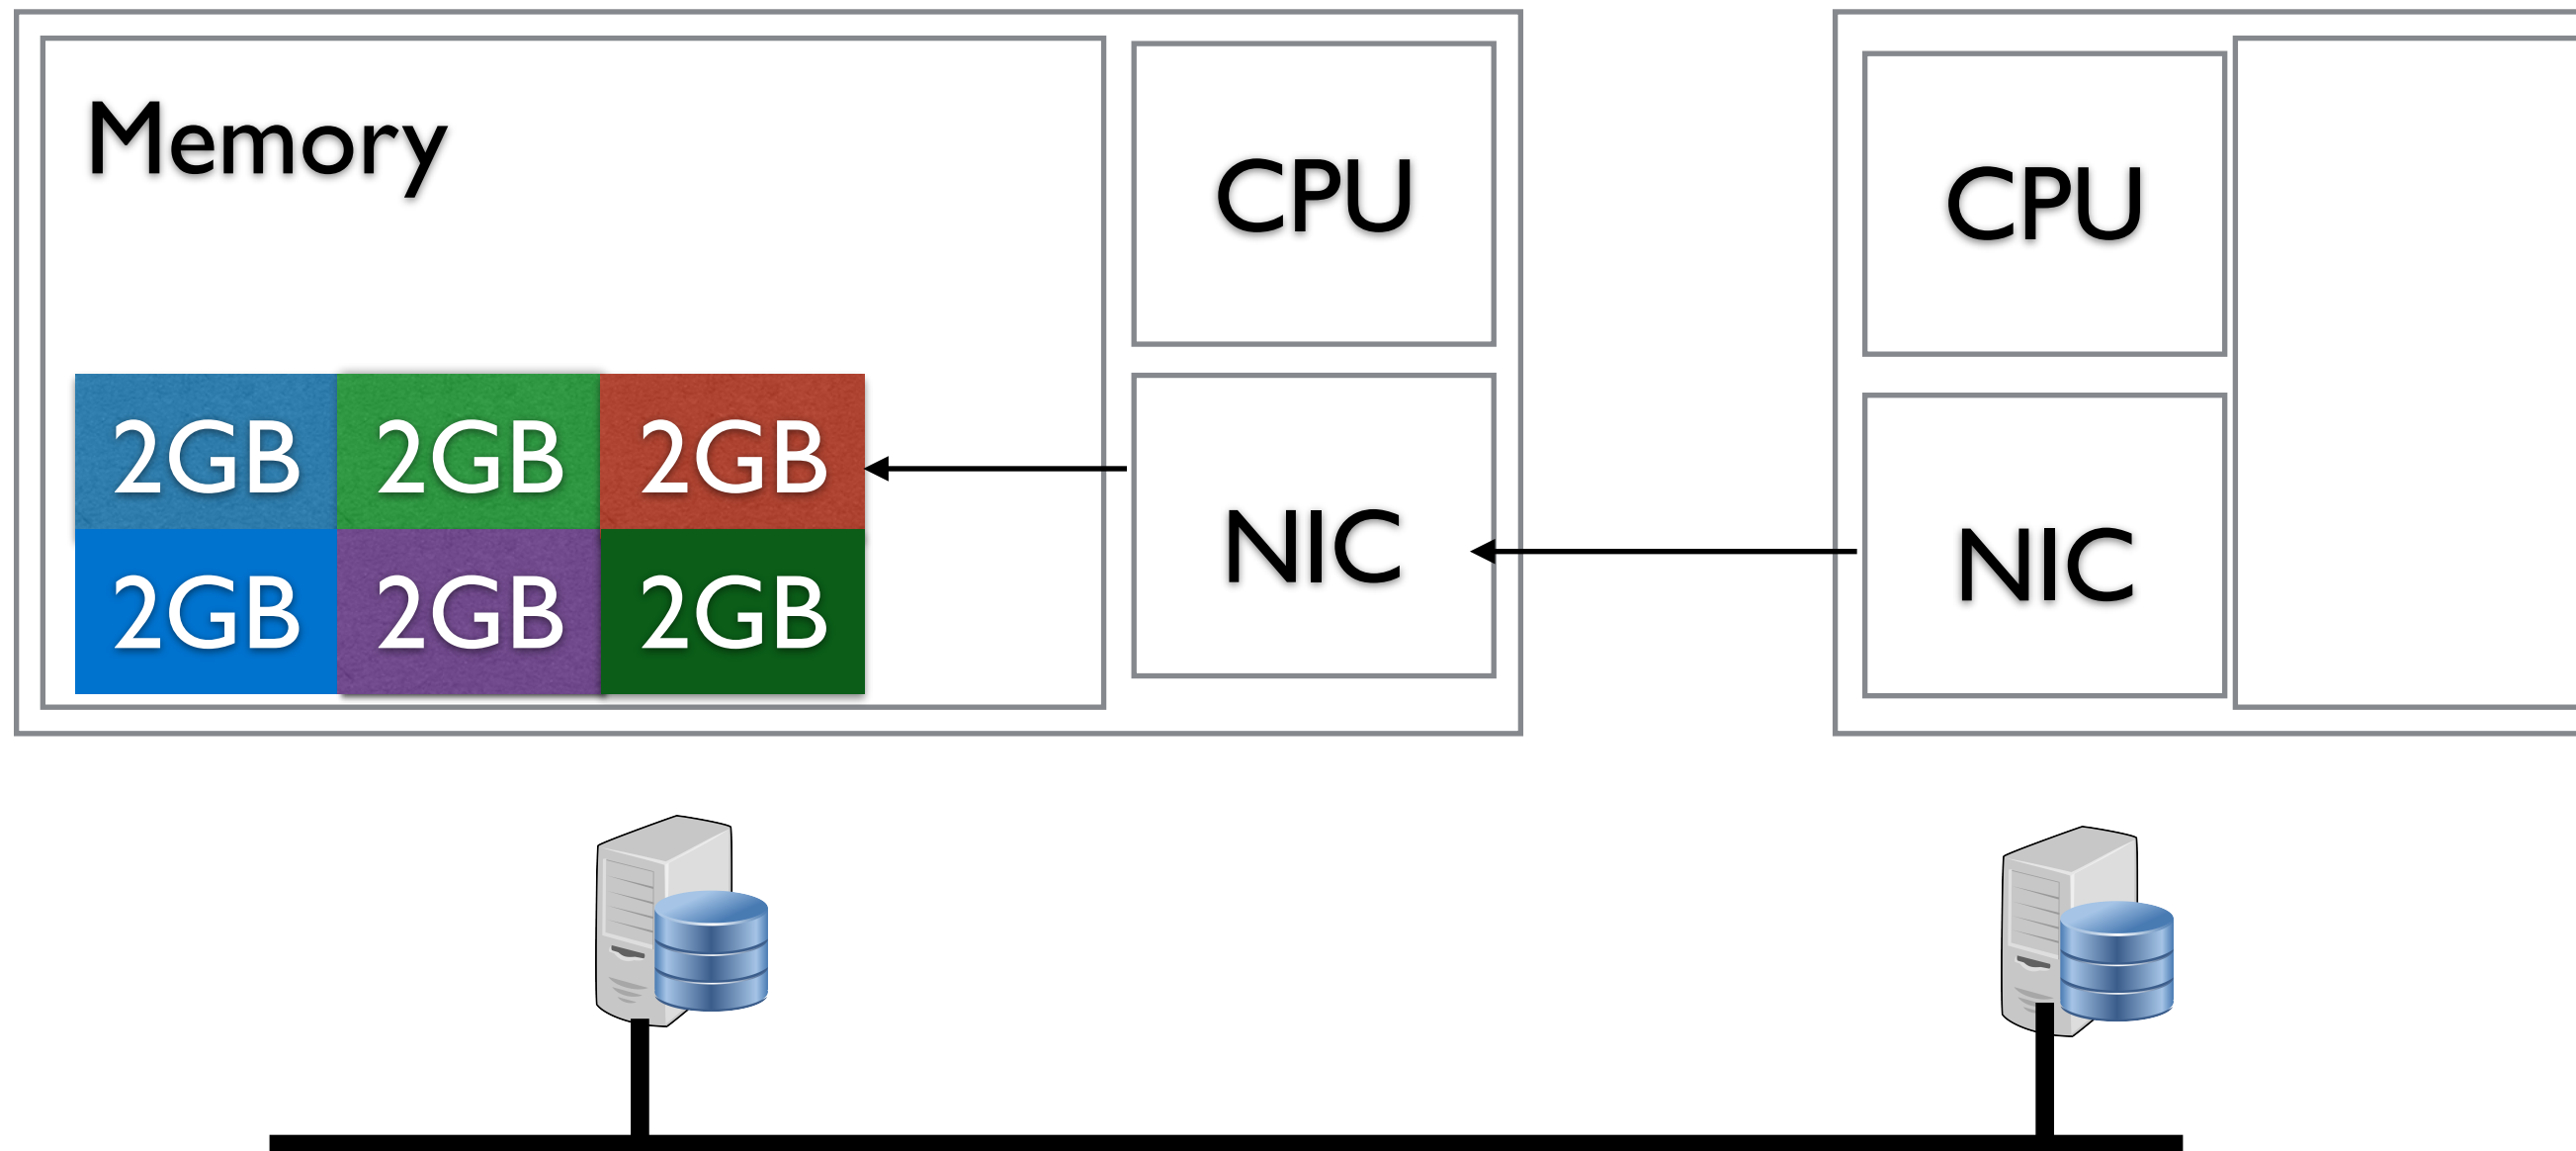
FaRM: txRead



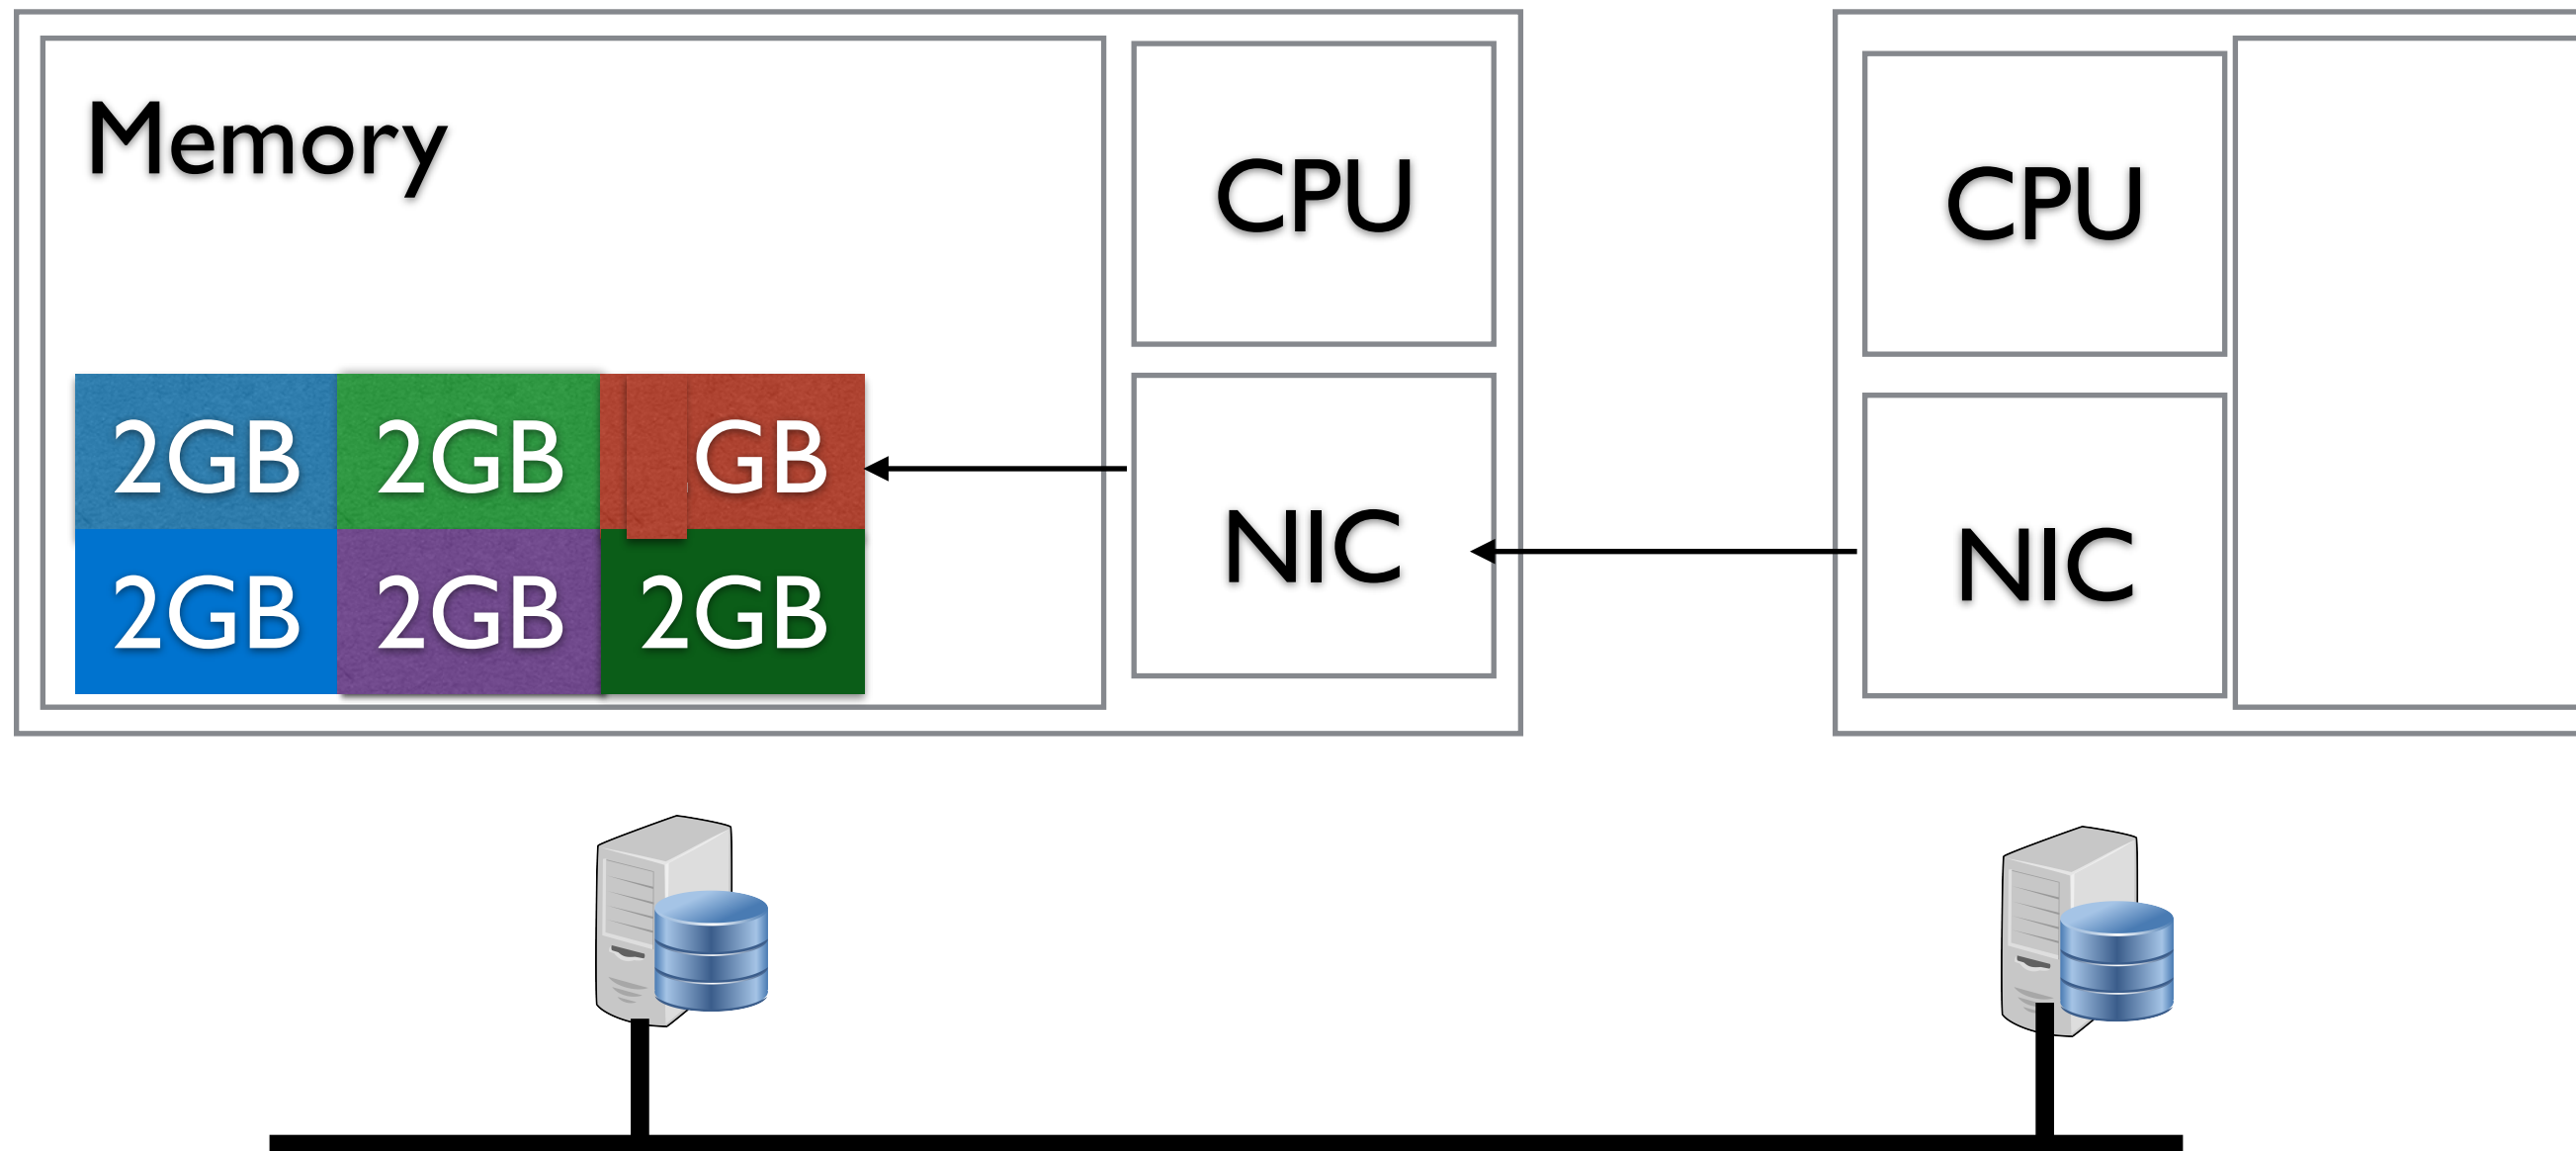
FaRM: txRead



FaRM: txRead




FaRM: txRead




No CPU involvement at “Server”

FaRM: txWrite

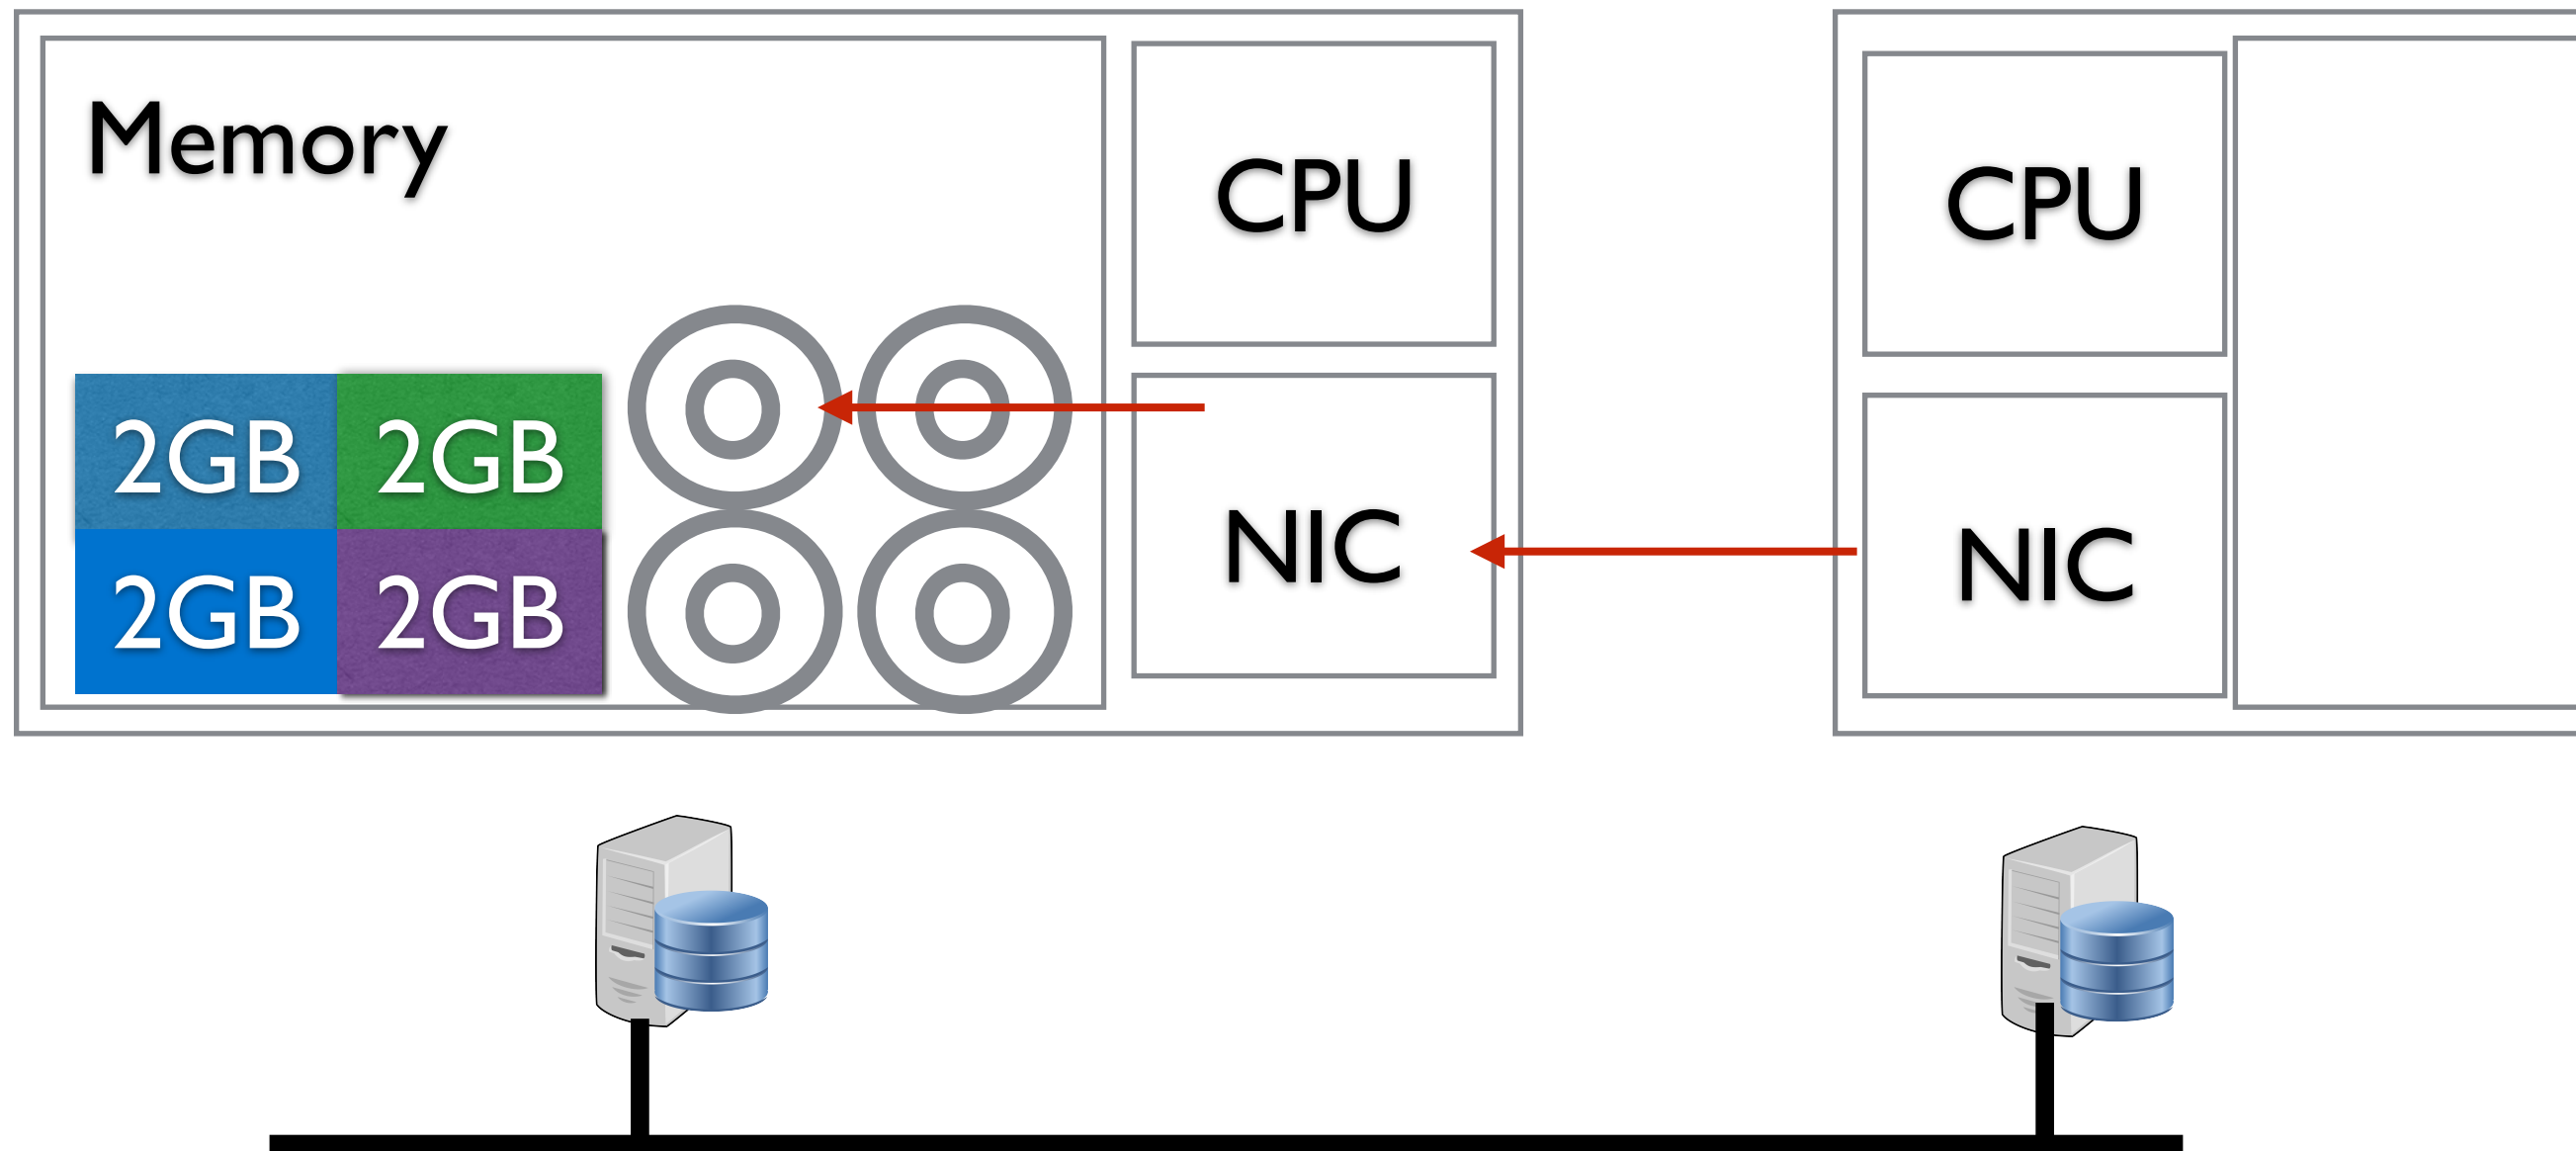


FaRM: txWrite



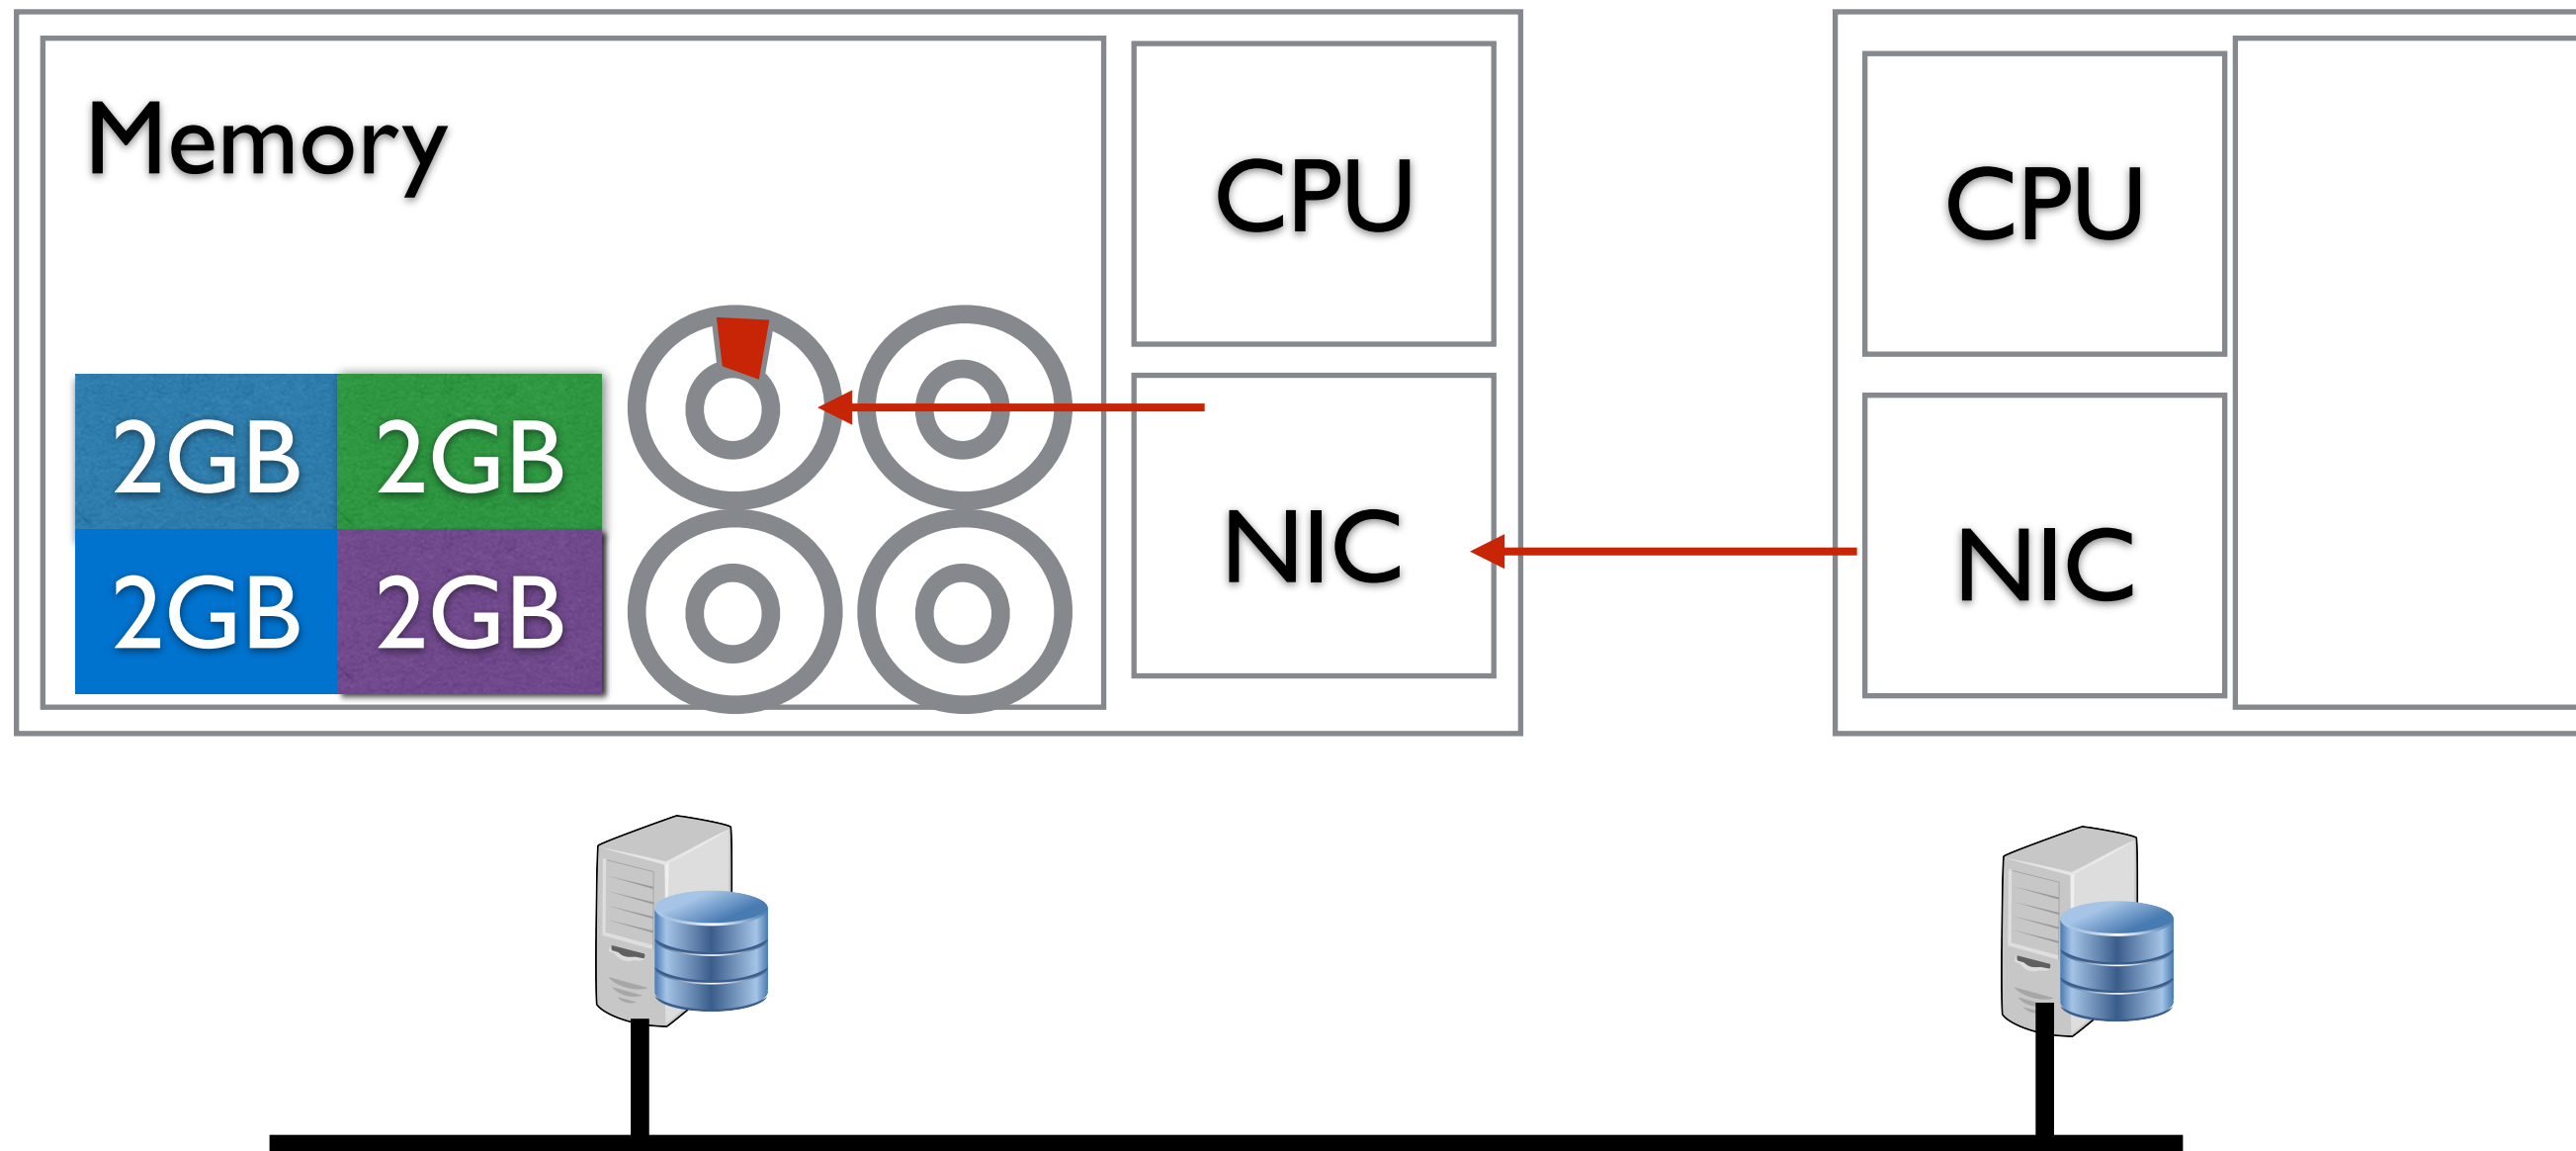
No CPU involvement at “Server”

FaRM: txWrite



No CPU involvement at “Server”

FaRM: txWrite



No CPU involvement at “Server”

Transactions

C



P1



B1




P2




B2




Transactions




Transactions




Transactions



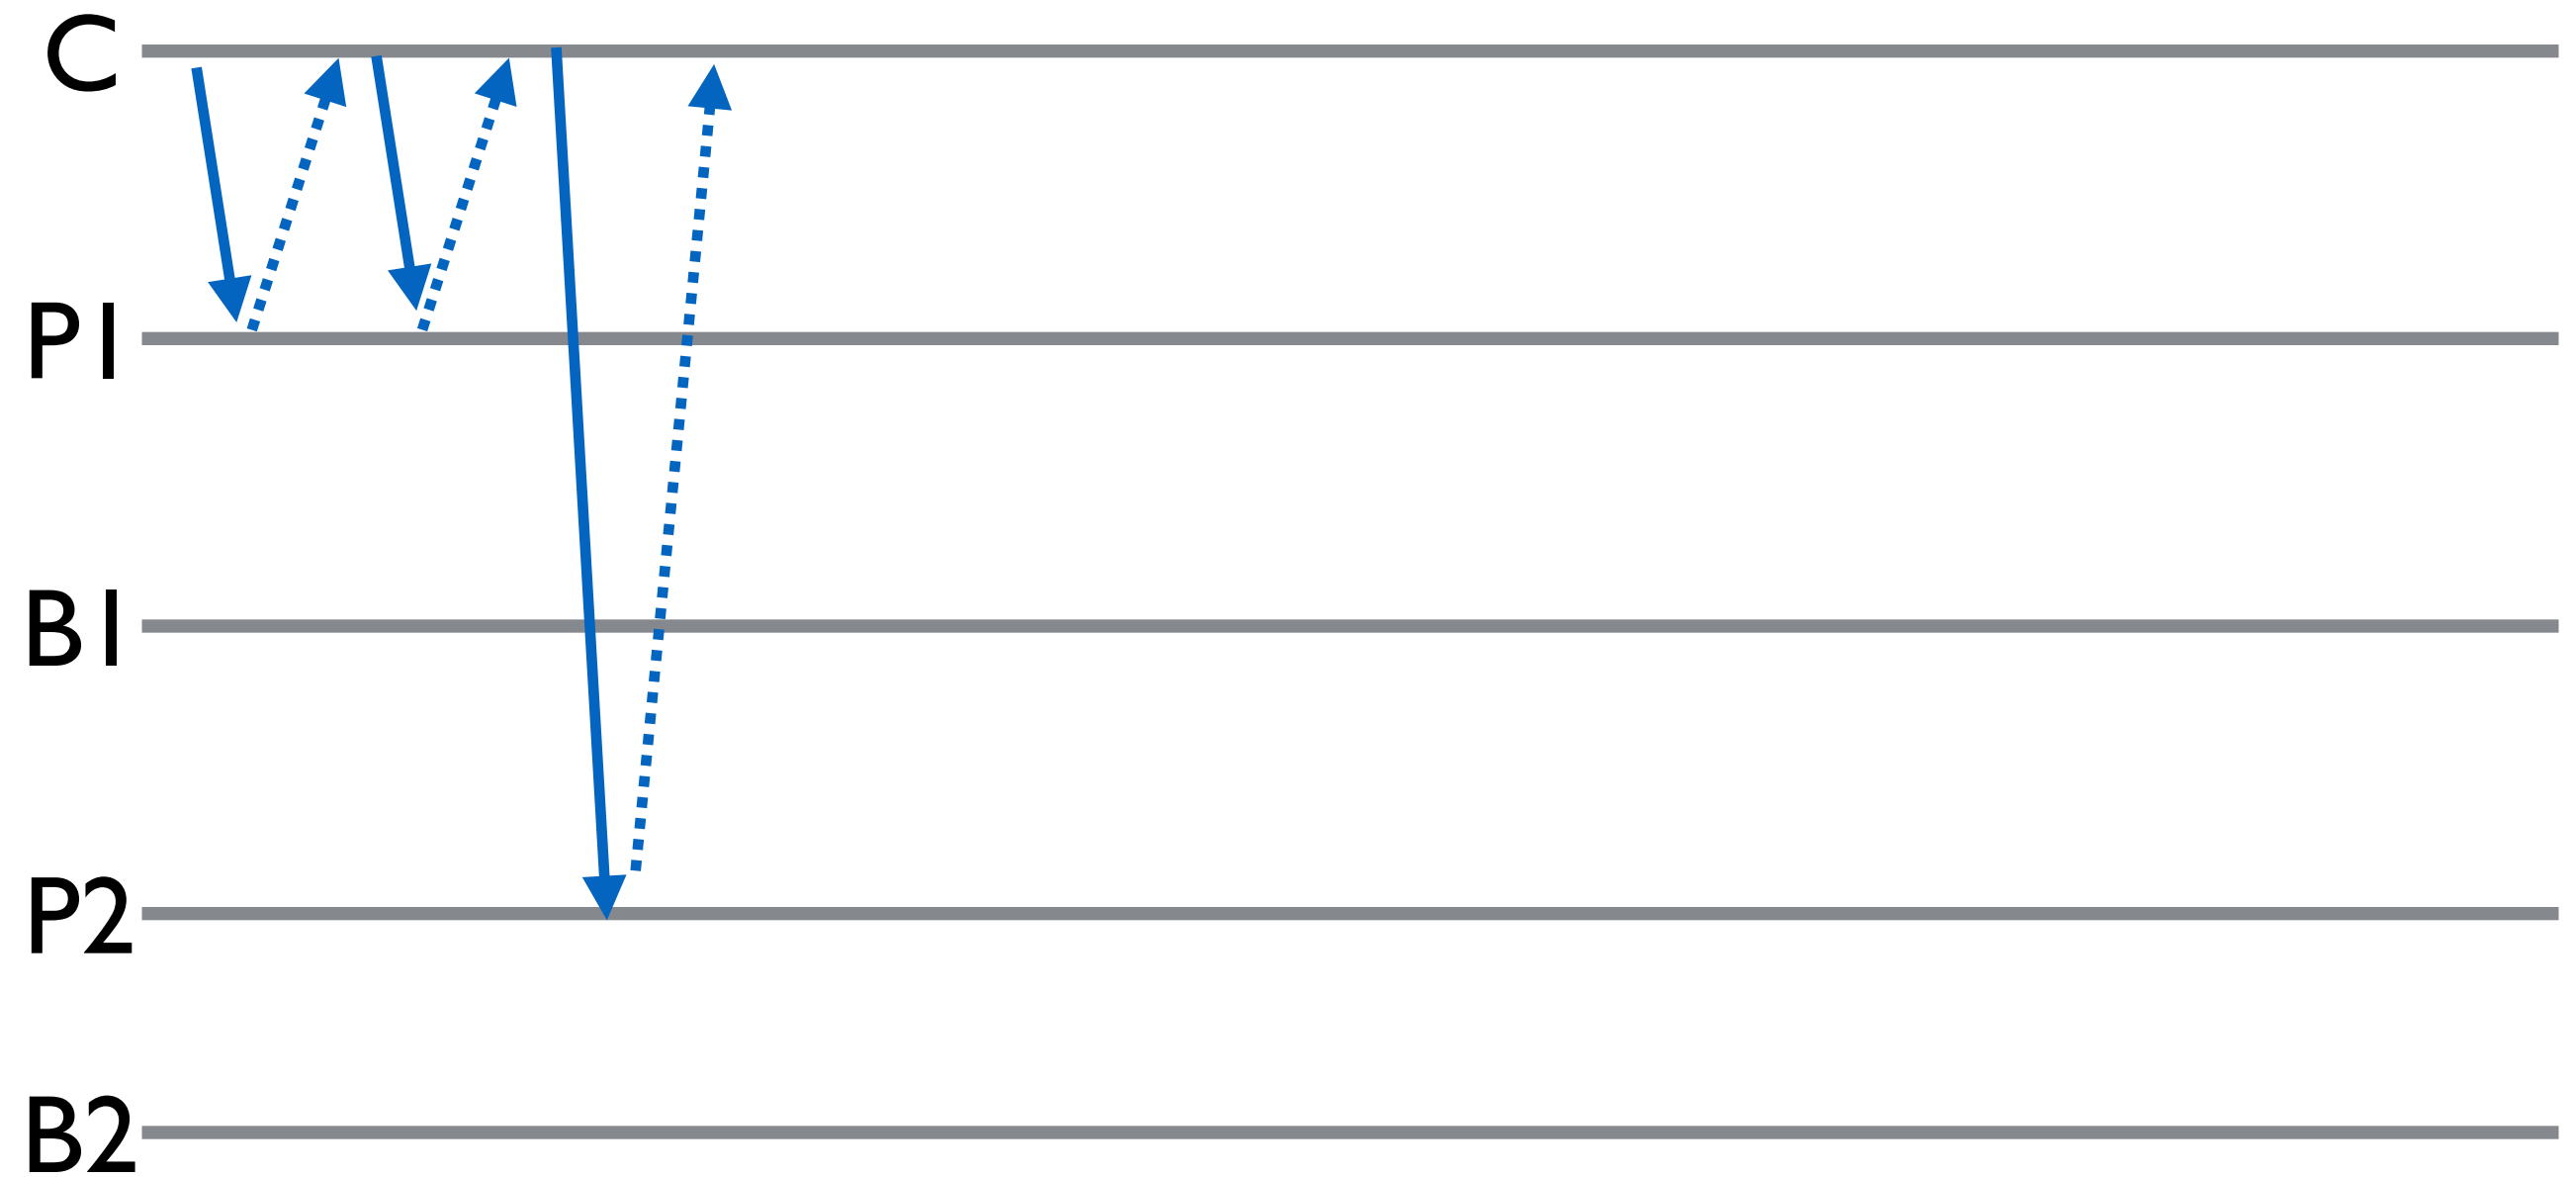
Transactions




Transactions




Transactions




Transactions




Transactions



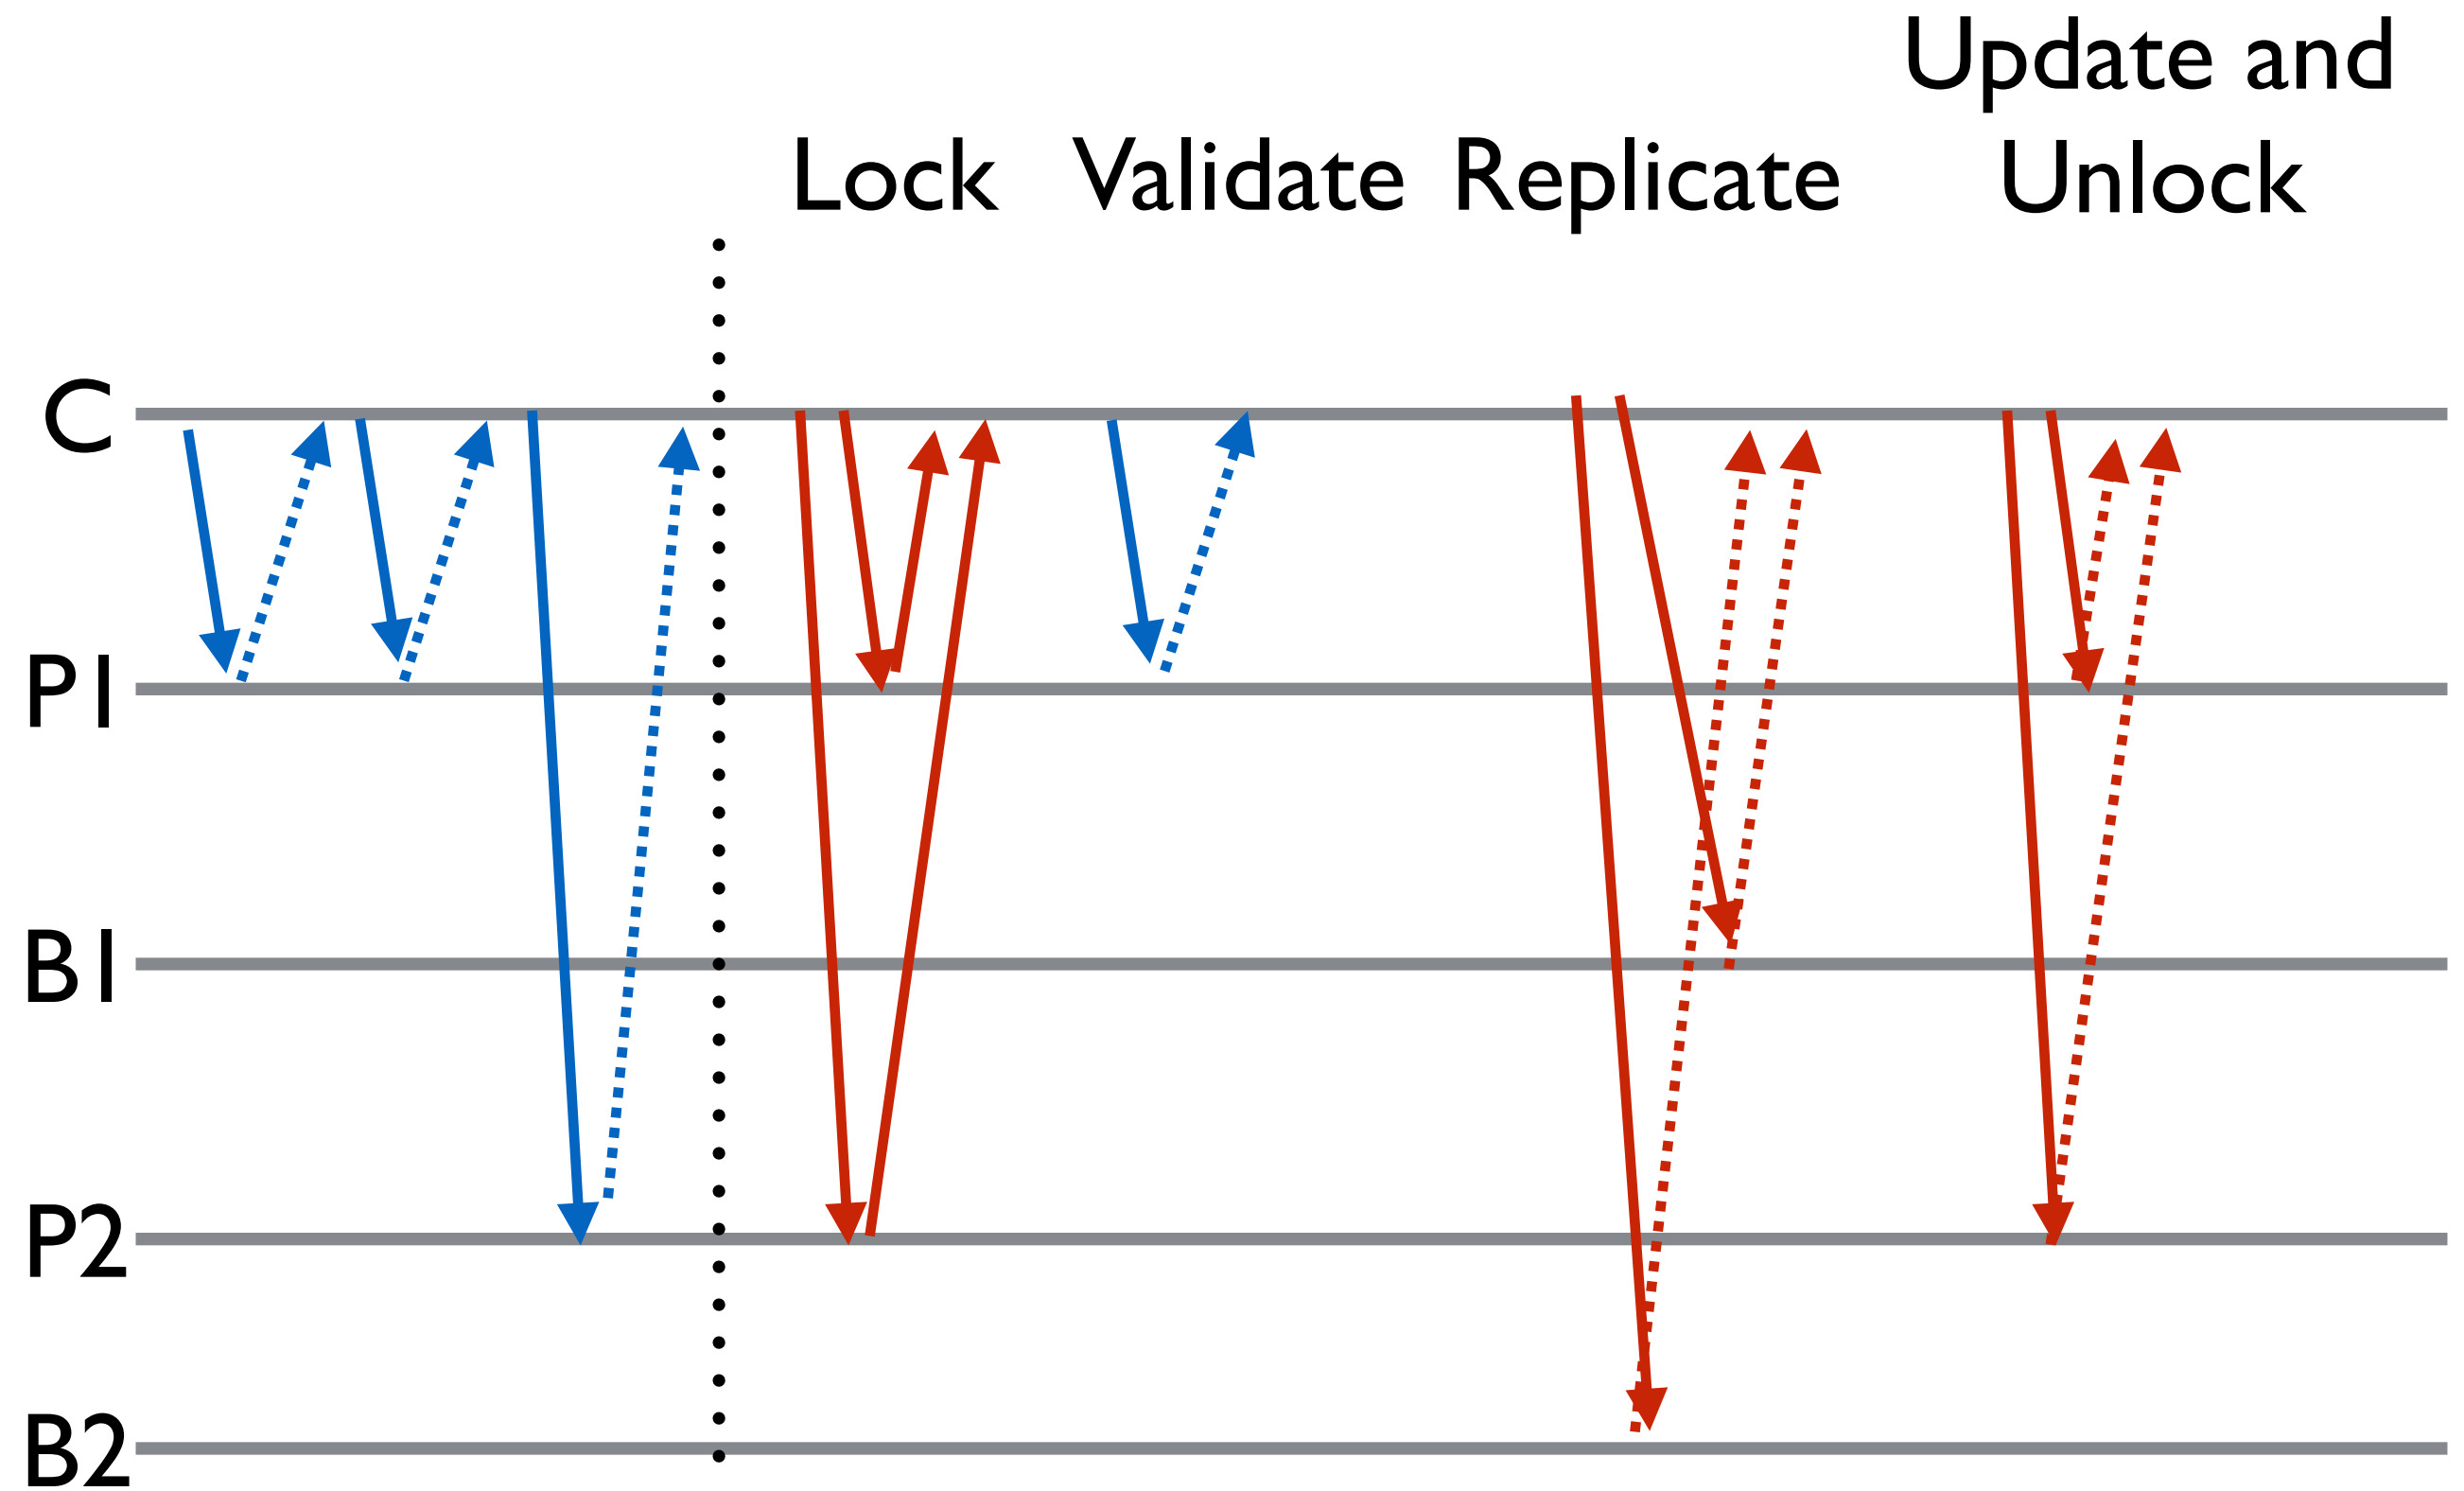
Transactions




Transactions



Transactions



Transactions



Failures and Recovery

- Complicated by one-side operations — no software running on message reception — can't reject
- They need to have a way of establishing exactly when and who can continue to send messages — “Precise membership”
- Recovery — transactions might have been reported as committed but the logs are not processed so must drain logs

High Availability

- Fast switch — backups are in memory so can switch fast
- Networks and implementation allows very small leases — 10ms
- lots of parallelism — unaffected regions can proceed and data recover in parallel