

Distributed Systems

Spring Semester 2020

Lecture 24: Ray



John Liagouris
liagos@bu.edu

Why this paper

- A general-purpose cluster computing framework
- Dynamic dataflow graphs
 - Needed for emerging AI applications
- Interesting programming model
 - Tasks — stateless functions
 - Actors — stateful computation
 - Futures — for communication between threads


Reinforcement Learning

- Process inputs from many sensors in parallel and in real-time
- Execute large number of simulations — 100s of millions
- Result of simulations and/or interaction with the environment can change the next step

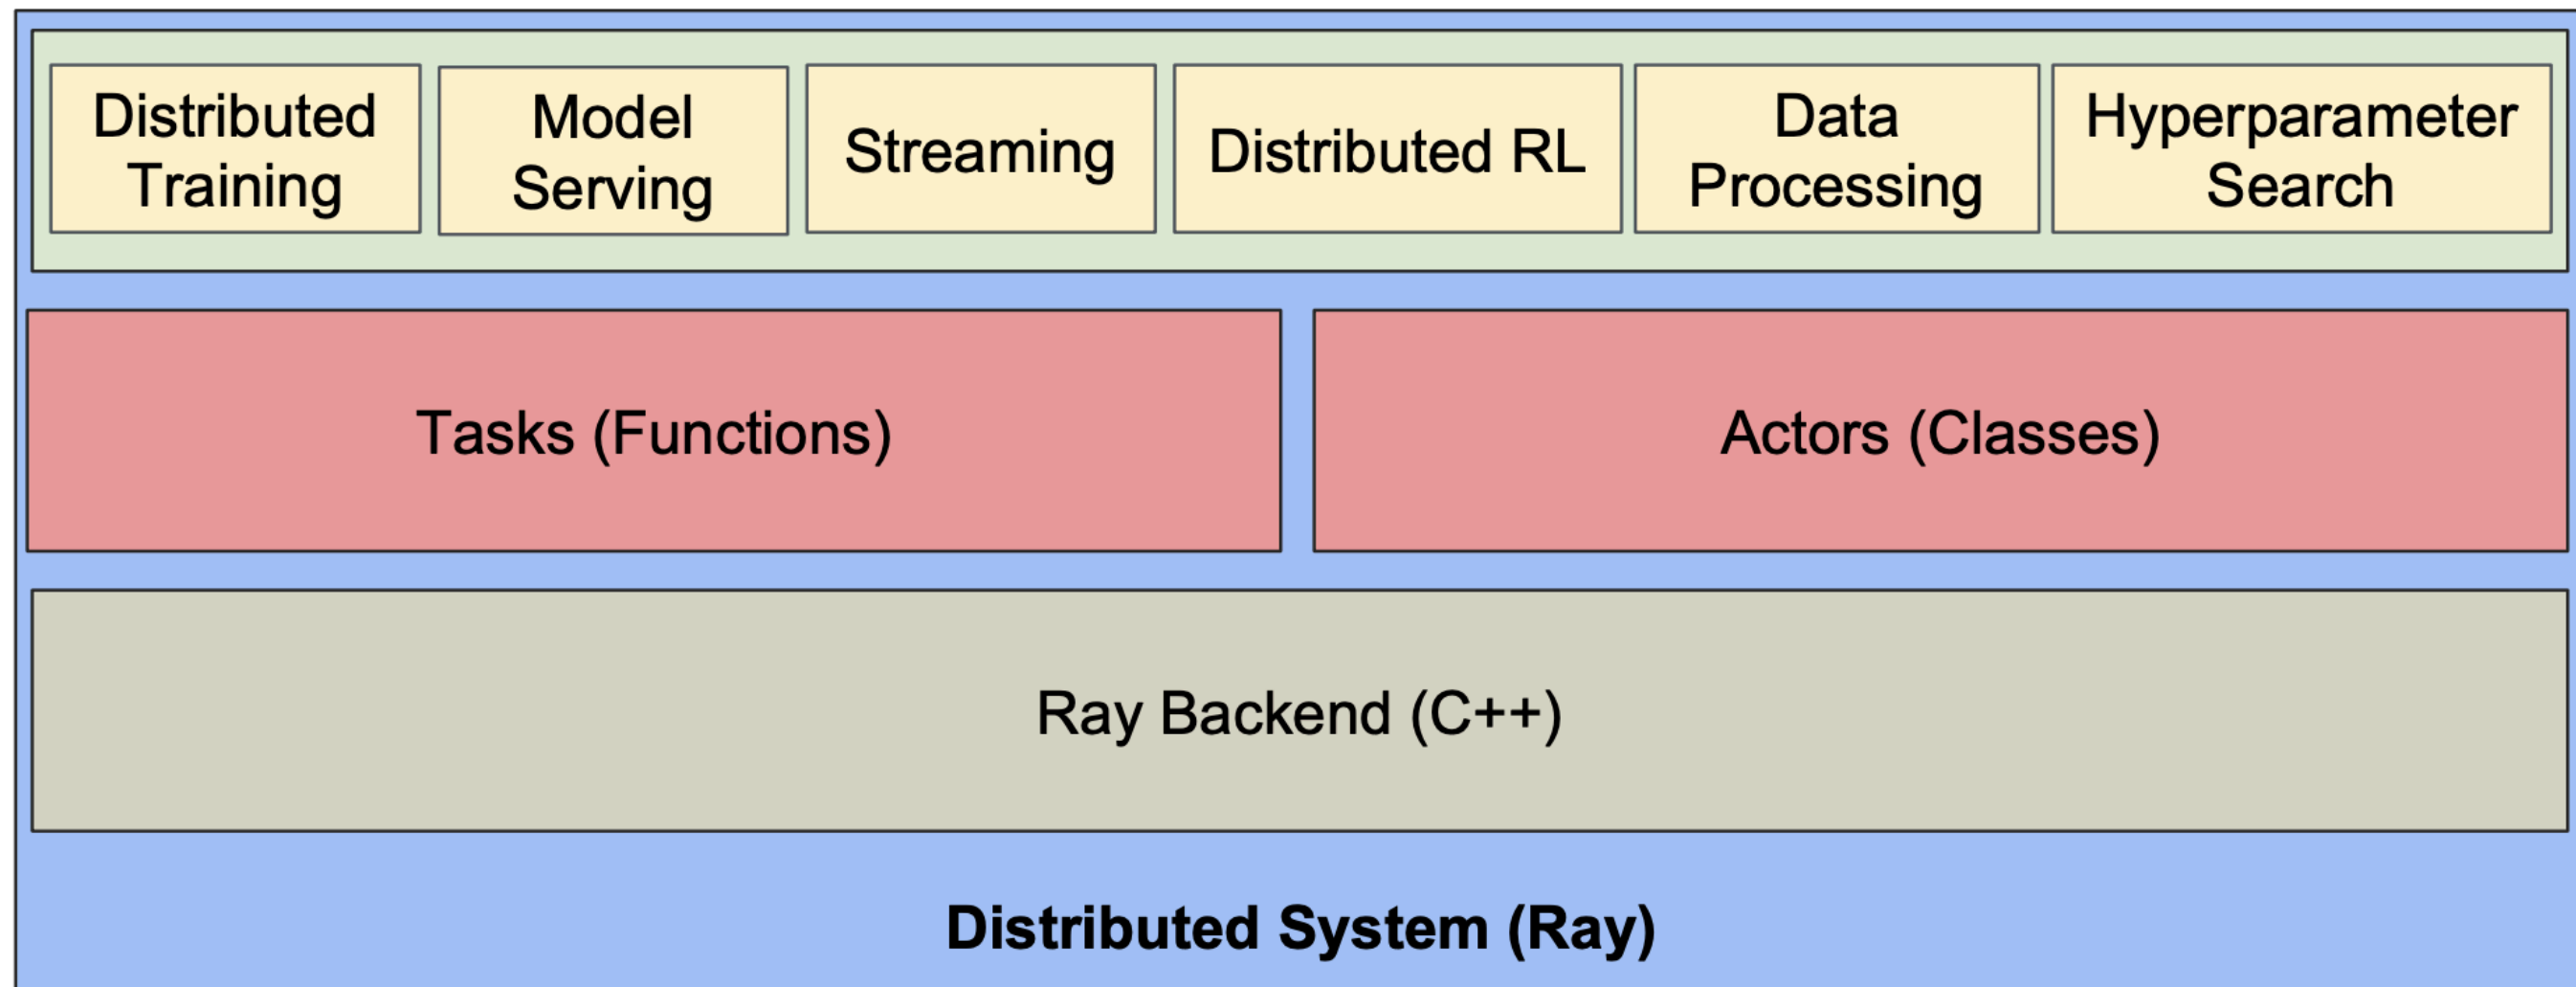


Reinforcement Learning

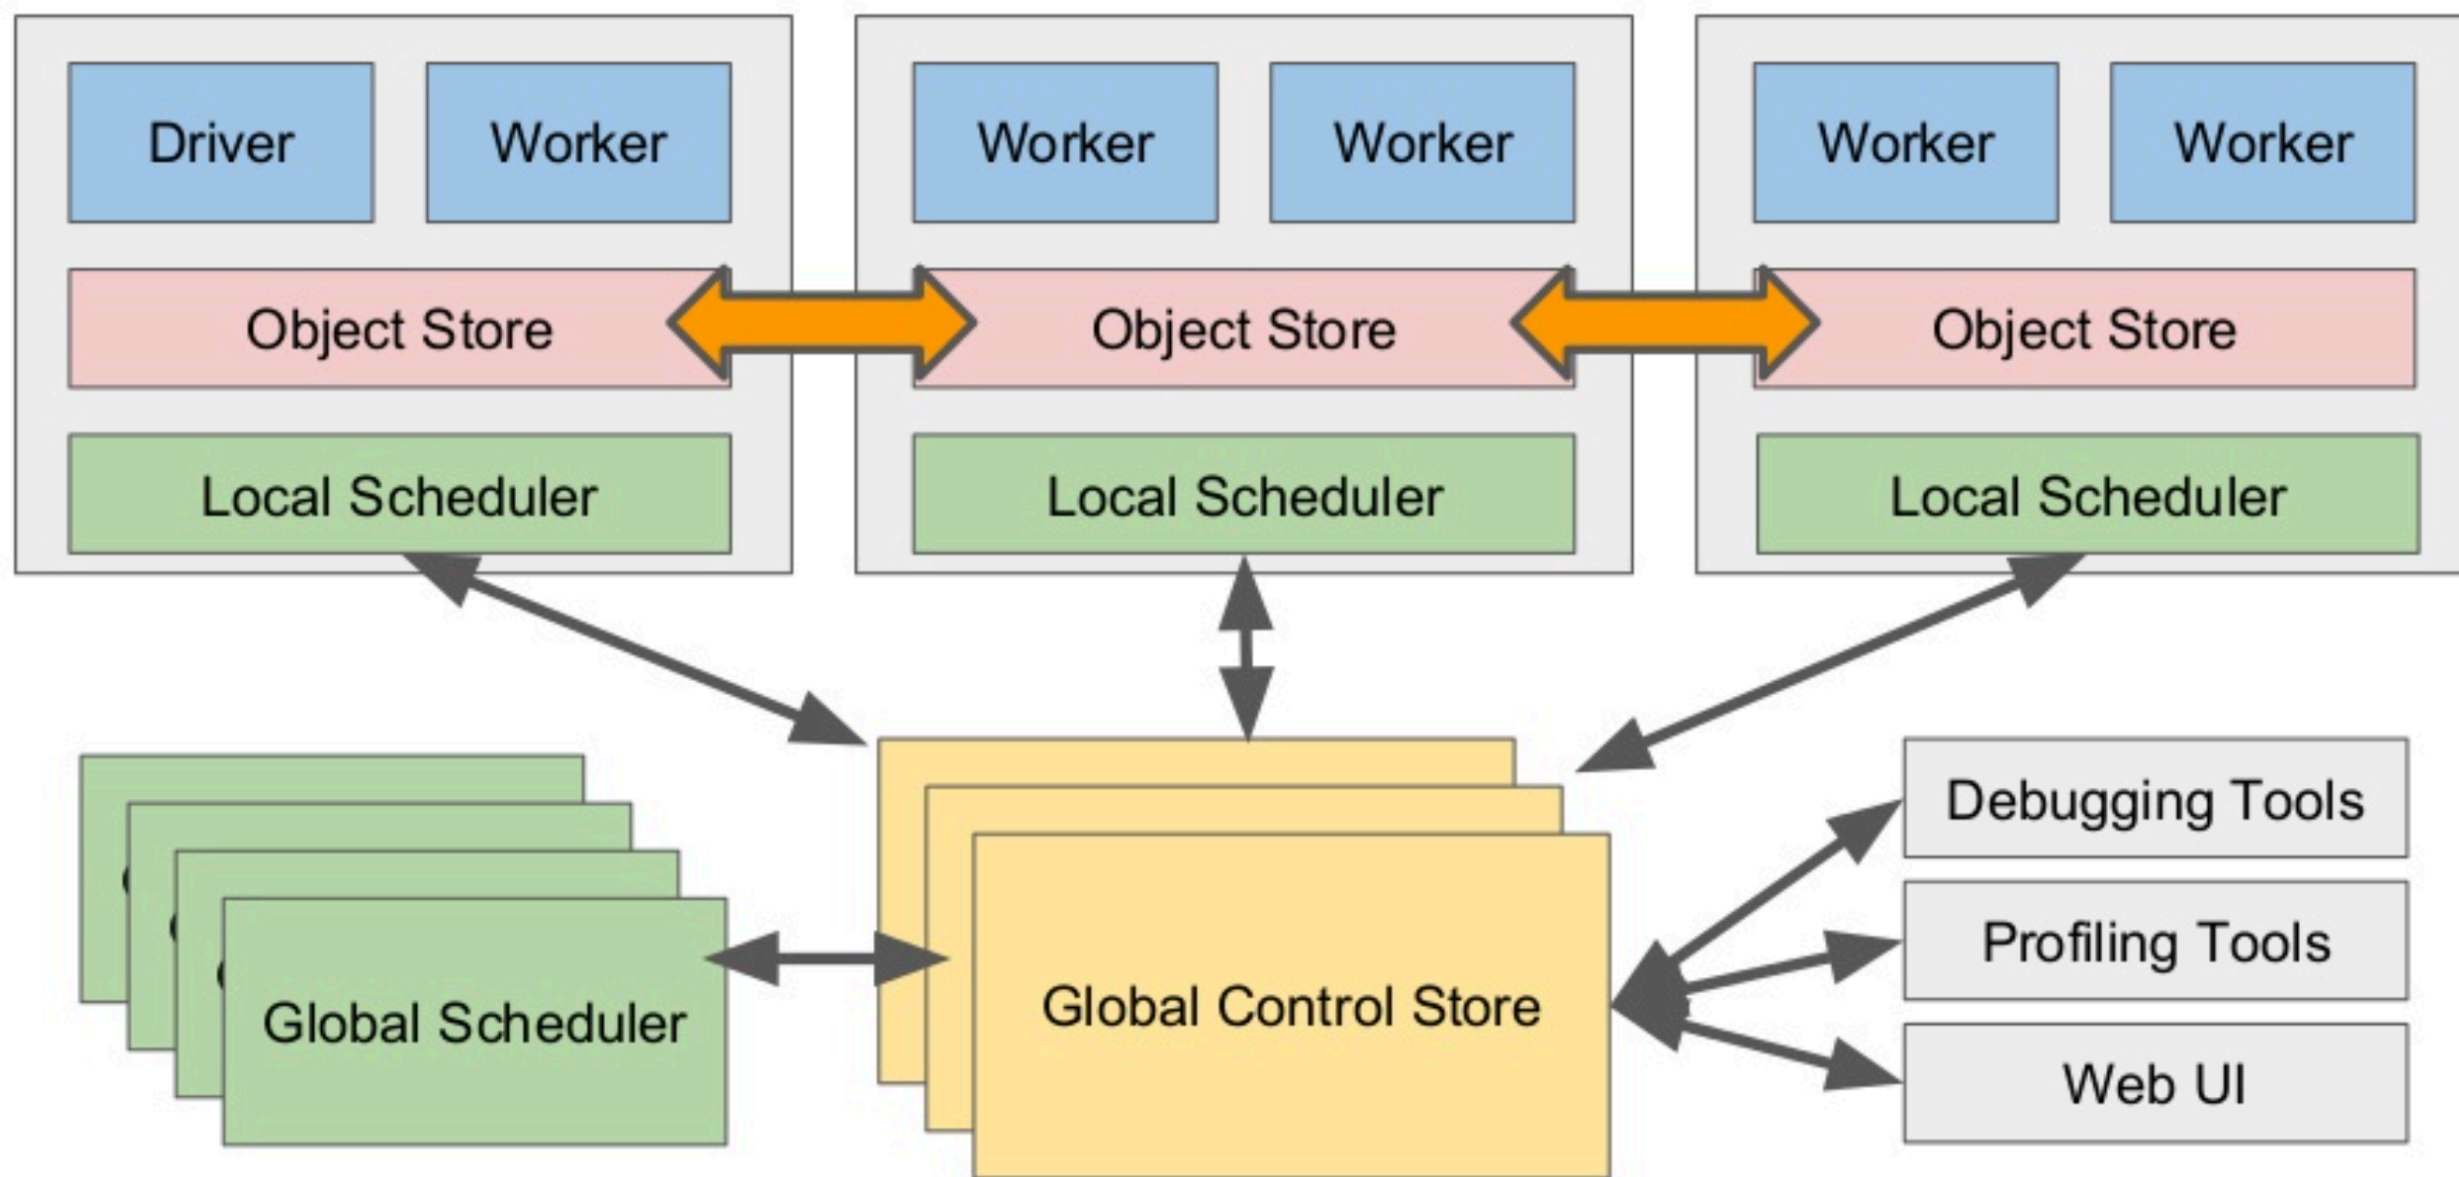
- Need to handle dynamic task graphs
 - Heterogenous durations
 - Heterogenous computations
- Scale millions of tasks per second



What is Ray?



Ray Architecture



Programming Model

- Tasks — stateless functions
- Actors — stateful objects
- Futures — for thread communication

...in Python!

Tasks

```
def zeros(shape):  
    return np.zeros(shape)
```

```
def dot(a, b):  
    return np.dot(a, b)
```


Tasks

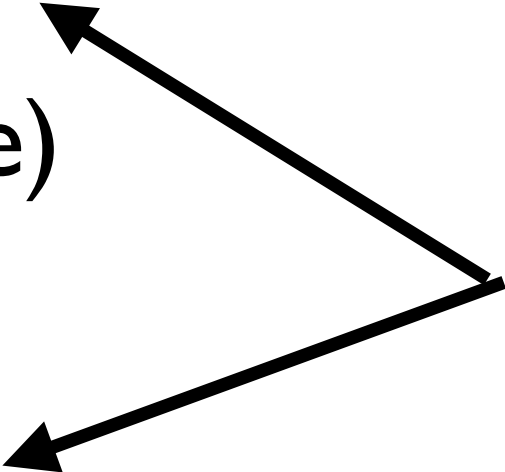
```
@ray.remote
```

```
def zeros(shape):  
    return np.zeros(shape)
```

```
@ray.remote
```

```
def dot(a, b):  
    return np.dot(a, b)
```

Can be also
placed explicitly
at a remote
machines



Tasks


```
@ray.remote
```

```
def zeros(shape):  
    return np.zeros(shape)
```

```
@ray.remote
```

```
def dot(a, b):  
    return np.dot(a, b)
```

```
id1 = zeros.remote([5, 5])
```



Future: Encapsulates a value
that is not available yet

Tasks


```
@ray.remote
```

```
def zeros(shape):  
    return np.zeros(shape)
```

```
@ray.remote
```

```
def dot(a, b):  
    return np.dot(a, b)
```

```
id1 = zeros.remote([5, 5])
```



Non-blocking call
zeros() is executed by a
different thread

Tasks

```
@ray.remote
```


```
def zeros(shape):  
    return np.zeros(shape)
```

```
@ray.remote
```

```
def dot(a, b):  
    return np.dot(a, b)
```

```
id1 = zeros.remote([5, 5])
```

```
id2 = zeros.remote([5, 5])
```



Tasks

```
@ray.remote
```

```
def zeros(shape):  
    return np.zeros(shape)
```

```
@ray.remote
```


```
def dot(a, b):  
    return np.dot(a, b)
```

```
id1 = zeros.remote([5, 5])
```

```
id2 = zeros.remote([5, 5])
```

```
id3 = dot.remote(id1, id2)
```

Futures can be chained!



Tasks

```
@ray.remote
```

```
def zeros(shape):  
    return np.zeros(shape)
```

```
@ray.remote
```


```
def dot(a, b):  
    return np.dot(a, b)
```

```
id1 = zeros.remote([5, 5])
```

```
id2 = zeros.remote([5, 5])
```

```
id3 = dot.remote(id1, id2)
```

```
ray.get(id3) ← Blocking call
```

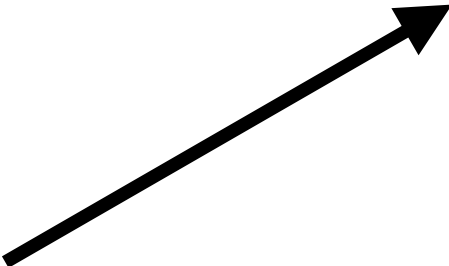


Actors

```
class Counter(object):  
    def __init__(self):  
        self.value = 0  
  
    def inc(self):  
        self.value += 1  
        return self.value
```

Actors

Can be also
placed explicitly
at a remote
machine



```
@ray.remote(num_gpus=1)
```

```
class Counter(object):
```

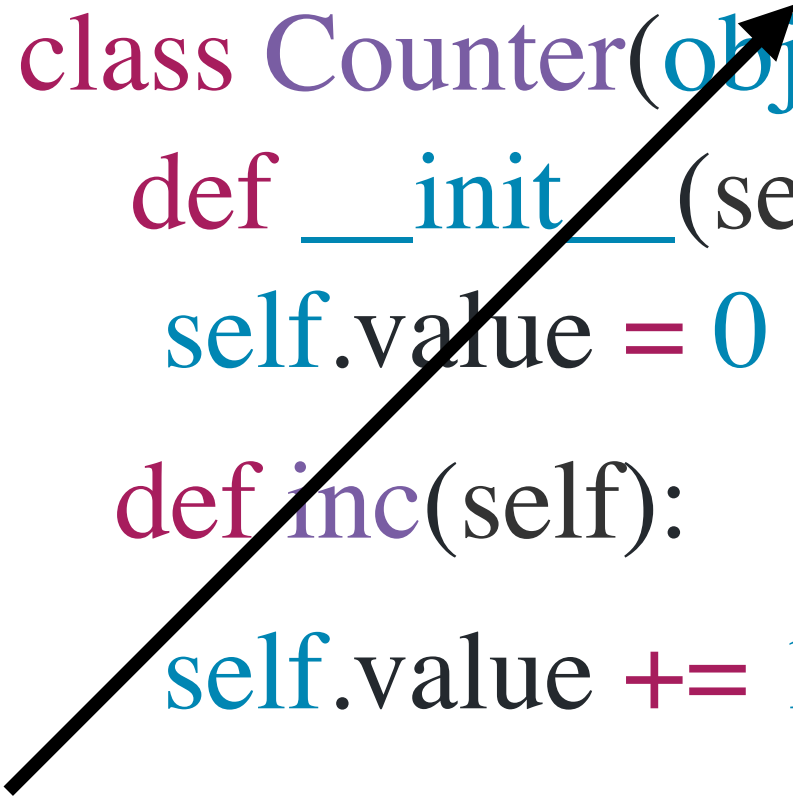
```
    def __init__(self):
```

```
        self.value = 0
```

```
    def inc(self):
```

```
        self.value += 1
```

```
    return self.value
```



We can specify
resources needed for
the actor, e.g., 1 GPU

Actors

Counter

Returns a
handle to the
actor instance

```
@ray.remote(num_gpus=1)
```

```
class Counter(object):
```

```
    def __init__(self):
```


```
        self.value = 0
```

```
    def inc(self):
```


```
        self.value += 1
```

```
    return self.value
```

```
c = Counter.remote()
```



Actors



➔ Stateful dependency

```
@ray.remote(num_gpus=1)
```

```
class Counter(object):
```

```
    def __init__(self):
```

```
        self.value = 0
```

```
    def inc(self):
```


```
        self.value += 1
```

```
        return self.value
```

```
c = Counter.remote()
```

```
id4 = c.inc.remote()
```

Actors



```
@ray.remote(num_gpus=1)
```

```
class Counter(object):
```

```
    def __init__(self):
```

```
        self.value = 0
```

```
    def inc(self):
```

```
        self.value += 1
```


```
        return self.value
```

```
c = Counter.remote()
```

```
id4 = c.inc.remote()
```

```
id5 = c.inc.remote()
```

Actors



```
@ray.remote(num_gpus=1)
```

```
class Counter(object):
```

```
    def __init__(self):
```

```
        self.value = 0
```

```
    def inc(self):
```

```
        self.value += 1
```

```
        return self.value
```

```
c = Counter.remote()
```

```
id4 = c.inc.remote()
```

```
id5 = c.inc.remote()
```

```
ray.get([id4, id5])
```

API

Name
<i>futures</i> = f.remote (<i>args</i>)
<i>objects</i> = ray.get (<i>futures</i>)
<i>ready_futures</i> = ray.wait (<i>futures</i> , <i>k</i> , <i>timeout</i>)
<i>actor</i> = Class.remote (<i>args</i>) <i>futures</i> = <i>actor.method.remote</i> (<i>args</i>)

How would you implement Raft in Ray?

```
class Raft(object):  
    def __init__(self):  
        self.peers = N  
  
        ...  
  
    def start_election(self):  
  
        ...  
  
    def request_votes(self):  
  
        ...
```

How would you implement Raft in Ray?

```
def start_election(self):  
    ...  
    majority = len(self.peers)/2+1  
    votes = []  
    for p in self.peers:  
        v = p.request_votes.remote()  
        votes.append(v)  
    ready, _ = ray.wait(votes, majority, election_timeout)
```

```
class Raft(object):  
    def __init__(self):  
        self.peers = N  
        ...  
    def start_election(self):  
        ...  
    def request_votes(self):
```

References

- Ray: <https://ray.io>
 - Documentation
 - Code
 - Tutorials
 - Use cases
- Anyscale: <https://anyscale.com>